

Zephyr™ Project

Developer Summit

June 8-10, 2021 • @ZephyrIoT

Logging subsystem overview

KRZYSZTOF CHRUSCIŃSKI

Agenda

- Variety of requirements
- Deferred mode concept
- Zephyr logging subsystem history
- Version 1
 - Features
 - Weaknesses
- Version 2
 - Challenges
 - Solutions
 - Outlook

Requirements - output

- Printf-like
 - `LOG_ERR("Error: %d", err)`
- Logging data
 - `LOG_HEXDUMP(buf, len, "packet:")`
- Readable
- Prefixed (level, source, function name)
- Timestamp
 - High precision
 - coarse, non-wrapping
 - Custom
- Formatted output vs compressed
 - Ease of use vs resource optimized

- UART console
 - Single direction vs interactive
- Filesystem
- Remote (net, bluetooth)
- Simultaneous multiple transports (*backends*)

Requirements - filtering

- Built-time
 - Global level (e.g. only warnings and errors)
 - Per source (e.g. debugs in i2c, warnings in spi)
 - Per instance (e.g. debugs on uart0, errors on uart1)
- Run-time
 - Per source and instance
 - Per backend

- Closest to non-intrusive
 - String formatting
 - Backends handling
- Log from any context
 - Even highest priority interrupts
 - Synchronous logging takes time, impossible for certain backends
- Critical last logs
- In some configurations we don't care (emulation, unit testing)

Requirements - RTOS

- ROM/RAM requirements
 - Speed
 - Size
 - Both
 - None

- Heavy processing deferred to the known context
- Buffering
 - full policy: saturate vs discard the oldest
- Challenges
 - Deceiving when single stepping
 - Flush before reset (fault, watchdog timeout)
 - Panic context

Panic mode

- Most important data!
- Exception context
- No more scheduler and interrupts
- Backends must switch to synchronous, interrupt-less mode (if possible)

- **Deferred**
 - Any backend
 - Any context
 - Close to non-intrusive
 - Delayed output
 - Build-time, run-time filtering
 - Rich formatting, timestamping
- **Immediate**
 - Intrusive
 - Limited backends supported
 - Immediate output
 - Filtering and formatting
- **Minimal**
 - Redirection to printk
 - No run time filtering, formatting, timestamping
 - Low footprint

- „Test string with 5 parameters %d %d %d %d %d”
 - Formatting **80us** (Cortex-M4 @64MHz)
 - 300us if %lld instead of %d
 - Transfer over UART @115.2k **4ms**
- Deferring
 - Store log message with string pointer, arguments, timestamp, etc. in few microseconds
 - Heavy weight processing in low priority thread (default or user context)

- Log call context
 - (runtime filtering) Check if any backend enables given (source,level)
 - Allocate and enqueue message
- In the known context (idle) for each backend:
 - (runtime filtering) check if backend enables given source,level
 - pass message to a log backend.

Deferred mode – corner cases

- „%s” with transient strings
- Storing arguments for formatting
 - Arguments dump vs va_list
- Offline processing
- Handling full buffer
- 2 processing contexts
 - Default low priority context
 - Panic interrupting default processing context

```
void foo(void)
{
    char name[] = "name"; /* on stack */
    LOG_INF("Name: %s", name);
}
```


- Introduced in June 2018 (version 1)
- Based on logging from Nordic Semiconductor nRF5 SDK
 - Primary bare metal, 32 bit ARM Cortex M only
 - Limited need for logging floats and 64 bit
 - Transient strings logging may be avoided, or limits accepted
 - Asynchronous backend processing to not block the main loop
- While Zephyr scope continues to grow...
- Overhaul merged in April 2021 (version 2)
 - Keeping version 1 as the default one for now

Log message (v1)

- Cast all string arguments to the machine word
 - No support for floating point or long long
- Calculate size needed and allocate message from the pool
- Add timestamp, level and source ID
- Add message to the list

Log message buffer management (v1)

- Pool of fixed size slabs
- Variable length message consists of linked slabs
 - Slab size aligned to fit standard message (up to 2 arguments)
- Head slab contains reference counter
 - Backend gets and puts back the message
 - Last one frees the message
 - Asynchronous backends (no circular buffer)
 - No control when backend puts back the message
- List of buffered messages



Log message buffer management (v1)

- Pros

- Fast allocation and enqueueing (single chunk)
- No 2 contexts problem
- Shared memory between backends

- Cons

- Fragmented
- Hard to self-contain (%s)
- Costly allocation when saturated
- No direct string formatting

%s arguments

- String arguments must be wrapped with *log_strdup()*
- *log_strdup()* copies to a buffer from dedicated memory pool
 - Size of the pool must be tuned to the application
- When message is freed, additional scan is performed to find duplicates and free them

```
void foo(void)
{
    char name[] = "name"; /* on stack */
    LOG_INF("Name: %s", log_strdup(name));
}
```

String formatting (v1)

- Convert fragmented message to `va_list`
- Complex to handle other types than machine word
- Code size (0.7k) and stack usage bloats

```
switch (nargs) {  
    case 0:  
        print(str); break;  
        break;  
    case 1:  
        print(str, args[0]); break;  
    ...  
    case 5:  
        print(str, args[0], args[1],  
              args[2], args[3], args[4]);  
        break;  
    ...  
    case 15:
```

Deferred logging v1 limitations

- *log_strdup()* PITA
- Limited types support
- Fragmented, distributed messages
 - Complex implementation
 - Complicates string formatting, offline processing
 - multi-domain support



Zephyr™ Project
Developer Summit

Time for the overhaul...

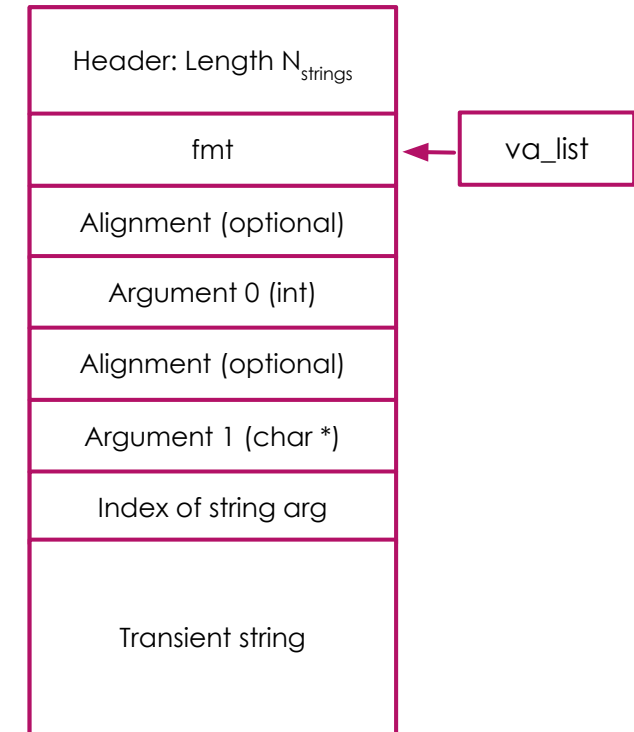
Overhaul goals

- Get rid of `log_strdup()`
- Support all format specifiers
- Self-contained message
- Keep v1 performance
- Others:
 - 64 bit timestamp
 - Simplify implementation
 - Unify standard and hexdump messages

Self-contained string package

- **cbprintf** – Zephyr string formatting library (since 11.2020)
- Way of deferring string formatting contrary to `va_list` or `__VA_ARGS__`
- Package
 - `int create_package(uint8_t *output, char *fmt, va_list ap)`
 - `output` is a buffer with opaque content. Can be moved and copied.
 - `snprintf_from_package(strbuf, sizeof(strbuf), output)`
- Self-contained
 - `fmt` pointer
 - Arguments
 - Any transient strings copied into the package body

- „Test %d %s”
- Argument alignment
- Package body as `va_list` stack frame
- When copying alignment must be maintained
- Processing
 - Update pointers to in-package strings
 - Call standard string formatting



- Walk over format string
- Copy arguments
- Maintain proper alignment
- If string is found, check if read only (rodata linker section)
- Multiple times faster than string formatting but still relatively slow (~20-30us)
- API
 - `int cbprintf_package(package, len, fmt, ...)`
 - `cbpprintf(cbprintf_cb out, void *ctx, void *packaged)`

A missing piece

- Build string package but keep v1 performance
- Hypothesis: String package can be built by assigning arguments to an array but
 - How handle special cases (e.g., float to double promotion)?
 - How to handle %s?
 - How to add proper alignment?
 - How to detect when it cannot be done?

- Introduced in C11
- C response to C++ overloading and templates ;)
- Supported by GCC, Clang

```
#define IS_INT(x) _Generic(x, \
    int: printk("yes"), \
    default: printk("no") \
)
```

_Generic new opportunities

- Static (compile time) string packaging
- Handle corner cases
- Compile-time detection when runtime packaging is necessary

Detect string pointers

- Detect char pointer

```
#define IS_CHAR_PTR(x) \
    _Generic(x, char *: 1, default: 0)
```

- Run it for each argument and sum the results

```
FOR_EACH(IS_CHAR_PTR, (+), __VA_ARGS__)
```

- Will return positive value if any char * in the argument list

```
MUST_RUNTIME_PACKAGE("%p", (char *)p)  
MUST_RUNTIME_PACKAGE("%s", "read only")
```

Static string packaging

- Detects if static packaging can be applied
- Calculates space needed for a package
- Write package
- Fallback to runtime package if `_Generic` not supported
- API:
 - `CBPRINTF_MUST_RUNTIME_PACKAGE(skip, ...)`
 - `CBPRINTF_STATIC_PACKAGE(packaged, inlen, outlen, align_offset, ...)`

Static string packaging use

```
#define STATIC_PACKAGE(...) \
int len; \
if (CBPRINTF_MUST_RUNTIME_PACKAGE(0, __VA_ARGS__)) { \
    len = cbprintf_package(NULL, 0, __VA_ARGS__); \
} else { \
    CBPRINTF_STATIC_PACKAGE(NULL, 0, len, 0, __VA_ARGS__); \
} \
uint8_t buf[len]; \
if (CBPRINTF_MUST_RUNTIME_PACKAGE(0, __VA_ARGS__)) { \
    len = cbprintf_package(buf, len, __VA_ARGS__); \
} else { \
    CBPRINTF_STATIC_PACKAGE(buf, len, len, 0, __VA_ARGS__); \
}
```

- Each static packaging resolves to hundreds of line of C code
- Compiler is smart and resolves it to few simple memory assignments or call to runtime packaging

```
int y = 5;  
int x = 5 - y;  
static int z = 100;  
if (x > 0) {  
    printk("removed %d", z);  
}
```


Compilation

```
void func(int i)
{
    STATIC_PACKAGE("test string: %d", i);
    memcpy(out, buf, len);
}
```

```
0000040E    LDR        R3, =0x00004A10
00000412    STRD       R3, R0, [SP, #8]
00000418    MOVS       R3, #3
0000041C    STR        R3, [SP, #4]
```

```
void func2(int i, uint8_t j, long long ll)
{
    STATIC_PACKAGE("test string: %d %hhu %lld", i, j, ll);
    memcpy(out, buf, len);
}
```

```
0000040E    LDR        R4, =0x00004A20
00000410    SUB        SP, SP, #24
00000412    STRD       R2, R3, [SP, #16]
00000416    STRD       R4, R0, [SP, #4]
0000041C    MOVS       R3, #6
00000424    STR        R3, [SP, #0]
0000041A    STR        R1, [SP, #12]
0000041E    MOVS       R2, #24
00000420    MOV        R1, SP
00000422    LDR        R0, =out
00000426    BL         memcpy
```

- Use *cbprintf* packaging
- At compile time determines message size (except hexdump case)
- Built log message as few memory assignments
- Stores the message
- Message has variable size and must not be fragmented

Log message v2

- Generic message (standard/hexdump)
- +12 bytes (32 bit arch)

Domain ID (3 bits)	Descriptor
Level (3 bits)	
String package length (10 bits)	
Data length (12 bits)	
Timestamp (32 or 64 bits)	
Source (pointer to static or dynamic struct)	
Package (0 – n bytes)	
Data (0 – n bytes)	

Log backend API change

- v1
 - Put(backend, msg)
 - Immediate
 - Put_sync_string()
 - Put_sync_hexdump()
- Separate handling of hexdump and standard messages
- v2
 - Process(backend, msg2)
- Single message type

Circular buffer vs fragmented chunks

- Variable length, continuous space preferable over fragmented chunks
 - Easier to build, move, process
- Thread can block
 - Asynchronous backends not that beneficial
- Circular buffer challenges
 - Efficiency – hard to bit memslab
 - No memcpy: Allocate, Commit, Claim, Free
 - Wrapping – padding needed
 - Multiple producer single consumer
 - ...unless panic
 - No space policy – saturate vs drop the oldest

Multi producer single consumer packet buffer

- No memcpy
 - Allocate, Commit
 - Claim, Free
- Memory optimized: 2 bit header
- User handlers
 - Get packet length
 - Notify about drop

valid	busy	
0	0	Free space
1	0	User packet
1	1	Claimed user packet
0	1	Internal padding

- Improvements
 - Any variable type supported
 - Soon to forget about `log_strdup()`
 - Subsystem code size 6.2k down to 3.5k
 - ~20% better buffer utilization (continuous vs fragmented)
 - Logging from user space 5x faster
- Logging macros
 - 2x bigger (average message size 15 to 32 bytes)
 - `CONFIG_LOG_SPEED=y` bigger (47 bytes) but close to v1 performance

- Self-contained messages opportunities
- Dictionary based logging (#30539)
 - Message send to host for offline processing
 - Log strings might be stripped from the binary (TBD)
- Flash/filesystem backend can be optimized
 - Currently stores formatted strings
- Multidomain logging
 - Multicore system
 - One master core with "real" backend
 - Other cores with ipc backends
 - TODO: String packaging alignment



Zephyr™ Project
Developer Summit

Thank you



ZephyrTM Project

Developer Summit

June 8-10, 2021 ▪ @ZephyrIoT