



Zephyr™ Project

Developer Summit

June 8-10, 2021 • @ZephyrIoT

Weaving Xtensa Yarns with Zephyr

NOTES ON AN IDIOSYNCRATIC
ARCHITECTURE



Zephyr™ Project

Developer Summit

June 8-10, 2021 • @ZephyrIoT

Andy Ross

SOFTWARE ENGINEER – INTEL

andrew.j.ross@intel.com

[andyross](#) @zephyrproject.slack.com

[andyross](#) @github

- Highly Configurable CPU IP from Tensilica, now Cadence
- Supported in Zephyr since 2018
 - Espressif ESP32
 - Intel Audio DSP Family (2013 onwards)
 - ODroid Go (Gameboy Emulator)
- Natural fit for applications with ISA extension needs
 - Vector DSPs
- Licensed as customer-configured IP only
 - No single "CPU" designs

- Traditional pipelined RISC
 - Load/store architecture
 - Single-issue
 - In-order execution
- 16 General Purpose Registers
 - A0-A16
 - A1 is stack pointer by convention
- 32 interrupts vectored to configurable (e.g. 7) number of priority levels
 - Fixed priority for each interrupt
 - Level 1 shared with exception processing

- Configurable "State Registers"
 - Move to/from GPRs with RSR/WSR instructions
 - Named with one byte in the instruction, space of 256 possible SRs
 - Easily enumerable by the OS for context management
 - (Mostly, the assembler makes this harder than it has to be)
 - Existence exposed by system headers (core-isa.h)
 - Examples:
 - PS – Program state register
 - LCOUNT/LBEG/LEND - LOOP instruction extension
 - SCOMPARE – Atomic test/set instruction operand
 - SAR – Shift and rotate instruction operand
- One toolchain per instantiated CPU design!

Xtensa ISA, weirder stuff...

- Register Windows!
 - 16 visible GPRs, but 16, 32 or 64 actual registers
 - Rotated in groups of 4 registers upon function call
 - Call can "save" 4, 8 or 12 of its own registers
 - (Or zero, which is a variant "CALL0" ABI for configurations with only 16 registers)
 - Cyclic with spill/fill managed by exception handlers
 - Trying to return to registers that are not present traps to fill them from the stack
 - Trying to use a GPR in a register block "owned" by an upthread function traps to spill
 - Trap happens at use time, not call time!
 - Faster function calls, in theory
 - In practice, generated code tends to be poor
 - "How many registers to use" is a global optimization (depends on callee behavior)
 - But register assignment happens locally
 - Upshot: code uses way too much stack for calls
 - Xtensa's proprietary compiler backend does a little better than gcc
 - Details very clever, but very complicated

- Hardware interrupt entry is mostly conventional
 - Interrupts jump automatically to one of N handlers
 - Mask interrupts unconditionally
 - Set EPCn to the interrupted program counter, EPSn to the interrupted state register
 - Don't touch GPR or window state
 - Provides EXCSAVEn for use as a scratch register
 - Zephyr doesn't use it!
- Saving context
 - We have to switch stacks (Zephyr has a separate interrupt stack)
 - The register windows assume a single stack
 - Oh well, must spill all 64 GPRs on interrupt
 - Or must we...

- Carefully construct a first call frame containing the interrupted stack frame
 - Will be restored on return when unspilled
- Push a second 4-register frame
 - Acts as a spill area for the new stack that won't clobber the first
- Forcibly set the stack pointer (old one will be restored on return)
- Call the interrupt handler (which is a standard C function)
- No registers need to be explicitly spilled at all!
 - Everything happens via natural windows exceptions
 - Interrupt entry overhead is on the order of two nested functions
 - If the interrupt handler is shallow, interrupted GPRs may never be spilled at all
 - Takes most overhead on interrupt exit, not entry

- One C function to handle each interrupt level
- Multiple interrupts per level
 - (Also level 1 includes exceptions)
- Of the 32 interrupts, each belongs to exactly one level
- Generate optimized C code with `xtensa_integen.py`!
 - Bit test only the interrupts that are possible at each level
 - Call the handler out of the Zephyr isr table
- Not as great as I would have hoped...
 - Function call overhead is expensive
 - Xtensa bit handling surprisingly slow
 - Most hardware turns out to have a second level interrupt controller anyway

- Call `z_get_next_switch_handle()` to find next context to run
- Common case: returning into same thread
 - No need to unspill windowed registers, just the interrupted GPRs
 - Even those will only have spilled the registers that were in use
 - Restore EPCn/EPSn and issue RFI to return
 - Zephyr RFI path is shared, always uses the same (highest) level for RFI
- Uncommon case: preemption
 - Leave the spilled interrupt context on the stack in place
 - But all the caller window frames need to be spilled! ...

Zephyr Window Spill

- Provided by HAL layer
 - Big (few hundred bytes of code)
 - Software-only solution, interprets window state.
- Zephyr wrote our own
 - Relies on exception handlers
 - Just poke registers to spill them
 - Tiny (-----> !!!)
 - Actually about twice as fast!
 - (Hardware exception processing is very speedy, use hardware)
- Also useful for exception handling
 - Debugger integration

```
/* And this is it! */  
.macro SPILL_ALL_WINDOWS  
#if XCHAL_NUM_AREGS == 64  
    and a12, a12, a12  
    rotw 3  
    and a12, a12, a12  
    rotw 3  
    and a12, a12, a12  
    rotw 3  
    and a12, a12, a12  
    rotw 3  
    and a12, a12, a12  
    rotw 4  
#elif XCHAL_NUM_AREGS == 32  
    and a12, a12, a12  
    rotw 3  
    and a12, a12, a12  
    rotw 3  
    and a4, a4, a4  
    rotw 2  
#endif  
.endm
```

- Xtensa cores often deployed in SMP configurations
 - Shared memory bus, same address space
 - ESP32: two cores, second dedicated to network in upstream use cases
 - Intel Audio DSP: 2-4 core configurations shipping
 - No MMU (yet)
- First Zephyr SMP Platform
 - Prototyped most of the SMP APIs
 - Made most of the early mistakes

- "What CPU am I?"
 - Kernel needs fast (!) access to per-CPU struct
 - Contains `_current` thread pointer
 - Use a dedicated SR
 - Must be "privileged" (for when we get an MMU) so user code can't mess it up
 - Not saved by context management, set once at boot
 - No architecture-specified register for this
 - Indirect via `CONFIG_XTENSA_KERNEL_CPU_PTR_SR`
 - On esp32, uses "MISC0", which is an excellent choice
 - No MISC on Intel, use EXCSAVE2 instead
 - Remember Zephyr doesn't use it! Nice to be frugal.
 - Maybe SR should store the thread pointer instead?

SMP Requirements...

- "Start CPU n"
 - No architectural interface for this at all.
 - ESP32 provides a HAL call
 - Intel DSP does this via the x86 host CPU!
 - Requires driver protocol, working this out with SOF right now
- "Interrupt other CPUs" (`arch_sched_ipi()`)
 - Needed for scheduling, to inform CPUs of changes to the set of runnable threads
 - Again, no architectural interface
 - Intel DSP has custom "IPC" (Inter-processor communication) hardware
 - Doorbell call/response registers
 - ESP32 just can't do it at all
 - At least not via a documented API

- Atomic memory access API
 - Core primitive from which spinlocks are built
 - Generally built on top of a single instruction "compare and swap"
 - Which xtensa has: S32C1I
 - But the compare operand lives in the SCOMPARE SR!
 - Led to a subtle bug: atomics are usually used in spinlocks, but not always
 - If SCOMPARE isn't saved with context, we can be interrupted before its used!
 - We went 2 years without SCOMPARE saved in context

- Xtensa often used in shared/slow/contended memory environments
- For example, the Intel Audio DSP:
 - Has two SRAM regions, one slower (and lower power) than the other
 - Also system DRAM access, which is VERY slow
 - Also a "TLB" translation lookup interposed in front of the SRAM access eating cycles
- So the CPU can be configured with L1 data and instruction caches
 - Typical 64 byte cache line size
 - Configurably N-way set associative
 - Configurable size, e.g. 16 kb
 - Straightforward writeback/flush/invalidate control via dedicated instructions
 - **AND IT IS INCOHERENT**

- Cache incoherence breaks basically all memory reasoning
 - Changes to memory from CPU0 stay in cache and aren't visible to CPU1
 - Changes from CPU1 can get flushed at any time and clobber CPU0 unexpectedly
 - Changes to "unshared" memory merely in the same cache line can be clobbered too
 - Even uncolliding changes may appear to other CPUs in any order
 - **EVERYING** in the kernel assumes universal visibility
- How to fix this?
- Idea #1: Don't
 - Disable the cache
 - Guarantees correctness at expense of performance
 - Works well if SRAM is fast relative to CPUs. ESP32 does this
 - Intel ADSP takes 12 cycles to read from SRAM!

- Idea #2: Harden the kernel
- Make all kernel code incoherence safe.
 - We already know when synchronization (spinlocks) happens, so leverage that:
 - Flush and invalidate the whole cache when lock is taken
 - Flush when lock is released
- This makes non-synchronizing CPU-bound code fast
 - But it does nothing for the kernel itself, which is actually much slower than Idea #1
 - Flushing the whole cache takes 256 instructions if done naively
- Still does nothing for the poor application, which has all the same issues

- Idea #3: CONFIG_KERNEL_COHERENCE
- Note one hardware trick:
 - The Xtensa memory protection option allows associating different cache behavior with different 512MB memory regions
 - The Intel ADSP maps addressible RAM twice, once cached and once uncached
- Recognize that almost all kernel data is pervasively shared
 - CPU structs, Thread records, Timeout objects, IPC primitives
 - All are designed to be used from any thread at any time
 - They really can't be cached anyway
 - Put them in uncached memory!

- Recognize that almost all stack accesses are intended to be unshared
 - And the stack makes up 80%+ of memory accesses in common code!
 - Put the stack in cached memory
- Do all this in the linker.
 - `__coherent` and `__incoherent` attributes for user-defined symbols
 - Make default data coherent (uncached) by default, for correctness
- Add runtime checks to catch errors
 - Code may want to place e.g. IPC objects on the stack, and it's legal elsewhere!
 - Means I have to play "coherence police" with regressions in the test suite

- What about context switch? This gets complicated
 - Must spill register windows first! (remember those?)
 - Need to flush out the old thread's stack, it may not run on this CPU next
 - Need to invalidate the new thread's stack, we may not have been the last to run it
 - Don't need to flush the dead area below the top of the old thread, it's not used
 - Don't forget the caller's spill region below the current stack pointer!
- This gets handled by `arch_cohere_stacks()` inside `z_swap()` and interrupt exit
 - But we can't do it all in C, because we're still running on the old stack.
 - Does most of the logic, but leaves the flush to assembly called immediately before switch
 - Stashes the computed pointer in EXCSAVE3 (useful again to have extra registers)

- Future optimizations:
- Some things in kernel data are unshared, or read-only, or change rarely
 - Ex.: CPU ID number never changes
 - Thread priority changes require an IPI and are heavyweight anyway
 - IPC configuration data (semaphore max size) only change at initialization
- Would be nice to have a framework to separate these

- Current Zephyr Xtensa platforms don't provide MPU/MMU hardware
 - Yet...
- Cadence offers both as options
 - Both are "oddly clever" as is Xtensa custom
 - Both are good targets for Zephyr userspace implementations
 - Neither are implemented in Zephyr
 - Yet...
- Xtensa MPU:
 - No address translation
 - 4k aligned region boundaries
 - 16 or 32 region registers partition all 4G of memory space
 - Means there are 8 or 16 possible "useful" regions, plus one for each two that are adjacent

- XtenSA MMU
 - Conventional paged memory translation unit, mostly
 - Read/write/execute and cacheability control per page
 - Set-associative TLB with configurable size
 - Also some special TLB ways with specific behavior for e.g. pinning and large page size
 - Software exception used to fill TLB on a page table miss
 - Hardware engine fills TLB from page table normally
 - Page table is a linear array in **virtual address space**
 - Means it has to map itself
 - But also means hardware has no special paths, or state machine
 - Software handling (per above) handles bootstrapping of paging

Future: Zephyr Userspace Requirements

- Syscalls: already have exception code dedicated, and instruction to trigger
 - Just need to decide on marshalling scheme for arguments
- Interrupts: Can't rely on interrupted user stack
 - Might be too small, and fault (!USERSPACE just overflows)
 - Might also be deliberately misconfigured with a kernel address!
 - Must switch stacks before context save
 - Complicates cross-stack calls
 - Maybe do only when kernel threads interrupted?
- Security: need to scrub kernel registers in windows
 - Not all Zephyr archs do this currently
 - Force spill and zero? (similar code to above, with writes)



Zephyr™ Project
Developer Summit

“It is one of man's curious idiosyncrasies
to create difficulties for the pleasure of
resolving them.”
- Joseph de Maistre



ZephyrTM Project

Developer Summit

June 8-10, 2021 ▪ @ZephyrIoT