



Zephyr™ Project

Developer Summit

June 8-10, 2021 • @ZephyrIoT

Zephyr Power Management - 101

FLAVIO CEOLIN

FLAVIO.CEOLIN@INTEL.COM



Zephyr™ Project
Developer Summit

“Power management is a feature that turns off the power or switches the system to a low-power state when inactive”

- Goals
- Power Management on Zephyr
 - System Power Management
 - Device Power Management
- Enabling
 - How to use it in a target
 - How to implement or customize Power Management
- Conclusions


- Consume **as little power as possible** for a given **use case** and **configuration**.
 - Do not waste energy when **idle**.
- **Customizable** - provide the mechanism, not the policy.
- Cross platform (architecture / board / SoC).
- Convenient to **use** and **deploy**.



Zephyr™ Project
Developer Summit

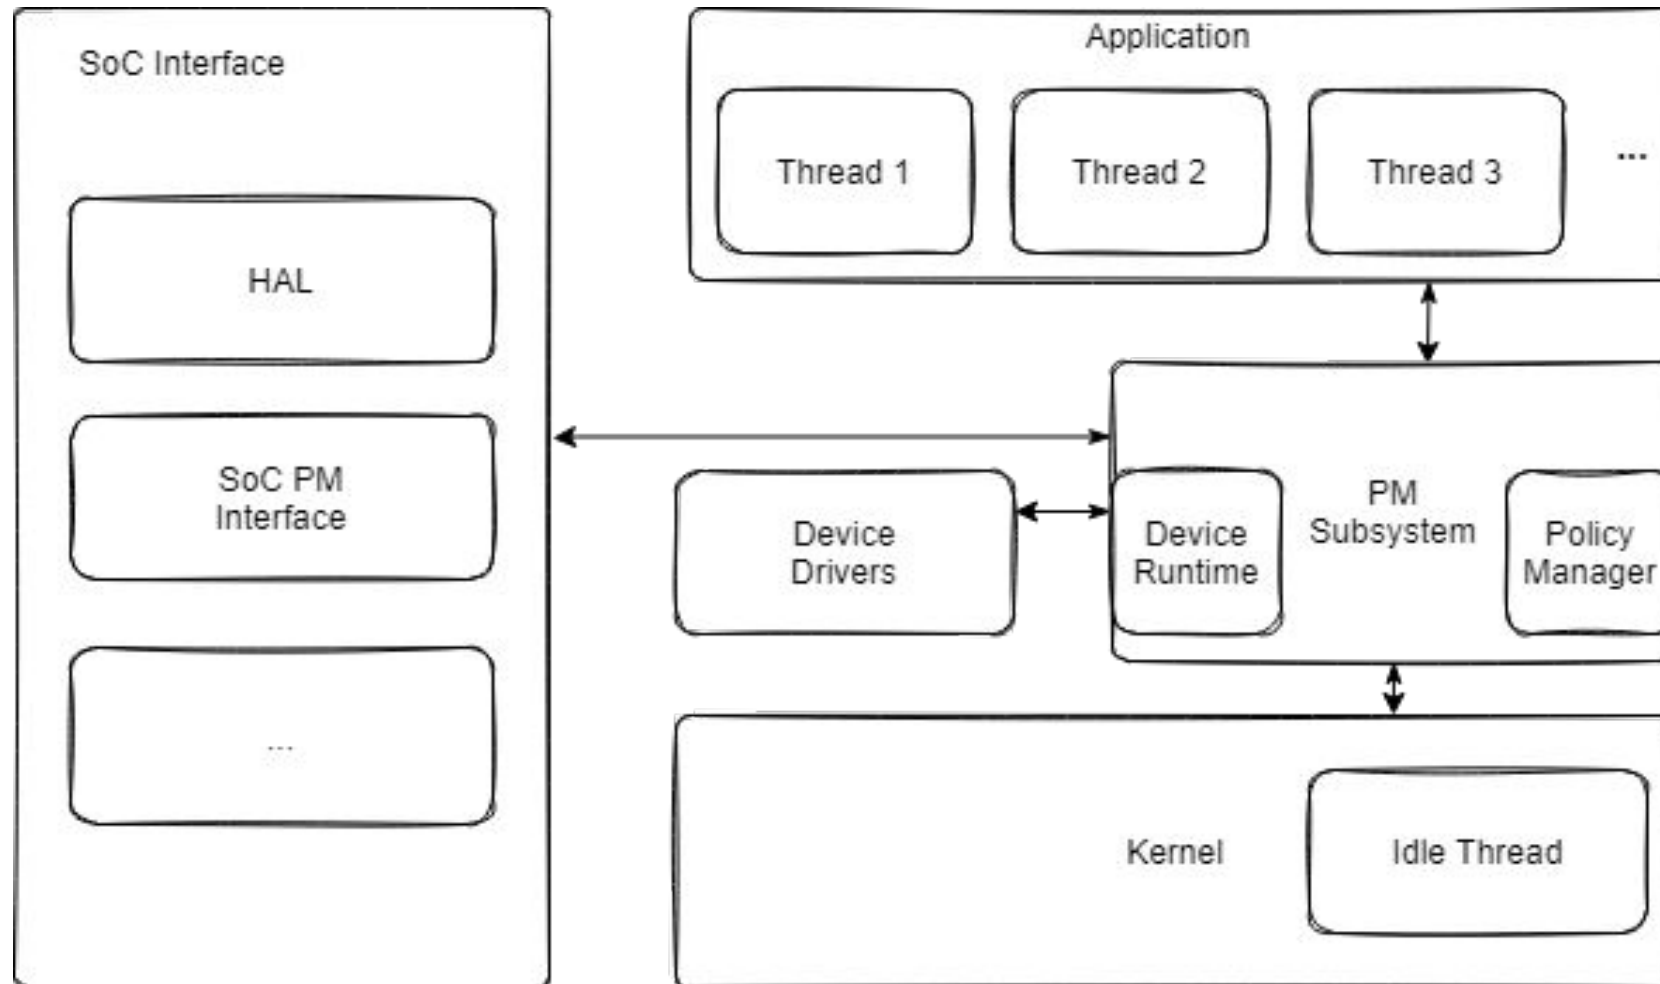
Power Management on Zephyr

Zephyr Components



- Kernel **Idling**
- **System** PM - Saves energy changing the SoC power state
 - **Policy Manager**
 - Power state **constraints**
- **Device** PM – Saves energy suspending devices
 - **Runtime** PM
 - Device busy / idle

PM High-level Overview



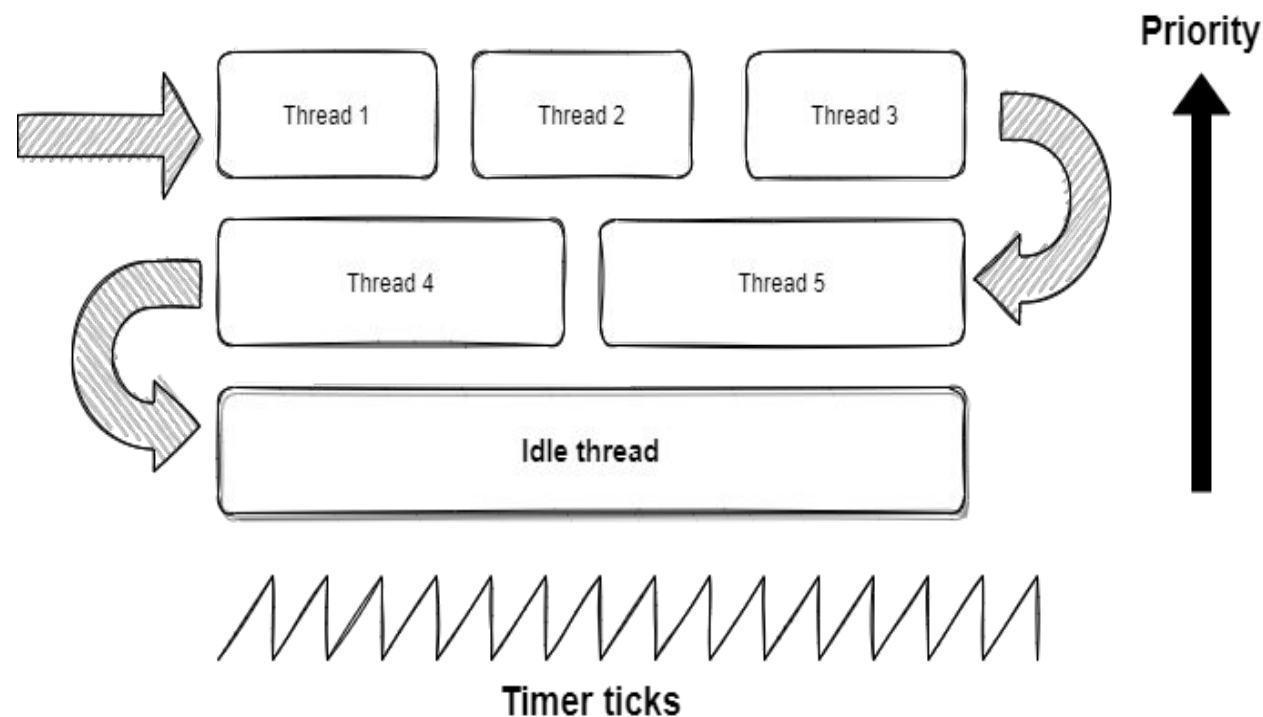


Zephyr™ Project
Developer Summit

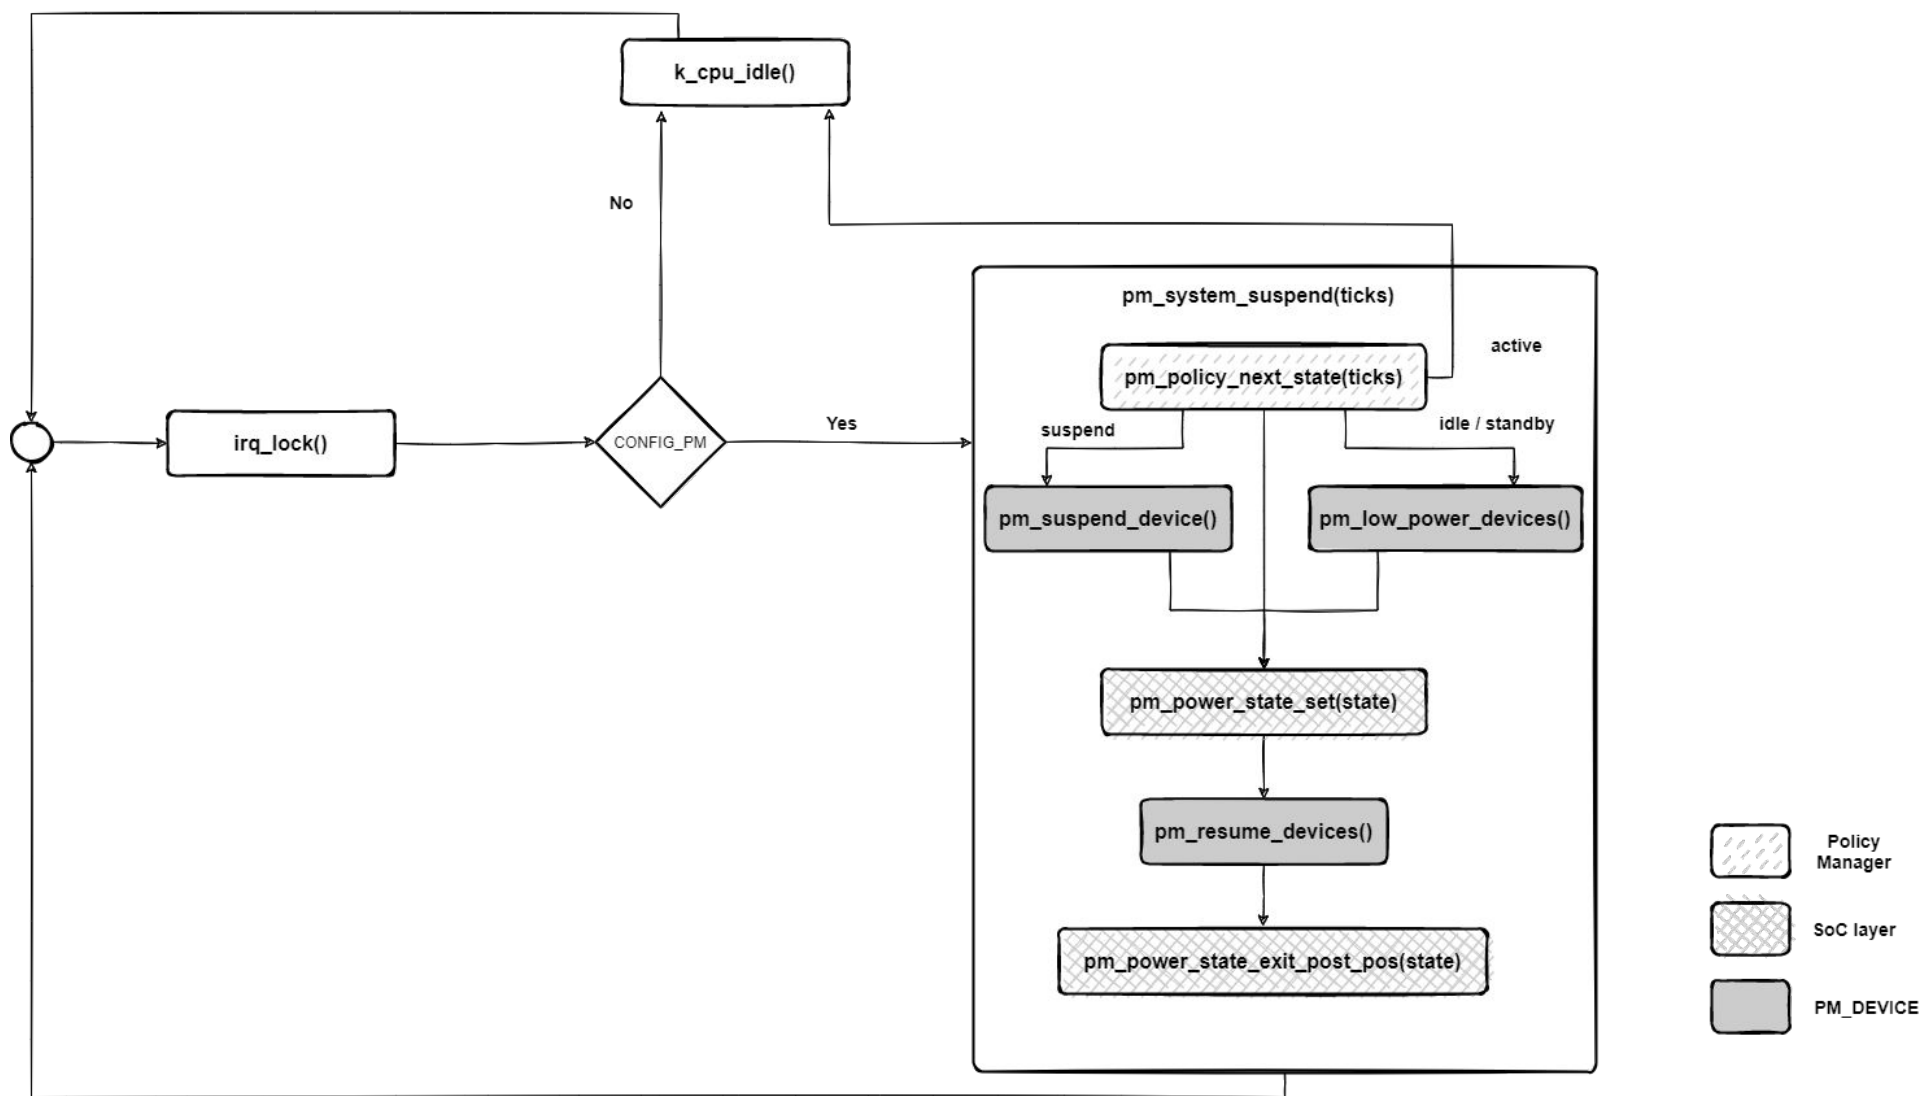
Power Management on Zephyr (System PM)

System IDLE

- Governed by the **idle** thread
 - Idle thread is the lowest priority thread and is executed when no other thread can run
 - Idle thread is scheduled again if no other thread is ready to run
- Tickless kernel
 - Interruptions only for registered events



Inside Idle thread



- States are **ACPI alike** and defined in devicetree
 - PM_STATE_ACTIVE
 - PM_STATE_RUNTIME_IDLE
 - PM_STATE_SUSPEND_TO_IDLE
 - PM_STATE_STANDBY
 - PM_STATE_SUSPEND_TO_RAM
 - PM_STATE_SUSPEND_TO_DISK
 - PM_STATE_SOFT_OFF
- Targets can declare only states that are meaningful for them
- It is fully customizable

```
compatible: "zephyr,power-state"
```

```
properties:
```

```
    power-state-name:
```

```
        type: string
```

```
        required: true
```

```
        description: indicates a power state
```

```
        enum:
```

```
            - "active"
```

```
            - "runtime-idle"
```

```
            - "suspend-to-idle"
```

```
            - "standby"
```

```
            - "suspend-to-ram"
```

```
            - "suspend-to-disk"
```

```
            - "soft-off"
```

```
    substate-id:
```

```
        type: int
```

```
        required: false
```

```
        description: Platform specific identification.
```

```
    min-residency-us:
```

```
        type: int
```

```
        required: false
```

```
        description: |
```

```
            Minimum residency duration in microseconds. It is the minimum time for a  
            given idle state to be worthwhile energywise. It includes the time to enter  
            in this state.
```

```
    exit-latency-us:
```

```
        type: int
```

```
        required: false
```

```
        description: |
```

```
            Worst case latency in microseconds required to exit the idle state.
```

- Choose the best state to the system for a given idle period
- Applications can define their own policy
- Zephyr comes with residency policy as default

```
struct pm_state_info pm_policy_next_state(int32_t ticks);
```

- Applications can set constraints in power states to inhibit the system to go to a particular state
- Devices are suspended and resumed **sequentially and synchronous** according to their dependencies
 - Right now the **time** spent in this task **is no accounted**
- Devices set busy are not suspended
- The system accounts the time necessary to the system wakes
- Notifications about power state changes are delivered just before and after the system changes the state
- Application is responsible to set up a wake up event



Zephyr™ Project
Developer Summit

Power Management on Zephyr (Device PM)

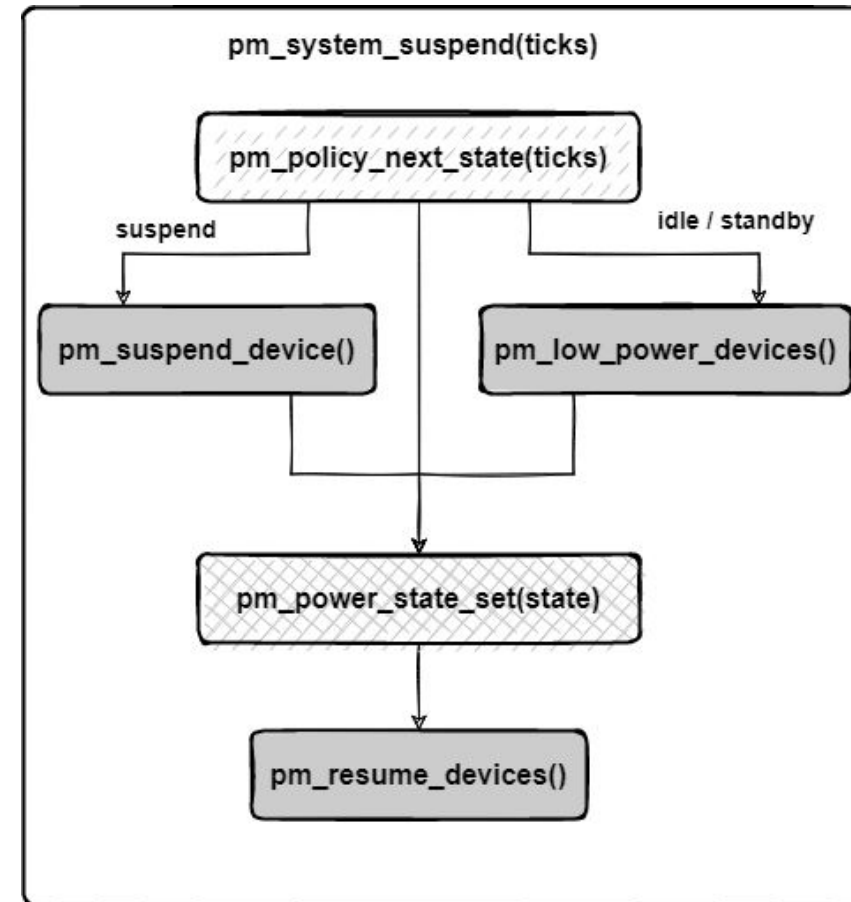
- There are two kinds of device power management:
 1. Central - The system suspends or resumes devices during system power management
 2. Runtime - Devices are suspended or resumed automatically according with their usage
- Devices are suspended and resumed according with their dependencies
- Drivers maintain power states per device
- Device driver PM interface
 - IOCTL style function implemented by drivers
 - Control codes `PM_DEVICE_STATE_SET` and `PM_DEVICE_STATE_GET`

Central Device PM

```
void pm_suspend|resume_devices(void) {  
    /*  
     * iterate over devices and check if they  
     * are not busy  
     */  
    ...  
  
    /*  
     * Wrapper for device driver interface.  
     * device->pm_control(device, PM_DEVICE_STATE_SET ...)  
     */  
    pm_device_state_set (device, state, NULL, NULL);  
    ...  
}
```

```
/* device states */
```

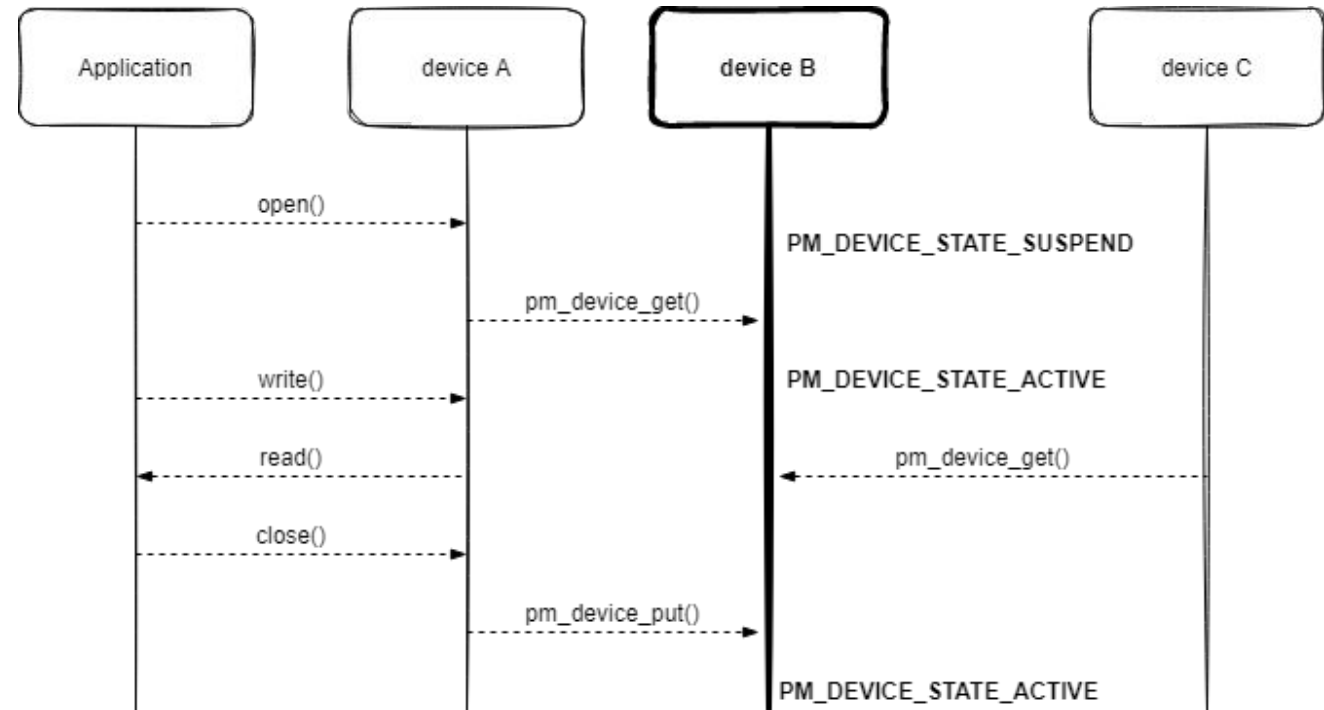
```
PM_DEVICE_STATE_ACTIVE  
PM_DEVICE_STATE_LOW_POWER  
PM_DEVICE_STATE_SUSPEND  
PM_DEVICE_STATE_FORCE_SUSPEND  
PM_DEVICE_STATE_OFF  
PM_DEVICE_STATE_RESUMING  
PM_DEVICE_STATE_SUSPENDING
```



Device runtime PM

- Devices are suspend / resumed in runtime
 - Independent of the system PM
 - System does need to be IDLE !
- Synchronous and Asynchronous APIs

```
int pm_device_get_async(const struct device *dev);  
int pm_device_get(const struct device *dev);  
  
int pm_device_put(const struct device *dev);  
int pm_device_put_async(const struct device *dev);
```





Zephyr™ Project
Developer Summit

Enabling PM

Using PM

- CONFIG_PM
 - Enable System PM.
- CONFIG_PM_DEVICE
 - Enable suspend and resume devices when the system is IDLE.
- CONFIG_PM_DEVICE_RUNTIME
 - Devices suspend / resumed in runtime
- They can be used together !

```
(Top) → Sub Systems and OS Services → Power Management
Zephyr Kernel Configuration
[*] System Power management
-- System Power Management --->
- - <HAS_NO_SYS_PM>
-- Device power management
[*] Device Power Management
[*] Runtime Device Power Management

show-all mode enabled
[Space/Enter] Toggle/enter  [ESC] Leave menu          [S] Save
[O] Load                  [?] Symbol info          [/] Jump to symbol
[F] Toggle show-help mode  [C] Toggle show-name mode [A] Toggle show-al
[Q] Quit (prompts for save) [D] Save minimal config (advanced)
```

Implementing PM – SoC layer

- Implement SoC hooks:

```
void pm_power_state_set(struct pm_state_info info);
```

- This function is called with interruptions locked and with the power state selected by the policy manager.

```
void pm_power_state_exit_post_ops(struct pm_state_info info);
```

```
void pm_power_state_exit_post_ops(struct pm_state_info info);
```

- It is called in the end of the IDLE thread loop. Zephyr expects interruptions to be restored after this function.

- Declare supported power states. E.g:

```
power-states {  
    idle: idle {  
        compatible = "zephyr,power-state";  
        power-state-name = "suspend-to-idle";  
        min-residency-us = <1000>;  
    };  
  
    standby: standby {  
        compatible = "zephyr,power-state";  
        power-state-name = "standby";  
        min-residency-us = <5000>;  
    };  
}
```

Implementing PM - Device

- Device drivers must implement

```
int (*pm_control)(const struct device *dev, uint32_t command,  
                  uint32_t *state, pm_device_cb cb, void *arg);
```

```
#define DEVICE_DEFINE(dev_name, drv_name, init_fn, pm_control_fn, \  
                      data_ptr, cfg_ptr, level, prio, api_ptr) \  
    Z_DEVICE_DEFINE(DT_INVALID_ID_NODE, dev_name, drv_name, init_fn, \  
                    pm_control_fn, \  
                    data_ptr, cfg_ptr, level, prio, api_ptr)
```

```
static int dummy_device_pm_ctrl(const struct device *dev,  
                                uint32_t ctrl_command,  
                                uint32_t *state, pm_device_cb cb, void *arg)  
{  
    int ret = 0;  
  
    switch (ctrl_command) {  
    case PM_DEVICE_STATE_SET:  
        if (*state == PM_DEVICE_STATE_ACTIVE) {  
            ret = dummy_resume_from_suspend(dev);  
        } else {  
            ret = dummy_suspend(dev);  
        }  
        break;  
    case PM_DEVICE_STATE_GET:  
        *state = dummy_get_power_state(dev);  
        break;  
    default:  
        ret = -EINVAL;  
    }  
  
    cb(dev, ret, state, arg);  
  
    return ret;  
}  
  
DEVICE_DEFINE(dummy_driver, DUMMY_DRIVER_NAME, &dummy_init,  
              dummy_device_pm_ctrl, NULL, NULL, APPLICATION,  
              CONFIG_KERNEL_INIT_PRIORITY_DEFAULT, &funcs);
```

Implementing - Policy Manager

- The policy just need to implement the following function:

```
/**  
 * Function to get the next PM state based on the ticks  
 *  
 * ticks -> the number of ticks to next scheduled event  
 */  
struct pm_state_info pm_policy_next_state(int32_t ticks);
```

- The state returned by the policy manager will be used by the system to figure out whether devices must be suspended and is passed to **SoC layer**
 - There are helpers to get power states declared on **Devicetree**




Zephyr™ Project
Developer Summit

Conclusions

- Highly customizable. Applications can tune every single bit on Zephyr PM:
 1. Customize power states
 2. Redefine the policy, defining its own or through constraints usage
 3. Override SoC implementation
 4. Enable / Disable device PM
- Cross platform architecture/board/SoC → great portability
- Offers great savings when combined
 - Still **few devices support PM** though !

Conclusions

```
./samples/boards/stm32/power_mgmt/blinkystm32l562e_dk
```





Zephyr™ Project
Developer Summit

Thank you !

Questions ?