



**Zephyr™**  
Project

Developer Summit

*Zephyros*

*Zephyros*

*Zephyros*

# Multi-domain Logging

Krzysztof Chruściński

# Ordered, human-readable logs from multiple domains on a single output

with low CPUs load

# Agenda

- Assumption, Preconditions, Challenges
- Single domain logging recap
- Multi-domain architecture
- Log message
  - Identification
  - Formatting
  - Ordering
- Potential improvements
- Other options



Krzysztof Chruściński

- Contributing to Zephyr since 2018
- Nordic Semiconductor
- Logging subsystem maintainer
- Other areas: shell, drivers, kernel

# Assumptions & Preconditions

- Domain = binary product (e.g., ARM TrustZone – 2 domains on a single CPU)
- Logging available from any context, ideally non-intrusive
- String formatting
  - Time consuming
  - Depends on string length and type and number of arguments
- Sending data over backend
  - Time consuming
  - Limited to thread context?
- Domain memory access limitations

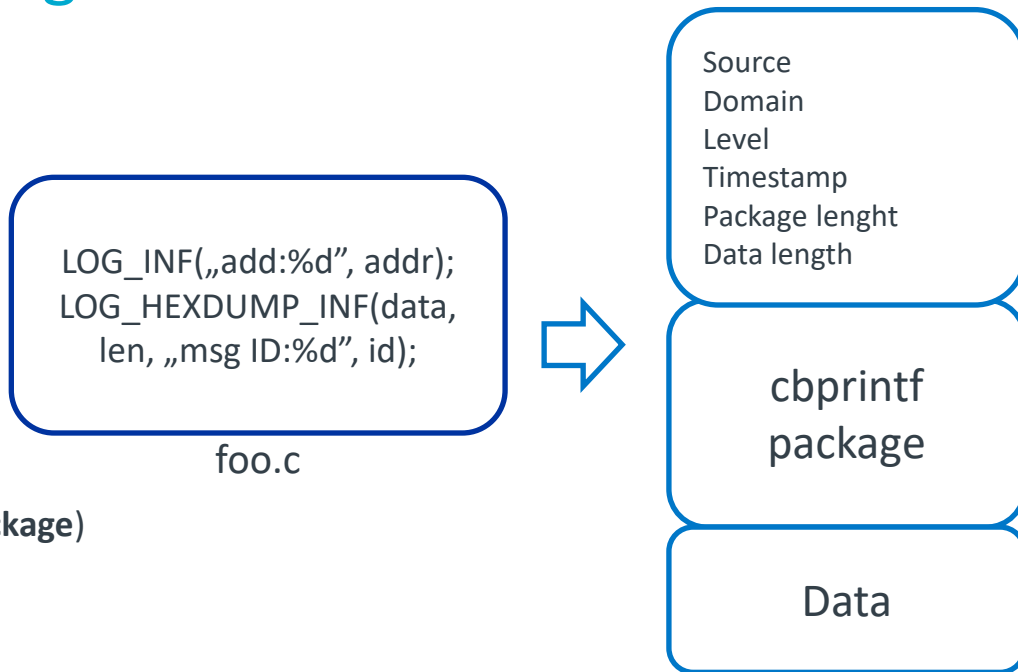
**LOG\_INF**(„Johann Gambolputty de von Ausfern-schplenden-schlitter-crasssrenbon-fried-digger-dingle-dangle-dongle-dungle-burstein-von-knacker-thrasher-apple-banger-horowitz-ticolensic-grander-knotty-spelltinkle-grandlich-grumblemeyer-spelterwasser-kurstlich-himbleeisen-bahnwagen-gutenabend-bitte-ein-nürnbürger-bratwustle-gerspurten-mitzweimache-luber-hunds-fut-gumberaber-shönendanker-kalbsfleisch-mittler-aucher von Hautkopft of %s“, „Ulm“)

# Challenges

- Performance
  - Processing and IPC communication
- Synchronization
  - Common timestamp source
- Ordering
  - by time of generation
- User friendliness
  - Same look&feel as for single domain

# Logging recap – message content

- Who?
  - Source ID, Domain ID
- When?
  - Timestamp
- What?
  - Level
  - \*String with arguments (**cbprintf package**)
  - \*Data (for hexdump messages)



# Cbprintf package – formatted string snapshot

- `Printflike(„foo %s %d”, str, val)`
  - „foo %s” – Variable in RO data. Pointer.
  - `str` – variable on stack. Pointer
  - `val` - Value
- Statically create using C11 `_Generic` keyword
  - Format string (*fmt*) not touched

```
#define FOO(x) _Generic(x, \
    int: printk(„int”), \
    default: printk(„else”))
```

**Header: `va_list` size, number of attached strings**

`fmt`

`*va_list frame`


`*RW strings prefixed with location index`

# Deferred logging in three stages

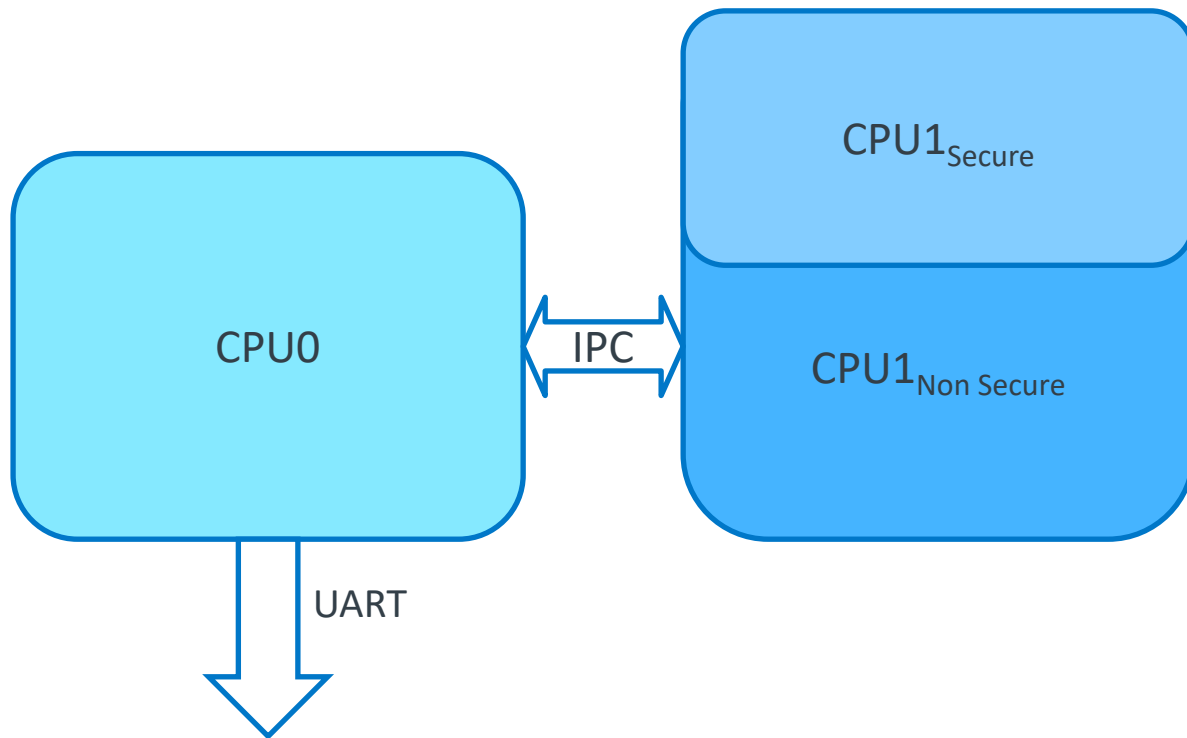
- Message creation
  - Any context
  - Create message with timestamp, allocate and commit
- Message processing
  - Known, low priority context
  - No real time requirements
  - Get message and pass it to all\* enabled backends
- Backend processing
  - \*Converting to human-readable string
  - Transferring (UART, flash, etc.)



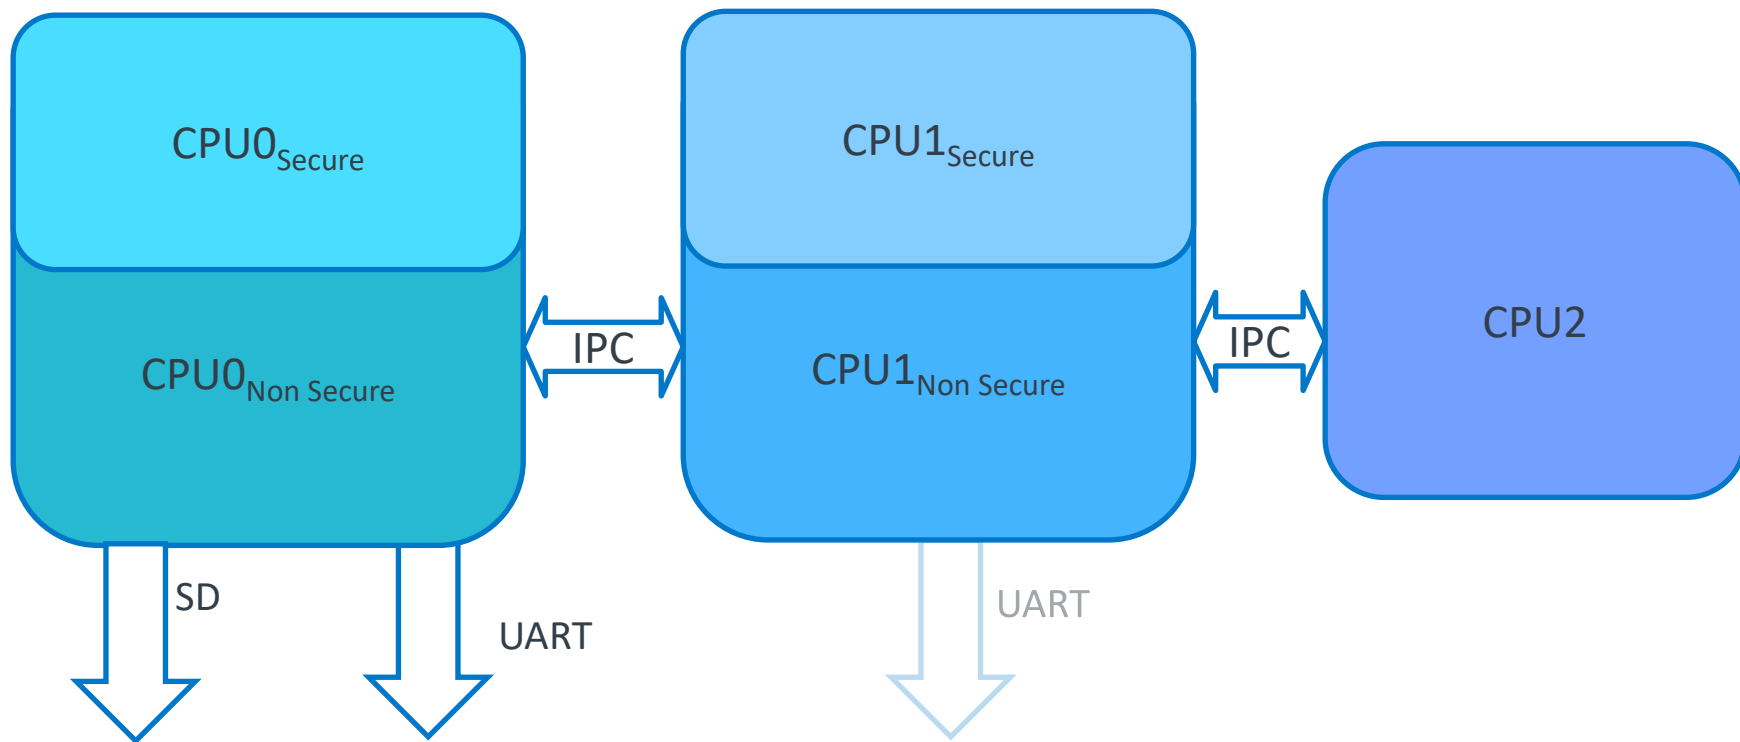
# Multi-domain architecture



# Multi-domain architecture




# Multi-domain architecture



# Logging Link

- Pairs with log backend
- Logging backend passes the message
- Link receives message from backend and enqueue locally

# Multi-domain architecture



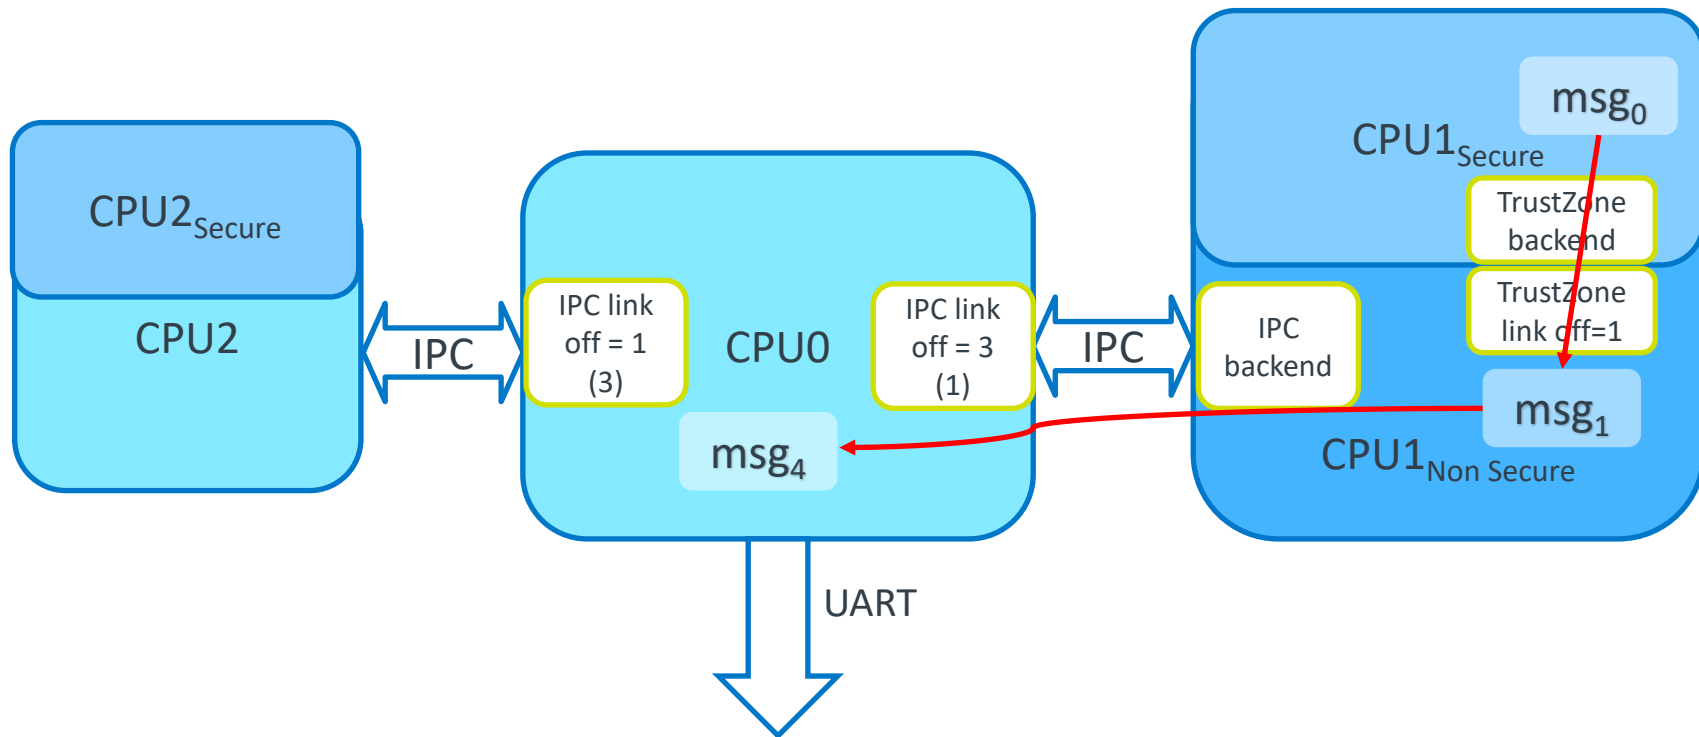
# DomainID addressing

- Static addressing
  - Domain ID assigned in Kconfig
  - More maintenance
  - Troublesome on-chip operations
- Dynamic addressing
  - Assigned during initialization
  - Less configuration maintenance but longer initialization
  - Link API for getting number of domains
- Dynamic approach used

# Dynamic domainID

- Initialization
  - For each link get number of domains
  - $\text{Link}_0 \text{ offset} = 1$
  - $\text{Link}_N \text{ offset} = \text{Link}_{N-1} \text{ offset} + \text{Link}_{N-1} \text{ domain count}$
- New message in link
  - Add link offset to the domainID
- Getting information from domainID
  - Decrement link offset before the call

# domainID





# Message formatting

- Adding prefixes:
  - Timestamp
  - Domain name
  - Source name !
  - Level
- String formatting
  - From cbprintf package !

# Getting source name

- `Get_source_name(domain_id, source_id)`
- How to reduce IPC communication?
- Cache
  - Dedicate memory
  - Store n most recently used `domain_id`, `source_id`
  - When entry not found in cache, evict the oldest and get new from link
- Alternative (not implemented)
  - Get number of all sources
  - Dynamically allocate space and fetch the names

# Fully self-contained package

## Transient

Size of va\_list

fmt

va\_list

## Self-contained

Size of va\_list

Number of rw strings

fmt

va\_list

RW strings (with locations)

## Fully self-contained



Size of va\_list

Number of rw+ro strings

fmt

va\_list

Strings (with locations)

# Fully self-contained package



# Conversion

- Conversion is simple:
  - Read address under location
  - Append string under that location to the package
- `LOG_INF(„lorem %s %s”, „ipsum”, p_char)`
  - Statically created transient package
  - In the context of log message converted to self-contained
  - In the IPC backend converted to fully self-contained

# Multi-domain logging

- Message received from the link
  - Is self contained
  - Source can be retrieved from cache or remote domain
  - Timestamp is in sync
  - But ...
  - Messages will be processed in the order of arrival to the final destination

# Ordering

- Message buffer
  - Circular
  - Messages can only be processed in the order of arrival
  - Solution?
- Buffer for each link (+ local)
  - Peek each buffer – process the oldest message
  - At RAM cost
  - Still unordered?
  - Delay processing by the maximum latency


# PR & Potential improvements

- *PR: #43797*
- *Fmt/source name* in the memory section accessible by the final destination
  - Not accessed anywhere else (if no formatting backends used in origin domain)
  - Fmt is redundant in the package
  - Requires linker tricks



# ARM Coresight - HW solution example

- For each core set of memory mapped register interfaces (STM System Trace Macrocell)
- Writing to registers generates stream of data
- Merged stream
  - Stored in RAM (ETR)
  - Send over the debug interface (TPIU)



# Thank you

Questions?