

Overview of Logging

- Aastha Grover, Intel Corporation

Agenda

- Overview
- Log API's
- Architecture of Logging Subsystem
- Log Message and Memory
- Log Filtering
- Modes of Logging
- Multi Domain Logging
- Types of Logging formats: pros, cons
- Switch Formats at Runtime



Aastha Grover

aasthagr

Intel Contributor to Zephyr Project.

Overview

- Logging outputs Human readable strings.
 - printf like results + more.
 - String formatting – cbprintf package
 - Supports all format specifiers.
- Identify source of Logging
 - Source module, severity level, time of generation (Timestamp)
- Timestamping
- Can log from any context – Interrupt context
- Functionality to dump data
 - LOG_HEXDUMP
- Log API for every Severity levels.
 - Information, Debugging, Warning & Error
- Supports Multiple output channels. Example:
 - UART
 - Filesystem,
 - Remote (Bluetooth, Net)
 - Multiple backends

How to Log using Zephyr?

Source Code:

```
#include <zephyr/logging/log.h>

LOG_MODULE_REGISTER(MODULE, Level);

LOG_INF("String");
```

Kconfig:

```
CONFIG_LOG=y
```

Example: Hello World Application
zephyr/samples/hello_world

```
#include <zephyr/kernel.h>
#include <zephyr/logging/log.h>

LOG_MODULE_REGISTER(hello_world, LOG_LEVEL_INF);

int main(void)
{
    printk("Hello World! %s\n", CONFIG_BOARD);
    LOG_INF("Hello please");
    return 0;
}
```

Output ->

```
Booting from ROM..
*** Booting Zephyr OS build zephyr-v3.3.0-3615-gb21496f9d0fa ***
Hello World! qemu_x86
[00:00:00.010,000] <inf> hello_world: Hello please
```

Timestamp <level> <module name> <Message>



Logger API's

X = ERR, WRN, DBG, INFO

- **LOG_X** : Standard printf-like messages.
Eg: LOG_ERR
- **LOG_INST_X** : Associated with the particular instance. Eg: LOG_INST_INFO
- **LOG_HEXDUMP_X** : Dump hex data.
- **LOG_INST_HEXDUMP_X** : Dump hex data associated with a particular instance.
- **LOG_PRINTK / LOG_RAW** :
Unconditionally print raw log message.
- **LOG_MODULE_REGISTER** : Creates module-specific state and register the module with Logger.
- **LOG_MODULE_DECLARE** : Declares a log module (not register it).

Architecture

▪ Frontend

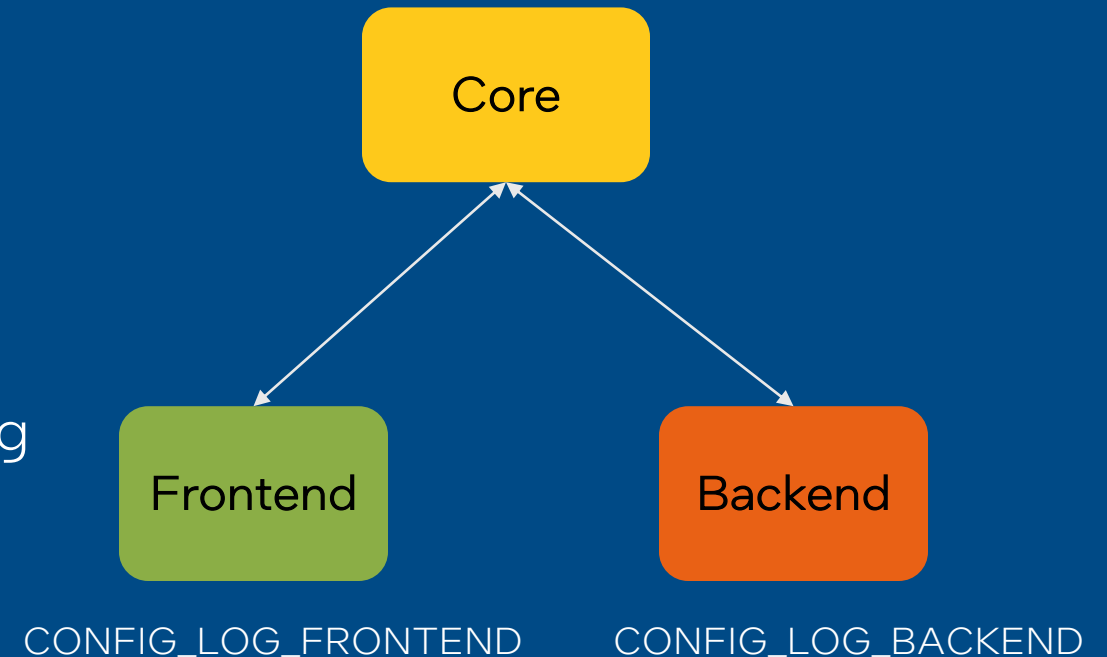
- Engaged by default when LOG API is called.
- Responsible for log filtering (Compile & Run-Time), buffer allocation, creating & committing the message.
- Optimized to log messages faster – interrupt context.

▪ Core

- Macros and functions needed for creating logging messages
- Enables/disables logging formats.

▪ Backends

- Outputs the log messages
- Cannot be selected by user.
- Messages are passed through a frontend and are then processed by active backends.



Frontend: Logging Process

- Filtering a message
 - Compile Time Filtering
 - Run Time Filtering
- Allocating buffer for the message
- Message Creation
- Commit the Message

Log Message

- Self-contained, continuous block of memory
- Log message contents:
 - Message Descriptor (Source, Domain & Level)
 - Timestamp
 - Formatted String details
 - Optional data

Log Message Format

- Message Header
 - MPSC packet buffer header: 2 bits
 - Trace/Log message flag : 1 bit
 - Domain ID: 3 bits
 - Level: 3 bits
 - Cbprintf Package Length: 10 bits
 - Data length: 12 bits
 - Reserved: 1 bit
 - Pointer: Pointer to the source descriptor
 - Timestamp: 32 or 64 bits
 - Optional padding
- Cbprintf package (optional) – Header, Arguments & Appended strings
- Hexdump data (optional)
- Alignment padding(optional)

Log Message Memory

- Messages stored in Multi Producer Single Consumer Packet Buffer (MPSC_PBUF) – circular buffer as continuous block of memory
- Messages must be sequentially freed – suited for copying the messages for offline processing.
- Backend processing is synchronous. Backend can make a copy for deferred processing.

Multi Producer Single Consumer Packet Buffer

- Circular Buffer - First-in First Out Order, Continuous block of memory.
- Allocation Policy, when requested space cannot be allocated:
 - Overwrite - [CONFIG_LOG_MODE_OVERFLOW](#) (Degrades Performance)
 - No-overwrite
- Producing the packet – 2 steps
 - Requested amount of data allocated, producer fills the data
 - Commits it
- Consuming the packet - 2 steps
 - Consumer claims the packet (gets pointer to it and length)
 - Packet is freed (Reduces memory copying)

Multi Producer Single Consumer Packet Buffer(Contd..)

- Each packet contains MPSC_PBUF specific header.
- Header – 2 bits
 - Valid
 - Busy (Packet being consumed)

Valid	Busy	Description
0	0	Space is Free
1	0	Valid Packet
1	1	Claimed valid Packet
0	1	Internal Skip Packet

Log Filtering

- Need?
 - To reduce Image size, to not overload the system
- Compile Time Filtering
 - Module
 - Severity Levels
- RunTime Filtering
 - Source : Module or specific instance of module
 - Independent for each backend

Run-Time Filtering

- Filter structure (Ten 3bit slots in RAM) for each source of logging:

Slot 0	INF	→ Aggregate maximal filter for given source
Slot 1	ERR	
Slot 2	INF	
Slot 3	OFF	
Slot 4	...	
Slot 5	...	
Slot 6	...	
Slot 7	...	
Slot 8	...	
Slot 9	OFF	

Each slot stores current filter for one backend in the system

Modes of Logging

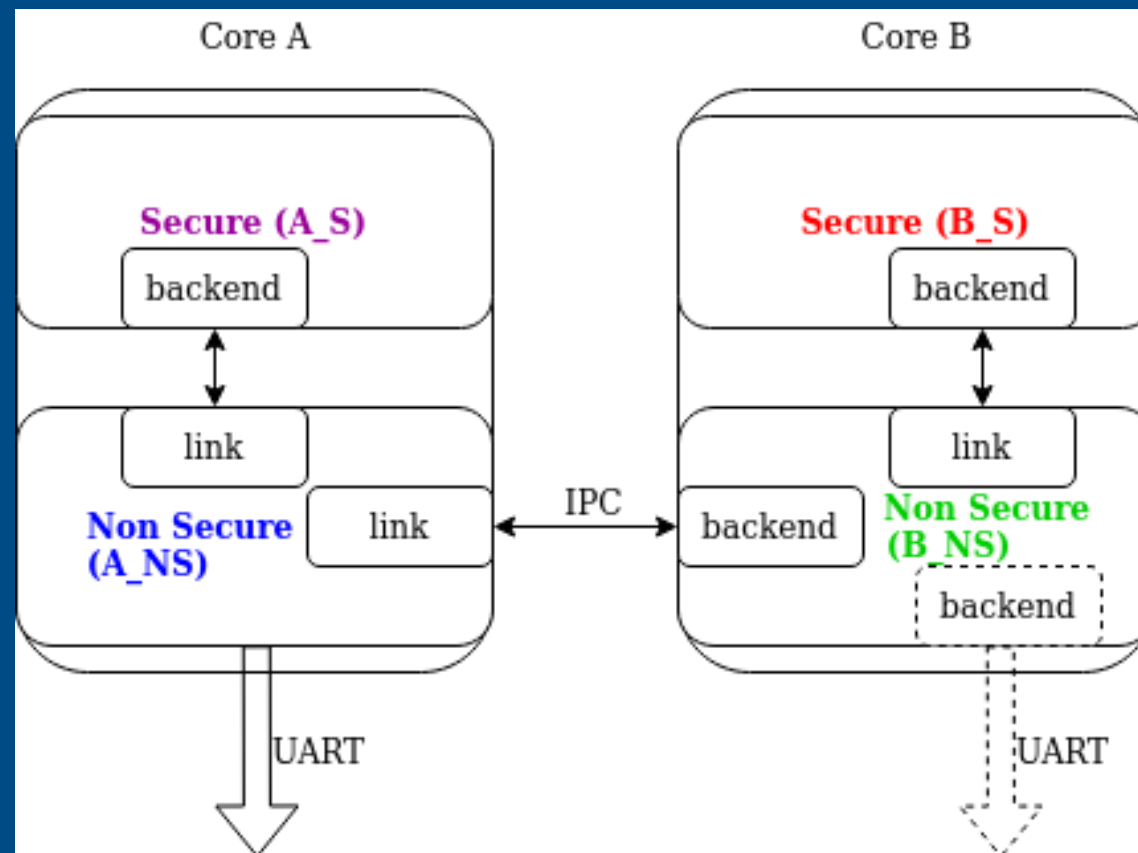
Deferred	Immediate	Minimal
<ul style="list-style-type: none">• Log messages are buffered and processed later.• Time consuming processing is deferred to the known context. Least impact on the application.• Build-time, run-time filtering• Rich formatting, timestamping.• Close to non-intrusive• Any backend	<ul style="list-style-type: none">• Log is processed in the context of the call. Immediate output.• Impacts performance, time consuming operations performed in context of the log entry (e.g. high priority interrupt)• Filtering and formatting• Intrusive• Limited backends supported	<ul style="list-style-type: none">• Redirected to printk• No run time filtering, formatting, timestamping• Low Footprint

Multi-Domain Logging

- Domain : Independent Binary Build
 - Restricted Memory access – cross domains
- Approaches:
 1. Inside each domain independently.
 - Not scalable
 2. Use a multi-domain logging system.
 - Log messages inside each domain goes to one root domain.
 - Passed using link between domains. One backend -> Other backend.
- Log Link Interface:
 - Receives log messages from another domain.
 - Creates copy & put local log message copy into message queue.
 - Matches complementary backend implementation.

Multi-Domain Logging (Contd..)

- Types of Domains:
 - End Domain:
 - Logging core implementation & a cross-domain backend.
 - Relay Domain
 - Links to other domains but does not have backends.
 - Has a cross-domain backend either to another relay or to the root domain.
 - Root Domain
 - One or multiple links and a backend that outputs logs to the user.
- Runtime - filtering works identically in both multi-domain and single-domain scenarios.



Types of Logging Formats

- MIPI-SyS-T: **CONFIG_LOG_MIPI_SYST_ENABLE**
 - Universal data format for sharing software debug and trace information between test systems and devices such as systems-on-chip (SoC) or platforms.
 - Simplify Integration of embedded software and debug hardware
 - Developing test products
- Dictionary: **LOG_DICTIONARY_SUPPORT**
 - Outputs in binary format.
 - Encodes arguments to formatted strings in native storage formats - compact
 - **long double** - Not supported.
- Text: Default
 - Requires less processing than dict.
- Custom Format: **CONFIG_LOG_CUSTOM_FORMAT_SUPPORT**

Dynamically Switch Logging formats

- User can pass the output format needed and use API's like
 - `log_format_set(const struct log_backend *backend, uint32_t log_type)`
 - `log_format_set_all_active_backends(size_t log_type)`
- Limitations – Deferred mode.
 - Results in dropped messages – large number of messages to be printed.
 - Won't be able to see complete results.
- Recommendations:
 - Use Immediate Mode

```
RAW DATA: 020A000B4200020000000000000000A2A202426F6F74696E67205A6570687972204F53206275696C64207A65706879722D76332E3320D333631352D7623231343936663964306661202A2A0A00
SYS-T RAW DATA: 220A000B17000200000000000000004572726F72206D657373616765206578616D706C652E00
SYS-T RAW DATA: 320A000B19000200000000000000005761726E696E67206D657373616765206578616D706C652E00
SYS-T RAW DATA: 420A000B1600020000000000000000496E666F206D657373616765206578616D706C652E00
SYS-T RAW DATA: 720A000B17000200000000000000004465627567206D657373616765206578616D706C652E00
SYS-T RAW DATA: 720A000B1E000200000000000000004465627567206D657373616765206578616D706C652C2025640001000000
SYS-T RAW DATA: 720A000B26000200000000000000004465627567206D657373616765206578616D706C652C2025642C202564000100000002000000
SYS-T RAW DATA: 720A000B2E000200000000000000004465627567206D657373616765206578616D706C652C2025642C2025642C20256400010000000200000003000000
SYS-T RAW DATA: 720A000B38000200000000000000004465627567206D657373616765206578616D706C652C2025642C2025642C2025642C20307825780001000000020000000300000004000000
SYS-T RAW DATA: 720A000B0C00020000000000000000636861722025630021000000
SYS-T RAW DATA: 720A000B1D000200000000000000007320737472202573202573007374617469632073747200632073747200
SYS-T RAW DATA: 720A000B150003000000000000000064207374722025730064796E616D69632073747200
SYS-T RAW DATA: 720A000B6B000300000000000000006D69786564207374722025732025732025732025732025732025730064796E616D696320737472002D2D2D0064796E616D696320737472002D2D2D00616E6F746865722064796E616D6963
20737472002D2D2D00616E6F746865722064796E616D69632073747200
SYS-T RAW DATA: 720A000B43000300000000000000006D6978656420632F732025632025732025732025732025630021000000737461746963207374720064796E616D69632073747200737461746963207374720021000000
SYS-T RAW DATA: 720A000B22000300000000000000004465627567206D657373616765206578616D706C652C20256600EA2E4454FB210940
SYS-T RAW DATA: 220A000B09000300000000000000002573006672616D6500
SYS-T RAW DATA: 220A000B47000300000000000000002573003033206432203034203030203030203038203031203032202030332030342030352030362030372030382020202020207C2E2E2E2E2E2E2E2E2E2E2E2E2E2E202000
SYS-T RAW DATA: 320A000B09000300000000000000002573006672616D6500
SYS-T RAW DATA: 320A000B4700030000000000000000257300303320643220303420303020303020303820303120303220203033203034203035203036203037203038202020202020207C2E2E2E2E2E2E2E2E2E2E2E2E2E2E202000
SYS-T RAW DATA: 420A000B09000300000000000000002573006672616D6500
SYS-T RAW DATA: 420A000B4700030000000000000000257300303320643220303420303020303020303820303120303220203033203034203035203036203037203038202020202020207C2E2E2E2E2E2E2E2E2E2E2E2E2E2E202000
SYS-T RAW DATA: 720A000B09000300000000000000002573006672616D6500
SYS-T RAW DATA: 720A000B4700030000000000000000257300303320643220303420303020303020303820303120303220203033203034203035203036203037203038202020202020207C2E2E2E2E2E2E2E2E2E2E2E2E2E2E202000
SYS-T RAW DATA: 020A000B220004000000000000000068656C6C6F207379732D74206F6E20626F6172642025730A0071656D755F78383600
[00:00:00.040,000] <err> syst: Error message example.
[00:00:00.040,000] <wrn> syst: Warning message example.
[00:00:00.040,000] <inf> syst: Info message example.
[00:00:00.040,000] <dbg> syst: Debug message example.
[00:00:00.040,000] <dbg> syst: Debug message example, 1
[00:00:00.040,000] <dbg> syst: Debug message example, 1, 2
[00:00:00.040,000] <dbg> syst: Debug message example, 1, 2, 3
[00:00:00.040,000] <dbg> syst: Debug message example, 1, 2, 3, 0x4
[00:00:00.040,000] <dbg> syst: char !
[00:00:00.040,000] <dbg> syst: s str static str c str
[00:00:00.040,000] <dbg> syst: d str dynamic str
[00:00:00.040,000] <dbg> syst: mixed str dynamic str --- dynamic str --- another dynamic str --- another dynamic str
[00:00:00.040,000] <dbg> syst: mixed c/s ! static str dynamic str static str !
[00:00:00.040,000] <dbg> syst: Debug message example, %f
[00:00:00.040,000] <err> syst: frame
03 d2 04 00 00 08 01 02 03 04 05 06 07 08 |.....
[00:00:00.040,000] <wrn> syst: frame
03 d2 04 00 00 08 01 02 03 04 05 06 07 08 |.....
[00:00:00.050,000] <inf> syst: frame
03 d2 04 00 00 08 01 02 03 04 05 06 07 08 |.....
[00:00:00.050,000] <dbg> syst: frame
03 d2 04 00 00 08 01 02 03 04 05 06 07 08 |.....
hello sys-t on board qemu_x86
SYS-T RAW DATA: 220A000B17000500000000000000004572726F72206D657373616765206578616D706C652E00
SYS-T RAW DATA: 320A000B19000500000000000000005761726E696E67206D657373616765206578616D706C652E00
SYS-T RAW DATA: 420A000B1600050000000000000000496E666F206D657373616765206578616D706C652E00
SYS-T RAW DATA: 720A000B17000500000000000000004465627567206D657373616765206578616D706C652E00
SYS-T RAW DATA: 720A000B1E000500000000000000004465627567206D657373616765206578616D706C652C2025640001000000
SYS-T RAW DATA: 720A000B26000500000000000000004465627567206D657373616765206578616D706C652C2025642C202564000100000002000000
SYS-T RAW DATA: 720A000B2E000500000000000000004465627567206D657373616765206578616D706C652C2025642C2025642C20256400010000000200000003000000
SYS-T RAW DATA: 720A000B38000500000000000000004465627567206D657373616765206578616D706C652C2025642C2025642C2025642C20307825780001000000020000000300000004000000
SYS-T RAW DATA: 720A000B0C00050000000000000000636861722025630021000000
SYS-T RAW DATA: 720A000B1D000500000000000000007320737472202573202573202573007374617469632073747200632073747200
SYS-T RAW DATA: 720A000B150005000000000000000064207374722025730064796E616D69632073747200
SYS-T RAW DATA: 720A000B6B000500000000000000006D69786564207374722025732025732025732025732025732025730064796E616D696320737472002D2D2D0064796E616D696320737472002D2D2D00616E6F746865722064796E616D6963
20737472002D2D2D00616E6F746865722064796E616D69632073747200
SYS-T RAW DATA: 720A000B43000600000000000000006D6978656420632F732025632025732025732025732025630021000000737461746963207374720064796E616D69632073747200737461746963207374720021000000
SYS-T RAW DATA: 720A000B22000600000000000000004465627567206D657373616765206578616D706C652C20256600EA2E4454FB210940
```

MIPI SyS-T

Text

MIPI SyS-T

References

- [Logging Documentation](#)
- Mailing Lists:
 - User List: users@lists.zephyrproject.org
 - Developer List: devel@lists.zephyrproject.org
- Discord Channel:
 - #logging
 - #tracing
- Contributing to Zephyr: See the [Contribution Guide](#)

Questions? Comments?

Thank You
- Aastha Grover

The Intel logo is centered on a solid blue background. It features the word "intel" in a white, lowercase, sans-serif font. A small, light blue square is positioned above the first vertical stroke of the letter "i". To the right of the word "intel" is a small white registered trademark symbol (®).

intel®