



Zephyr[®] Project

Developer Summit


Flexible system design via RPC: Embracing distributed computing in Zephyr

Yuval Peress, *Google*

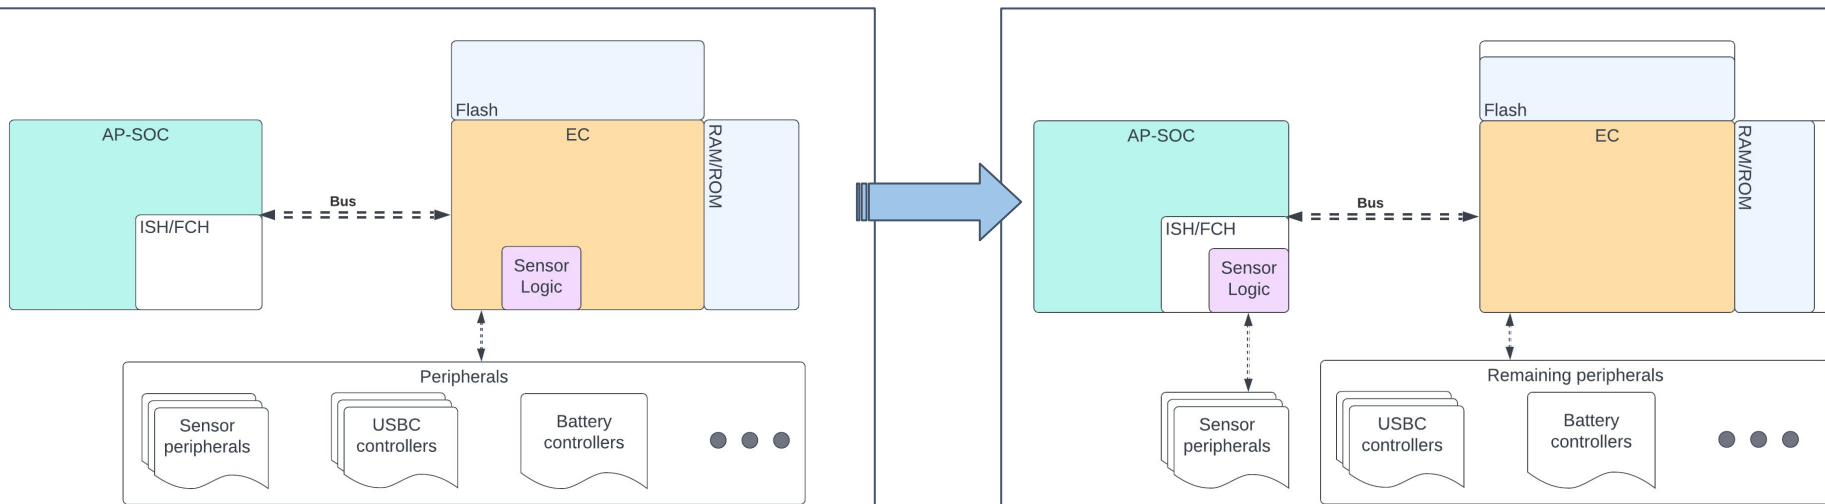
Email: peress@google.com

Discord: [yuval#5515](#)

GitHub: [yperess](#)



- Chromebooks ship with an EC and a separate Application Processor (AP)
- Many APs (Intel and AMD) come with a dedicated sensor core
- We wanted to move the sensor logic to reduce the cost of the EC
- But how? Dependencies, tests, and prior designs were broken



Agenda

- Portable design
- Pigweed RPCs and protobufs
- Transitioning from headers to services
- An example

Portable design

hint: they're microservices

Chromium's EC has many tasks

SYSWORKQ

SHELL

TASK_TOUCHPAD

TASK_CHG_RAMP

TASK_USB_CHG

TASK_DPS

TASK_CHARGER

TASK_CHIPSET

TASK_MOTIONSENSE

TASK_USB_MUX

TASK_HOSTCMD

TASK_KEYPROTO

TASK_POWERBTN

TASK_KEYSCAN


TASK_PD_C<port_num>

TASK_PD_INT_C<portn_num>

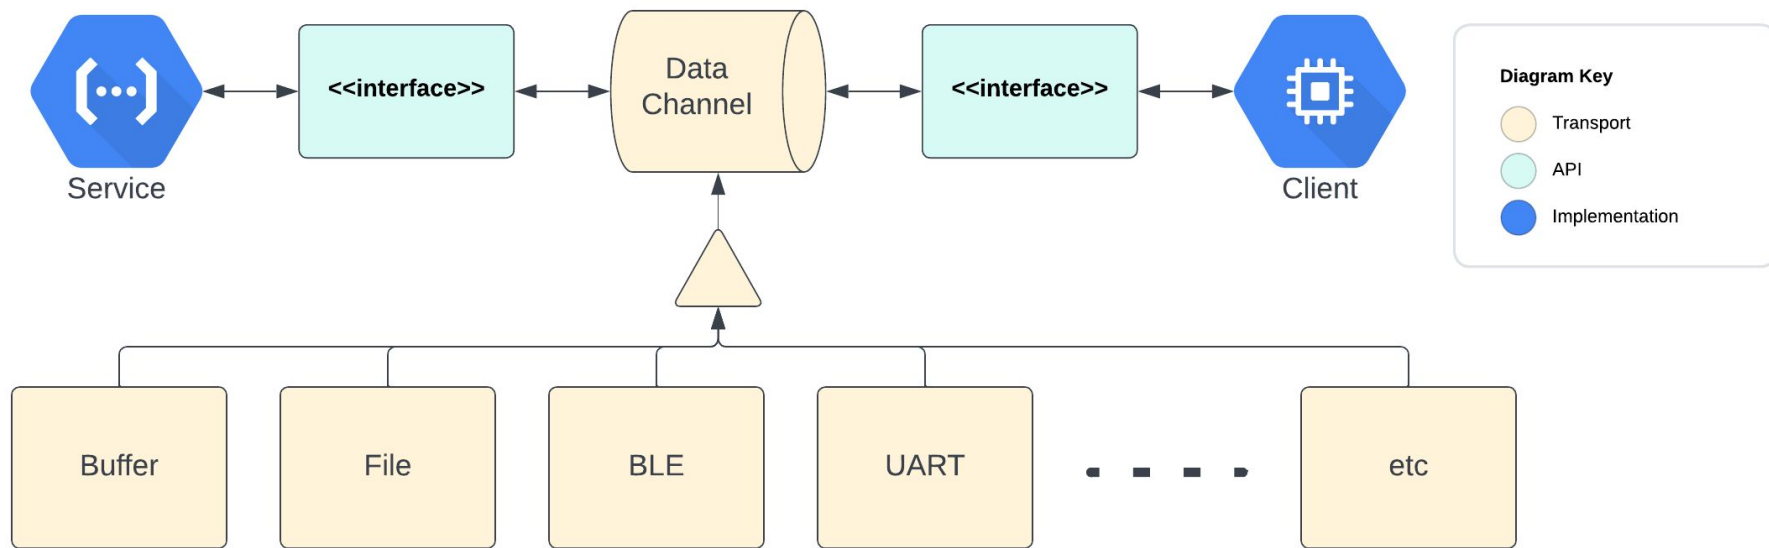
- Sensors
- Power Delivery

Does it matter where the task lives?

- Sensor logic can be on a dedicate core such as Intel's ISH or AMD's SFH.
- Power delivery (PD) logic can be on PD chip



With the right design,
the service and clients never need to be rewritten





Zephyr® Project
Developer Summit

Pigweed RPCs and protobufs

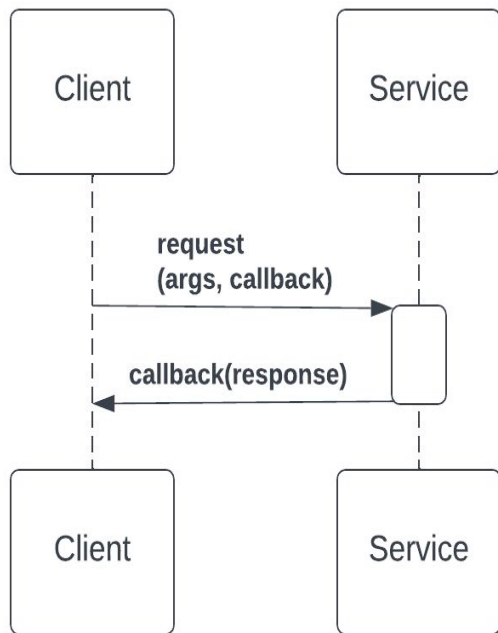
#EMBEDDEDOSSUMMIT

What is Pigweed?

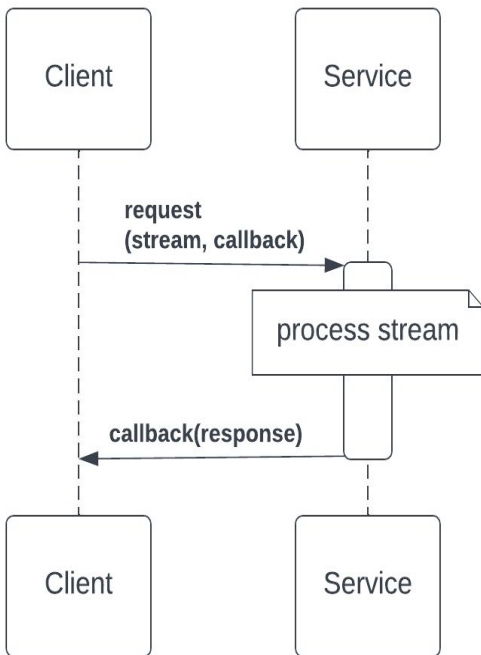
- Pigweed is a collection of tools/modules
- Highly tuned for embedded applications
- Modules in this talk:
 - pw_rpc
 - pw_hdlc

RPC concepts

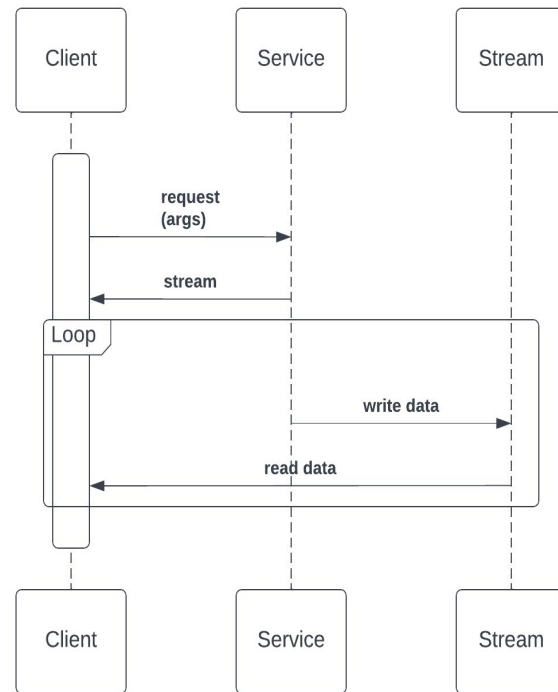
Unary RPC



Client Streaming




Server Streaming



A simple service

- Asynchronous
- `set_val`
 - passes an int and saves it on the service
 - returns a status when finished
- `get_val`
 - passes no args
 - returns a status and the current value



When to use streams?

- Pigweed uses server streams for logging (`pw_log_rpc`)
 - Client makes a request to the service (EC) and gets back a stream
 - Each message on the stream is a log message
- When data is generated with some latency

How to set it up (west.yml)?

Add the remote

remotes:

- name: pigweed
url-base: <https://pigweed.googlesource.com/pigweed>

Add pigweed to the project

projects:

- name: pigweed
remote: pigweed
revision: main

How to set it up (kConfig)?

```
(Top) → Zephyr → C++ Language Support
[*] C++ support for the application
    C++ Standard (C++ 17) --->

(Top) → Zephyr → Modules
*** Available modules. ***
pigweed (/home/peress/workspace/zds2023/pigweed) --->

(Top) → Zephyr → Modules → pigweed (/home/peress/workspace/zds2023/pigweed)
pw_assert --->
pw_base64 --->
pw_bytes --->
pw_checksum --->
pw_chrono --->
pw_containers --->
```

Enable C++

How to set it up (kConfig)?

```
(Top) → Zephyr → C++ Language Support
[+] C++ support for the application
    C++ Standard (C++ 17) --->

(Top) → Zephyr → Modules
*** Available modules. ***
pigweed (/home/peress/workspace/zds2023/pigweed) --->

(Top) → Zephyr → Modules → pigweed (/home/peress/workspace/zds2023/pigweed)
pw_assert --->
pw_base64 --->
pw_bytes --->
pw_checksum --->
pw_chrono --->
pw_containers --->
```

Find the Pigweed
module

How to set it up (kConfig)?

```
(Top) → Zephyr → C++ Language Support
[*] C++ support for the application
    C++ Standard (C++ 17) --->

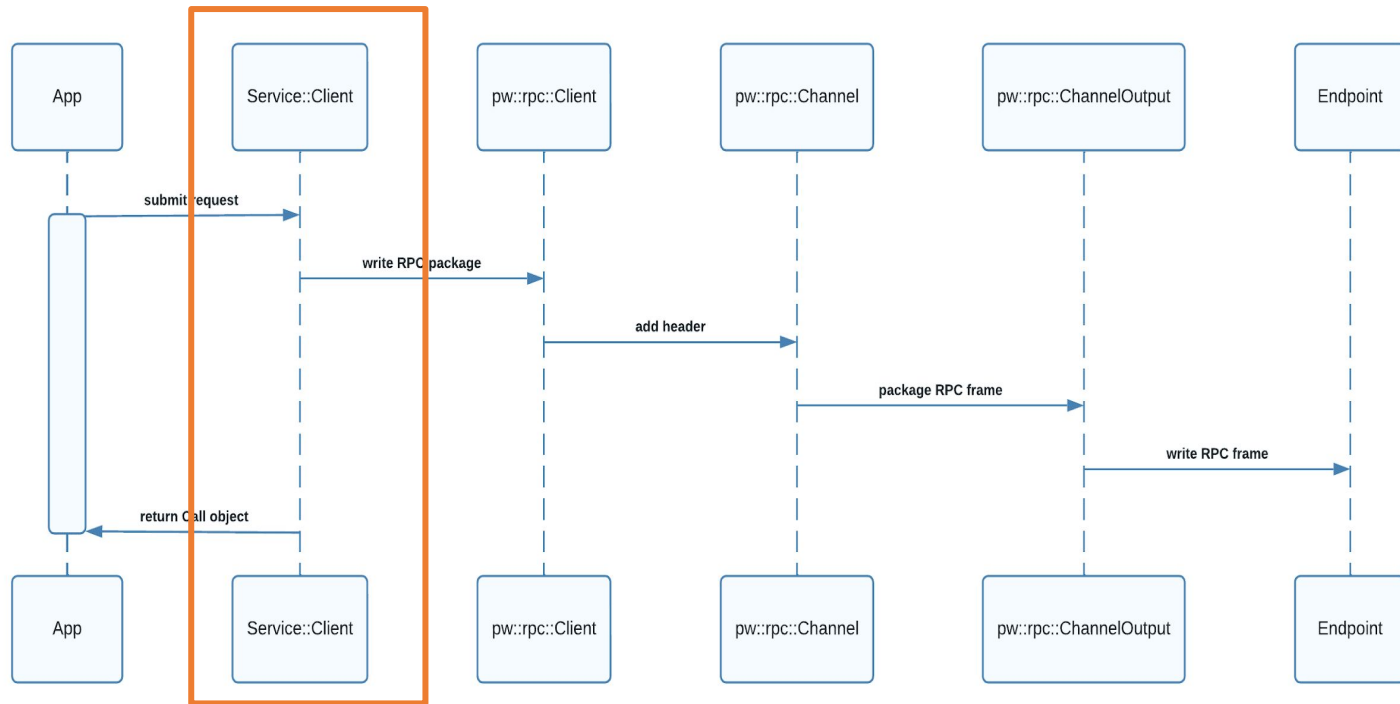
(Top) → Zephyr → Modules
*** Available modules. ***
pigweed (/home/peress/workspace/zds2023/pigweed) --->

(Top) → Zephyr → Modules → pigweed (/home/peress/workspace/zds2023/pigweed)
pw_assert --->
pw_base64 --->
pw_bytes --->
pw_checksum --->
pw_chrono --->
pw_containers --->
```

Enable the pw_*
libraries you need

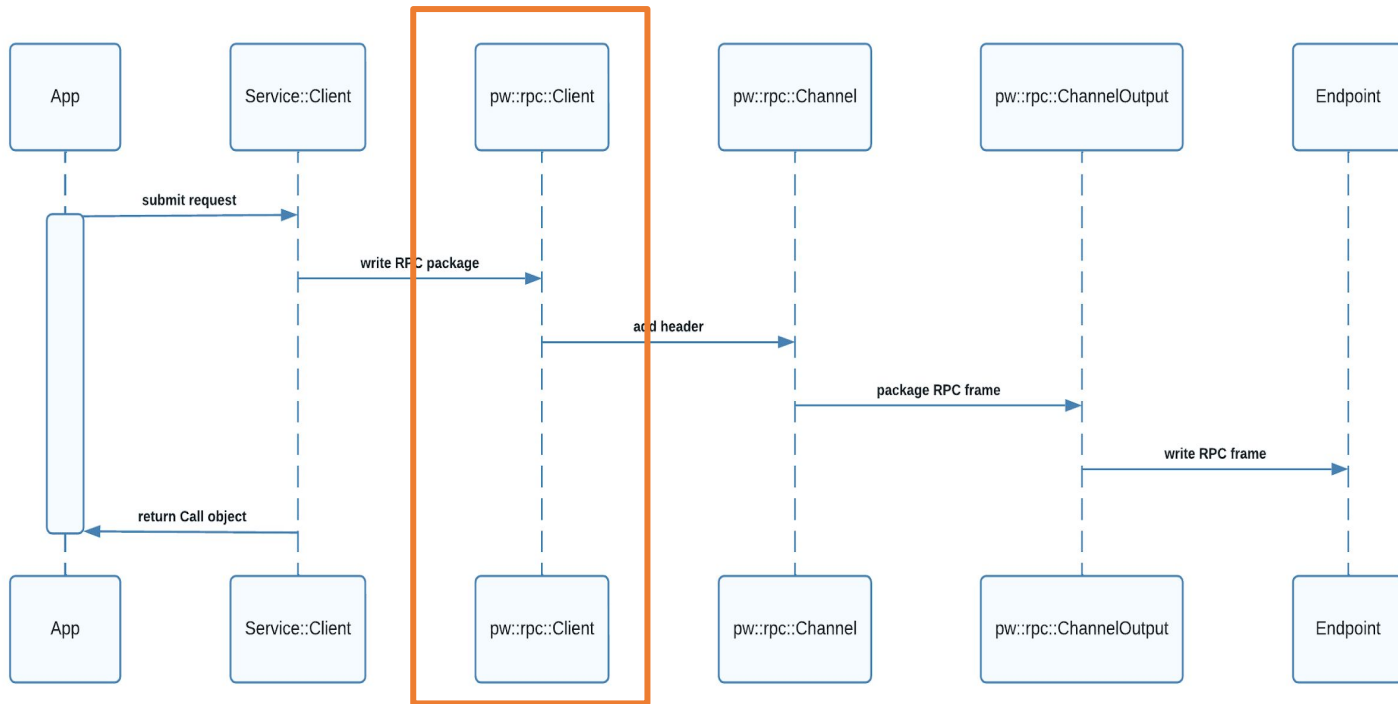
Data visualization of pw_rpc

Custom client that injects the service specific information automatically, like service and method IDs.



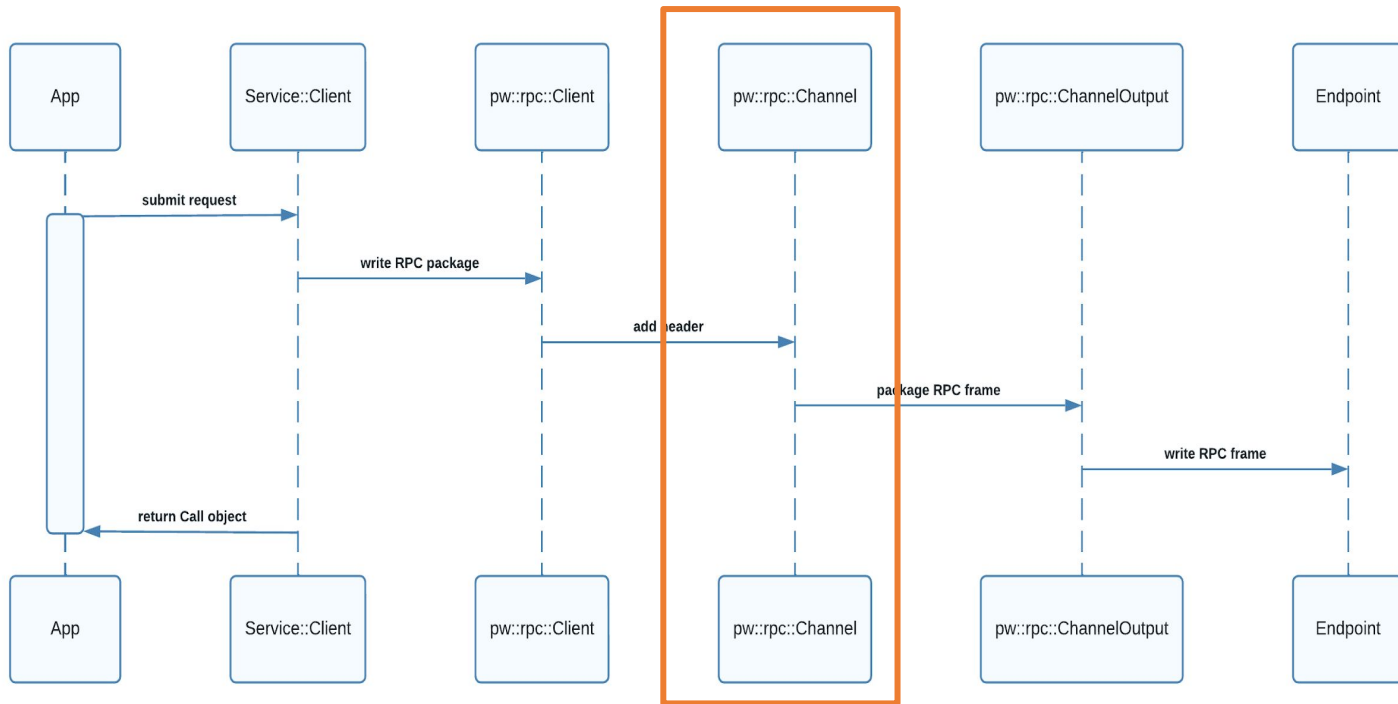
Data visualization of pw_rpc

General purpose client which routes the package to the correct channel



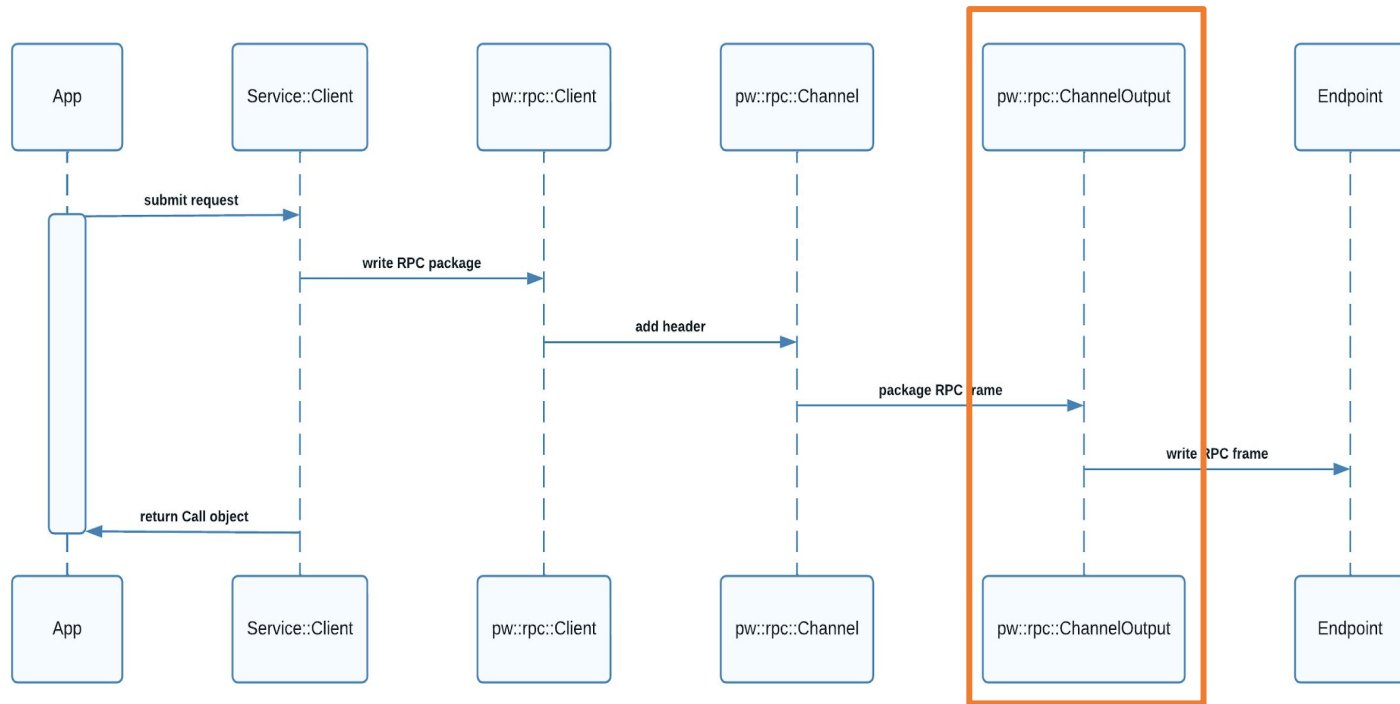
Data visualization of pw_rpc

Custom channel
implementation
controlling the
RPC protobuf
wire format



Data visualization of pw_rpc

Writes the RPC frame to the destination and potentially add another level of encoding





Zephyr® Project
Developer Summit

Transitioning from headers to services

#EMBEDDEDOSSUMMIT

Can't I just use a .h file? Yes, but...

- Protobuffers are more flexible (client and service can be in different languages)
- Protobuffers make it easy to test by creating a mock client/service
- Protobuffers extend easier than structs (support multiple versions)

Proto files force us to think about API boundaries

- APIs + documentation provide a contract
- Keep our code interaction confined so refactors are easier to manage
- Boundaries make writing tests easier and faster
- Protos are designed to be extensible

What would the header look like?

```
/* The data structures used for the API */
```

```
struct SetValueRequest {  
    int32_t value;  
};  
struct SetValueResponse {};  
struct GetValueRequest {};  
struct GetValueResponse {  
    int32_t value;  
};
```

```
/* APIs for setting and getting the value */
```

```
int client_set_value(  
    const struct SetValueRequest *request,  
    void(*callback)(int status, const struct SetValueResult *result)  
);  
int client_get_value(  
    const struct GetValueRequest *request,  
    void(*callback)(int status, const struct GetValueResult *result)  
);
```

So what's the problem?

1. Data structs are hard to maintain as new arguments are added
2. The API is lacking a lot of features still
 - a. How is the data is sent?
 - b. What is the wire format?
 - c. What thread is the callback called on?
 - d. Can we cancel a request?
3. How do we test this service?

`pw_rpc` solves these issues

Example

Let's build a simple service

- SetValue needs:
 - Passing the value as an argument

```
message SetValueRequest {  
    int32 value = 1;  
}
```

Let's build a simple service

- SetValue needs:
 - Passing the value as an argument
 - No return value

```
message SetValueRequest {  
    int32 value = 1;  
}  
  
message SetValueResponse {}
```

Let's build a simple service

- SetValue need:
 - Passing the value as an argument
 - No return value
- GetValue needs:
 - Passing nothing

```
message SetValueRequest {  
    int32 value = 1;  
}  
  
message SetValueResponse {}  
  
message GetValueRequest {}
```

Let's build a simple service

- SetValue need:
 - Passing the value as an argument
 - No return value
- GetValue needs:
 - Passing nothing
 - Return the value

```
message SetValueRequest {  
    int32 value = 1;  
}  
  
message SetValueResponse {}  
  
message GetValueRequest {}  
  
message GetValueResponse {  
    int32 value = 1;  
}
```

Let's build a simple service

- SetValue need:
 - Passing the value as an argument
 - No return value
- GetValue needs:
 - Passing nothing
 - Return the value
- Add the service

```
message SetValueRequest {  
    int32 value = 1;  
}  
  
message SetValueResponse {}  
  
message GetValueRequest {}  
  
message GetValueResponse {  
    int32 value = 1;  
}
```

```
service Cache {  
    rpc SetValue(SetValueRequest) returns (SetValueResponse) {}  
    rpc GetValue(GetValueRequest) returns (GetValueResponse) {}  
}
```


Service implementation header

```
class Cache : public pw_rpc::nanopb::Cache::Service<Cache> {  
    public:  
        Cache() : value_(0) {}  
        ::pw::Status SetValue(const ::SetValueRequest& request, ::SetValueResponse& response);  
        ::pw::Status GetValue(const ::GetValueRequest& request, ::GetValueResponse& response);  
    private:  
        int32_t value_;  
};
```

Service implementation

```
::pw::Status Cache::SetValue(const ::SetValueRequest& request, ::SetValueResponse& response) {  
    value_ = request.value;  
    return ::pw::OkStatus();  
}
```


```
::pw::Status DemoService::GetValue(const ::GetValueRequest& request, ::GetValueResponse& response) {  
    response.value = value_;  
    return ::pw::OkStatus();  
}
```

Abstracting away the ChannelOutput

- The ChannelOutput controls the wire format
- Generally, uses a `pw::stream::Writer` to write the final bytes
- Can enable us to efficiently switch how the service communicates with the client
- Example ChannelOutputs:
 - `pw::hdlc::RpcChannelOutput`
 - A custom `SimpleChannelOutput` used in this talk
- Both examples will use the same stream Writer to write to a ring buffer

ChannelOutput options

- pw_hdlc provides a ChannelOutput implementation which packs the data in an HDLC frame
- For local writes (between threads) I've implemented a simple ChannelOutput which simply writes frames as:



RingBufferReaderWriter

- Transactional
- Wraps a ring buffer
- Uses a mutex and condvar to control data availability

Performance?

- Creates 2 threads (client -> service & service -> client)
- On the main thread run 1,000 iterations of:
 - Call SetData, wait for response
 - Call GetData, wait for response
- Comparison setups:
 - [control] Read/write a plain serialized struct using SimpleChannelOutput to a plain service implementation
 - [experiment] Uses pw_rpc to write simple RPC frames using SimpleChannelOutput to RPC service implementation and server
- Things to consider:
 - The control is an oversimplification (no priority control, no call cancel, doesn't account for extensibility)
 - Some code paths of pw_rpc were identified as bottlenecks are actively being optimized

Performance?

- [control] took 65,644,840 nanoseconds (~33 μ s / call)
- [experiment] took 233,103,156 nanoseconds (~116 μ s / call)

Overall slowdown 255%, but... with `pw_rpc` you get free upgrades :)

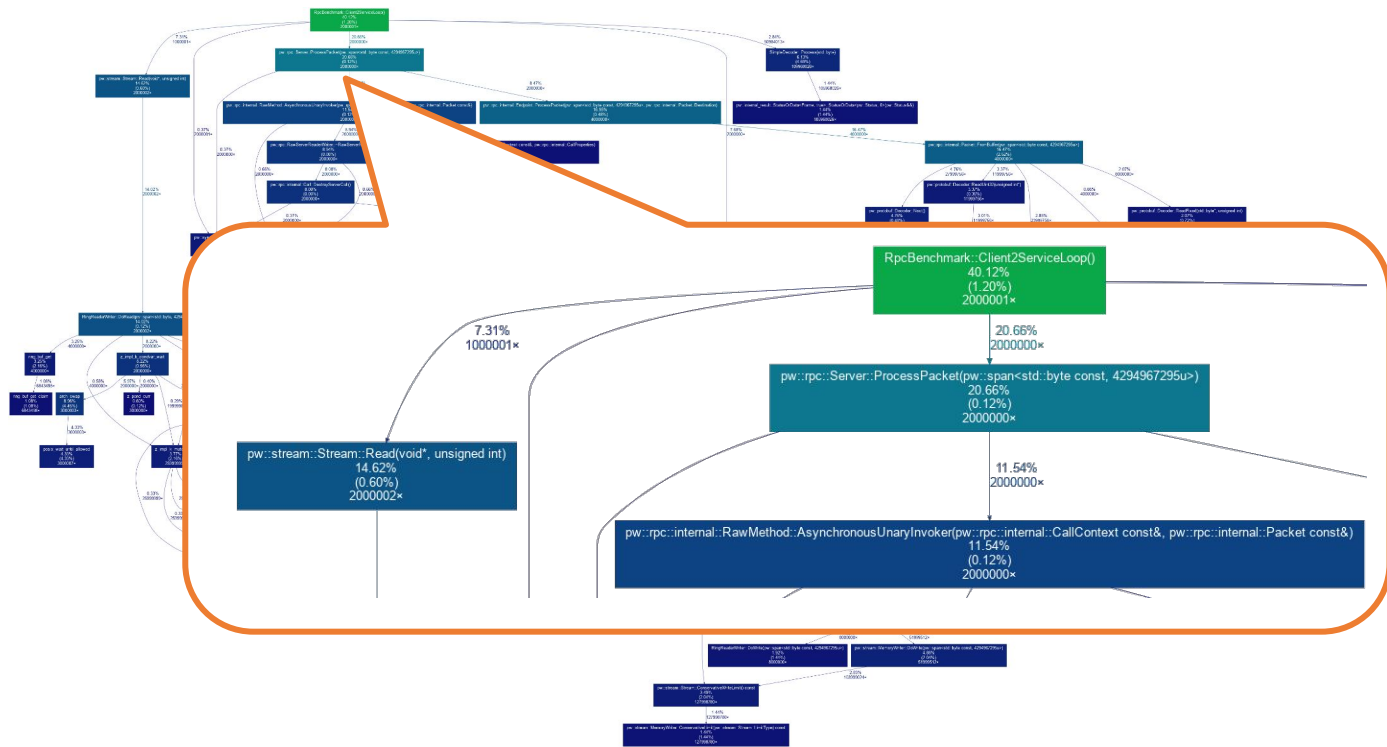
Future performance improvements

- Call graphs in the appendix show several points to improve on (when communicating between threads)
 - a. Don't use nanopb to serialize the RPC header (cost is 21%)
 - b. Remove the need for a disconnect RPC packet on the Call destructor (38%)
- These would bring overall RPC cost closer to 45% (14 μ s / call)

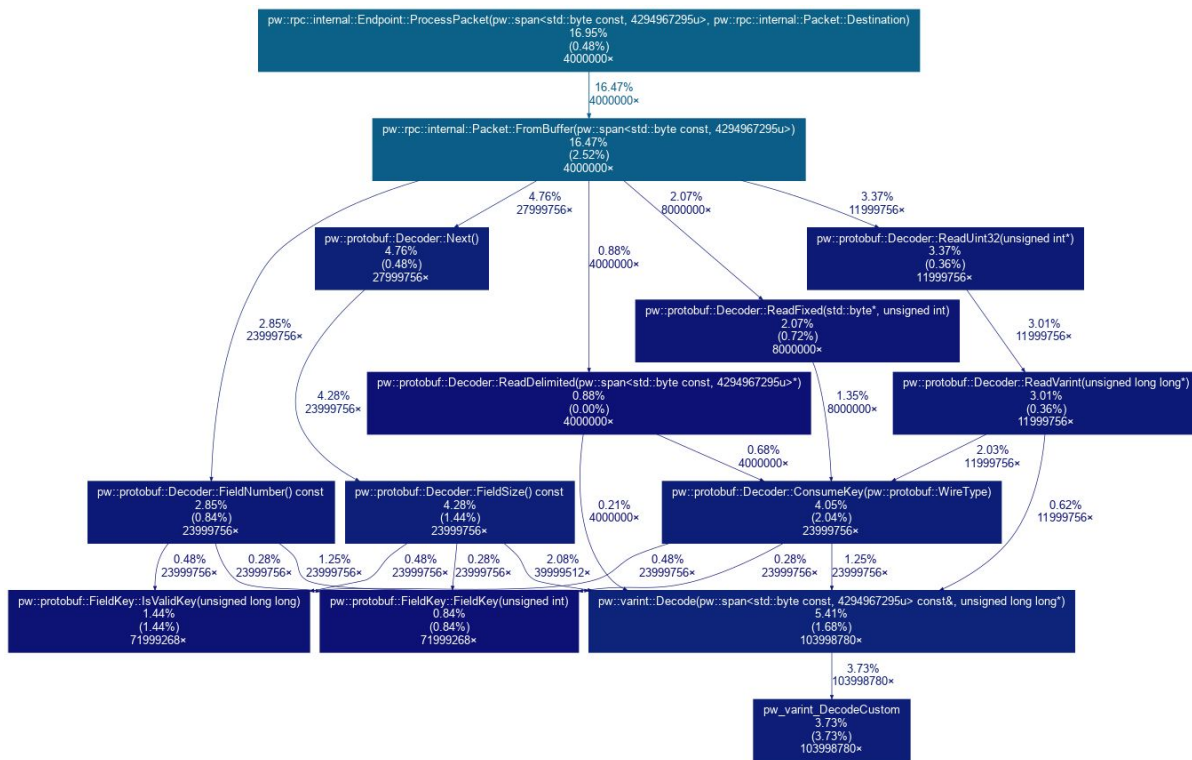
Questions?

Call graph of the client2service handler





Processing the header (21% of the cost)



Call destructor (38% of the cost)

