



Zephyr[®] Project

Developer Summit

Introduce Hardware-Level Device Isolation to Zephyr

Jaxson Han & Huifeng Zhang – Arm

jaxson.han@arm.com


GitHub: <https://github.com/povergoing>

huifeng.zhang@arm.com

GitHub: <https://github.com/SgrrZh>

Introduce Hardware-Level Device Isolation to Zephyr

Authors



Jaxson Han



Huifeng Zhang

Contents

- Background
- SMMU
- Zephyr device model
- Zephyr HW-level device isolation
- Conclusion

Background


More and more DMA device drivers on RTOS

- An observation:
 - The number of DMA devices on Low-power platforms is increasing.
 - IoT industry
 - ...
 - More RTOS on high-performance platforms with variety of DMA devices
 - Automotive Industry (high-performance & safety)

- New challenges for Zephyr:

DMA device bypass the system access control?

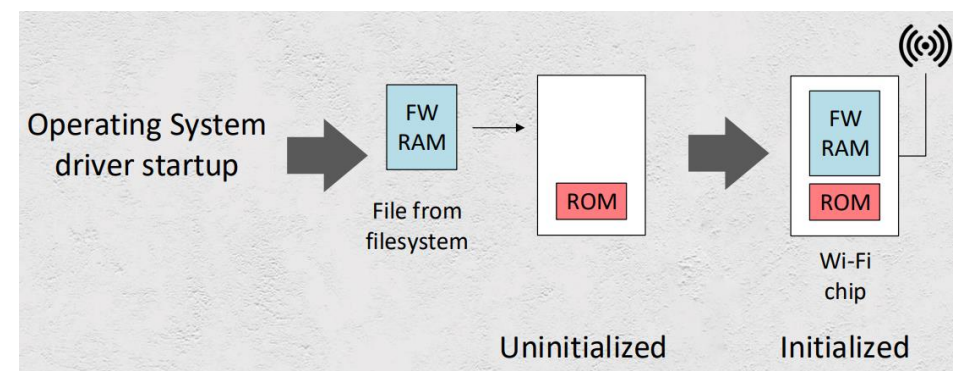
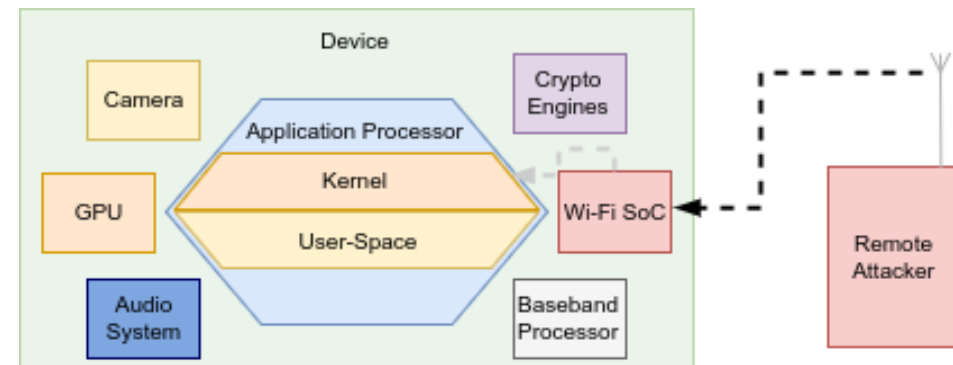
How to restrict DMA devices?



Background

DMA devices might be buggy or even malicious

- DMA devices can break the system
 - WiFi chip bug[1], [2], [3]
 - permission leaks
 - remote control
 - DMA attack[4],[5],[6]
 - steal data or cryptographic keys
 - install or run spyware and other exploits
 - modify the system to allow backdoors or other malware
- More DMA drivers added into Zephyr



How to restrict DMA devices on Zephyr?

[1] https://googleprojectzero.blogspot.com/2017/04/over-air-exploiting-broadcoms-wi-fi_4.html

[2] https://googleprojectzero.blogspot.com/2017/04/over-air-exploiting-broadcoms-wi-fi_11.html

[3] <https://www.bleepingcomputer.com/news/security/vulnerabilities-found-in-highly-popular-firmware-for-wifi-chips/>

[4] <https://web.archive.org/web/20160304055745/http://www.hermann-uwe.de/blog/physical-memory-attacks-via-firewire-dma-part-1-overview-and-mitigation>

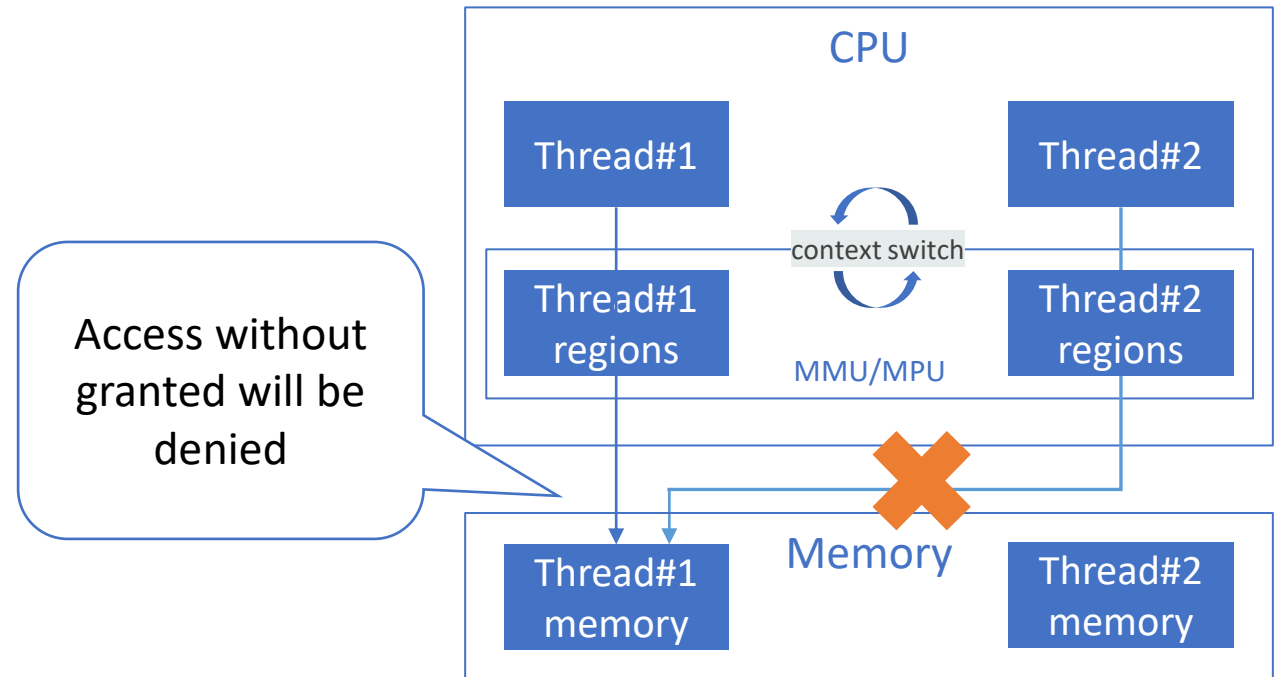

[5] <https://www.manageengine.com/device-control/prevent-dma-attacks.html>

[6] https://en.wikipedia.org/wiki/DMA_attack

Background

Why HW-level device isolation is needed

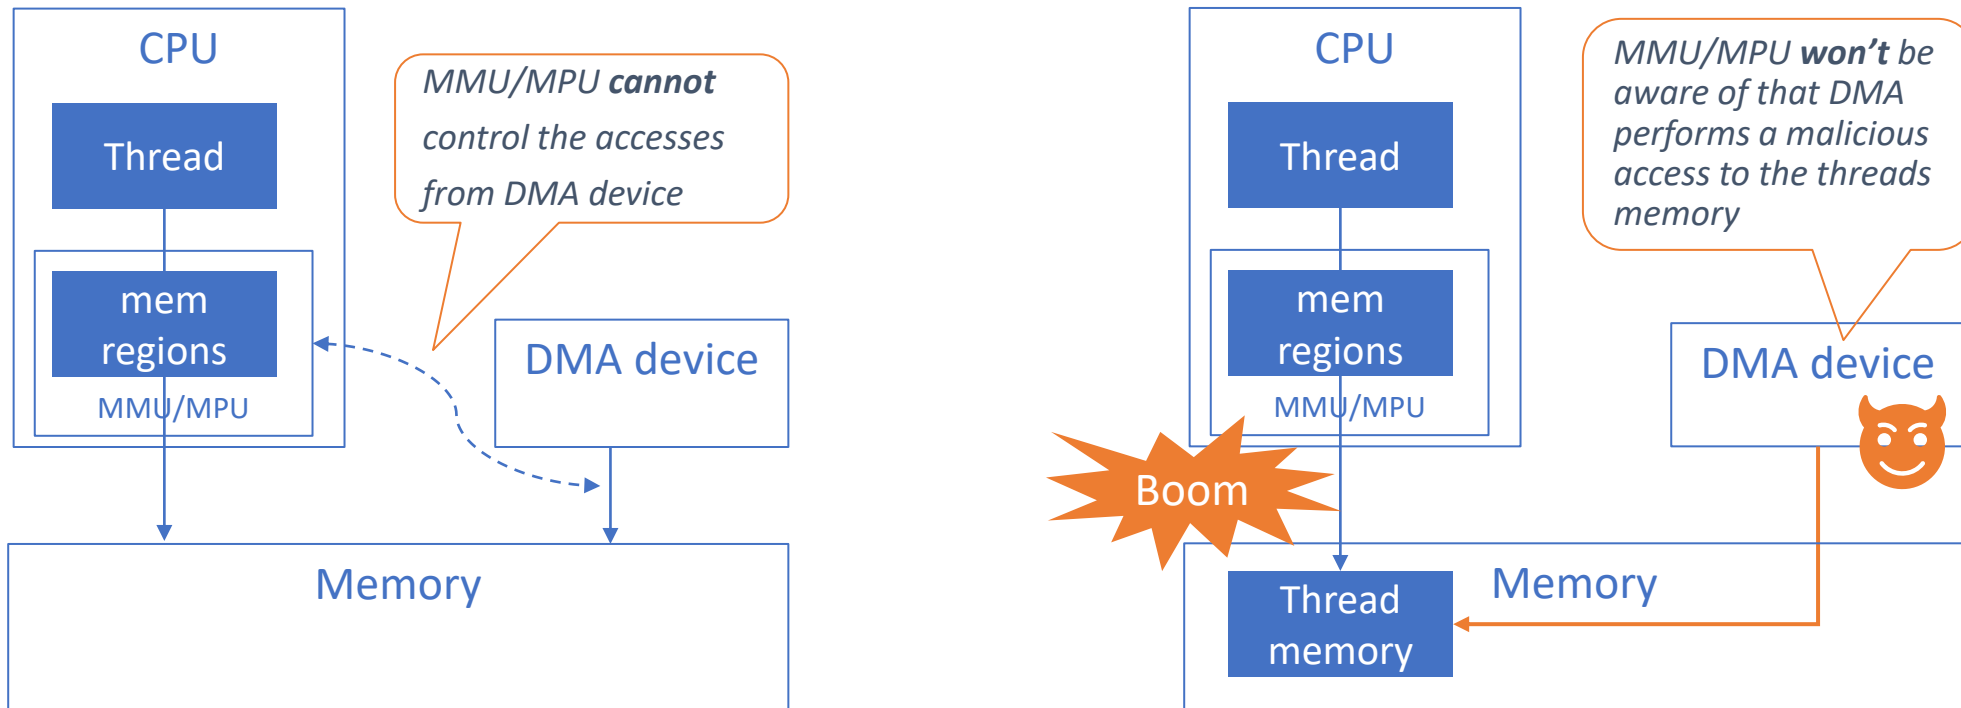
- Zephyr uses MMU/MPU to isolate the thread memory regions to protect the system.



Background

Why HW-level device isolation is needed

- However, MMU/MPU can only restrict memory accesses from CPUs.
- Memory accesses from DMA are **NOT** protected by MMU/MPU
- May cause system crash or security issues



Background

How to mitigate the issue

- Hardware-level isolation can mitigate the issue
 - SMMU (System MMU, Arm)
 - IOMMU (IO MMU, Intel)
 - xPPU (Xilinx peripheral protection unit)
 - Other vendor FireWall
 - ..
- SMMU support allows systems to share A-profile page tables with peripherals
 - widely used in Rich OSes (e.g. Linux) and hypervisor
 - isolate the DMA access
 - eliminate the requirement for physically contiguous pages for DMA buffers
 - extend old DMA device access
 - accelerate virtualization
 - be flexible to switch passthrough or virtualization

Background

How to mitigate the issue

- But too powerful
 - Fully supporting SMMU/IOMMU increases overhead for Zephyr
 - **Inappropriate** to add SMMU/IOMMU for low-power platforms.
 - increased cost
 - power consuming
- What we do
 - Partially enable SMMUv3 (as an impl example) for DMA device isolation
 - lower the overhead by using linear mapping (lower TLB miss)
 - isolate every DMA devices to improve the security
 - Add a Subsystem interface to manage DMA device isolation
 - easier to use for driver, components and applications
 - to support more hardware-level isolation technologies

Background

What's more apart from the isolation


- Zephyr as a Virtual Machine Manager (VMM) or a hypervisor?
 - needs SMMU/IOMMU driver to support the virtualization
 - passthrough, VirtIO or ...
 - accelerate the DMA access
- Zephyr to support more platforms?
 - devices require SMMU/IOMMU in some high-performance platform
 - To support some 32-bit DMA devices on 64-bit platforms

Contents

- Background
- **SMMU**
- Zephyr device model
- Zephyr HW-level device isolation
- Conclusion

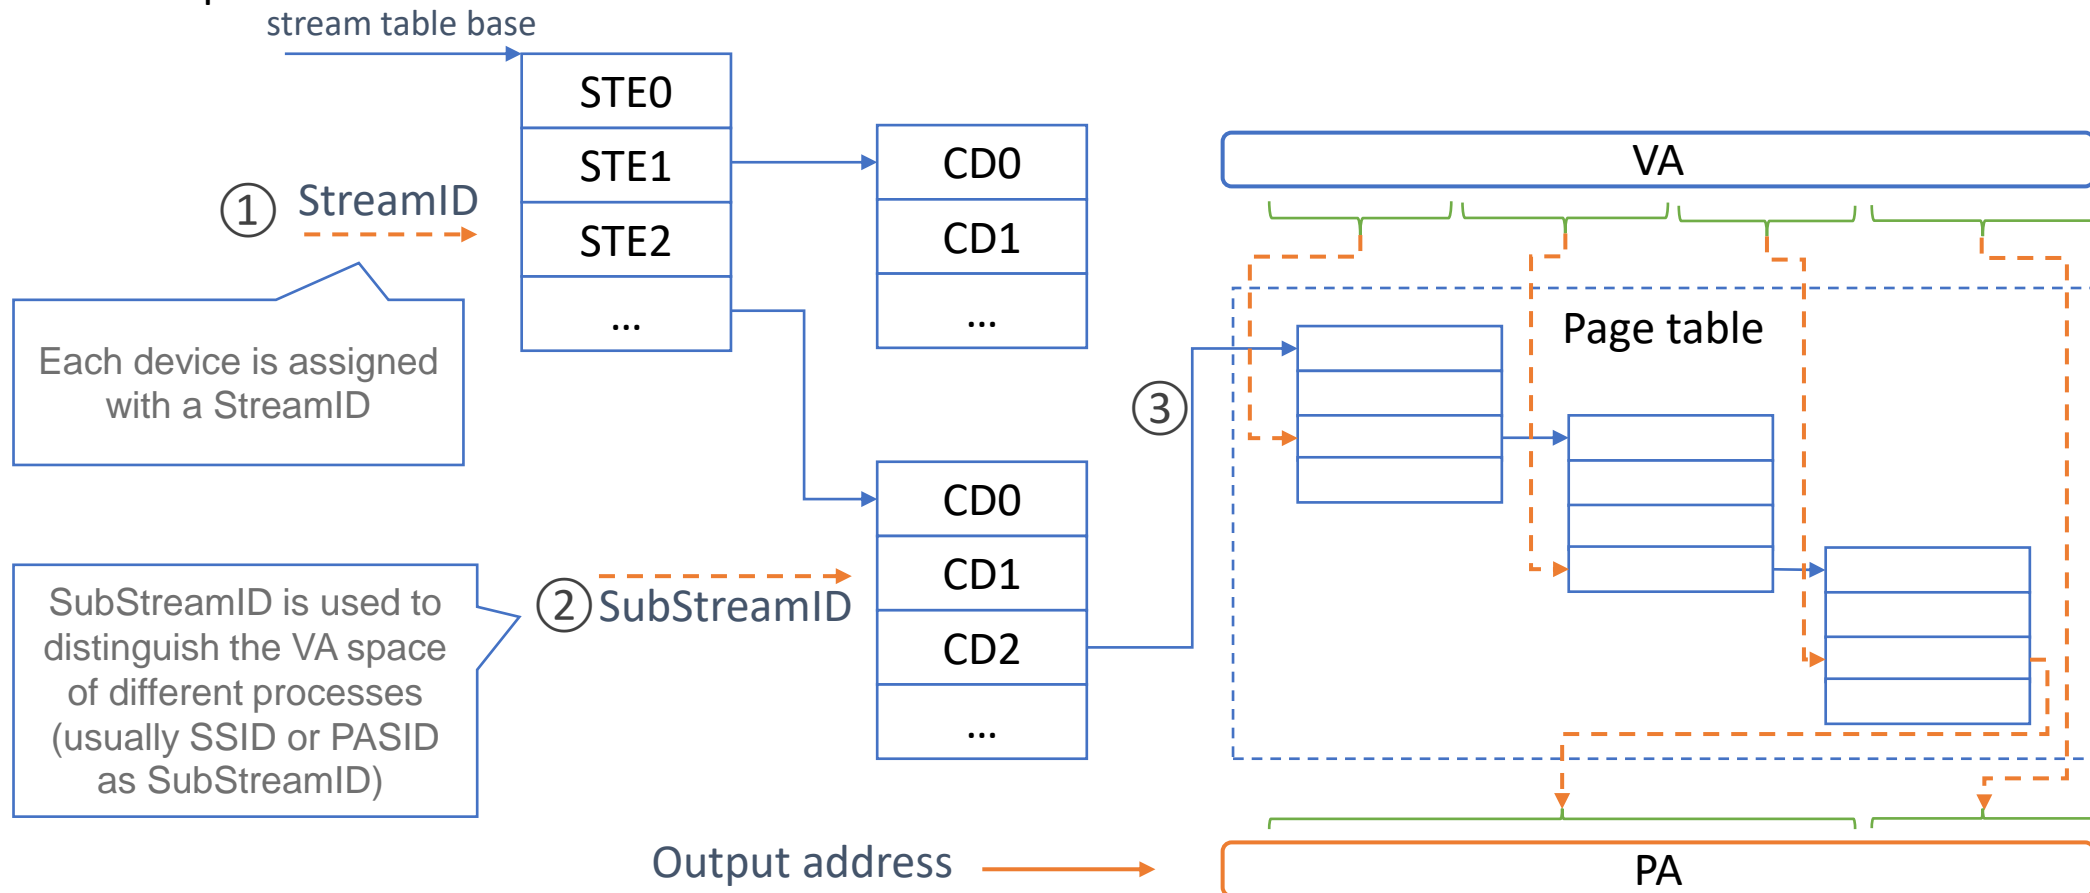
Introduction

- Similar to MMU:
 - performs address translation and access control
- But differently:
 - for bus initiators external to the CPU (DMA devices).
- Multi page-tables to support multi DMA devices




Address translation for DMA devices

- SMMU has a **stream table**, **context desc tables** and **page tables** (same with MMU)
- Access issued from a device contains **StreamID**, **SubStreamID** and **VA**
- Will output a **PA** after translation



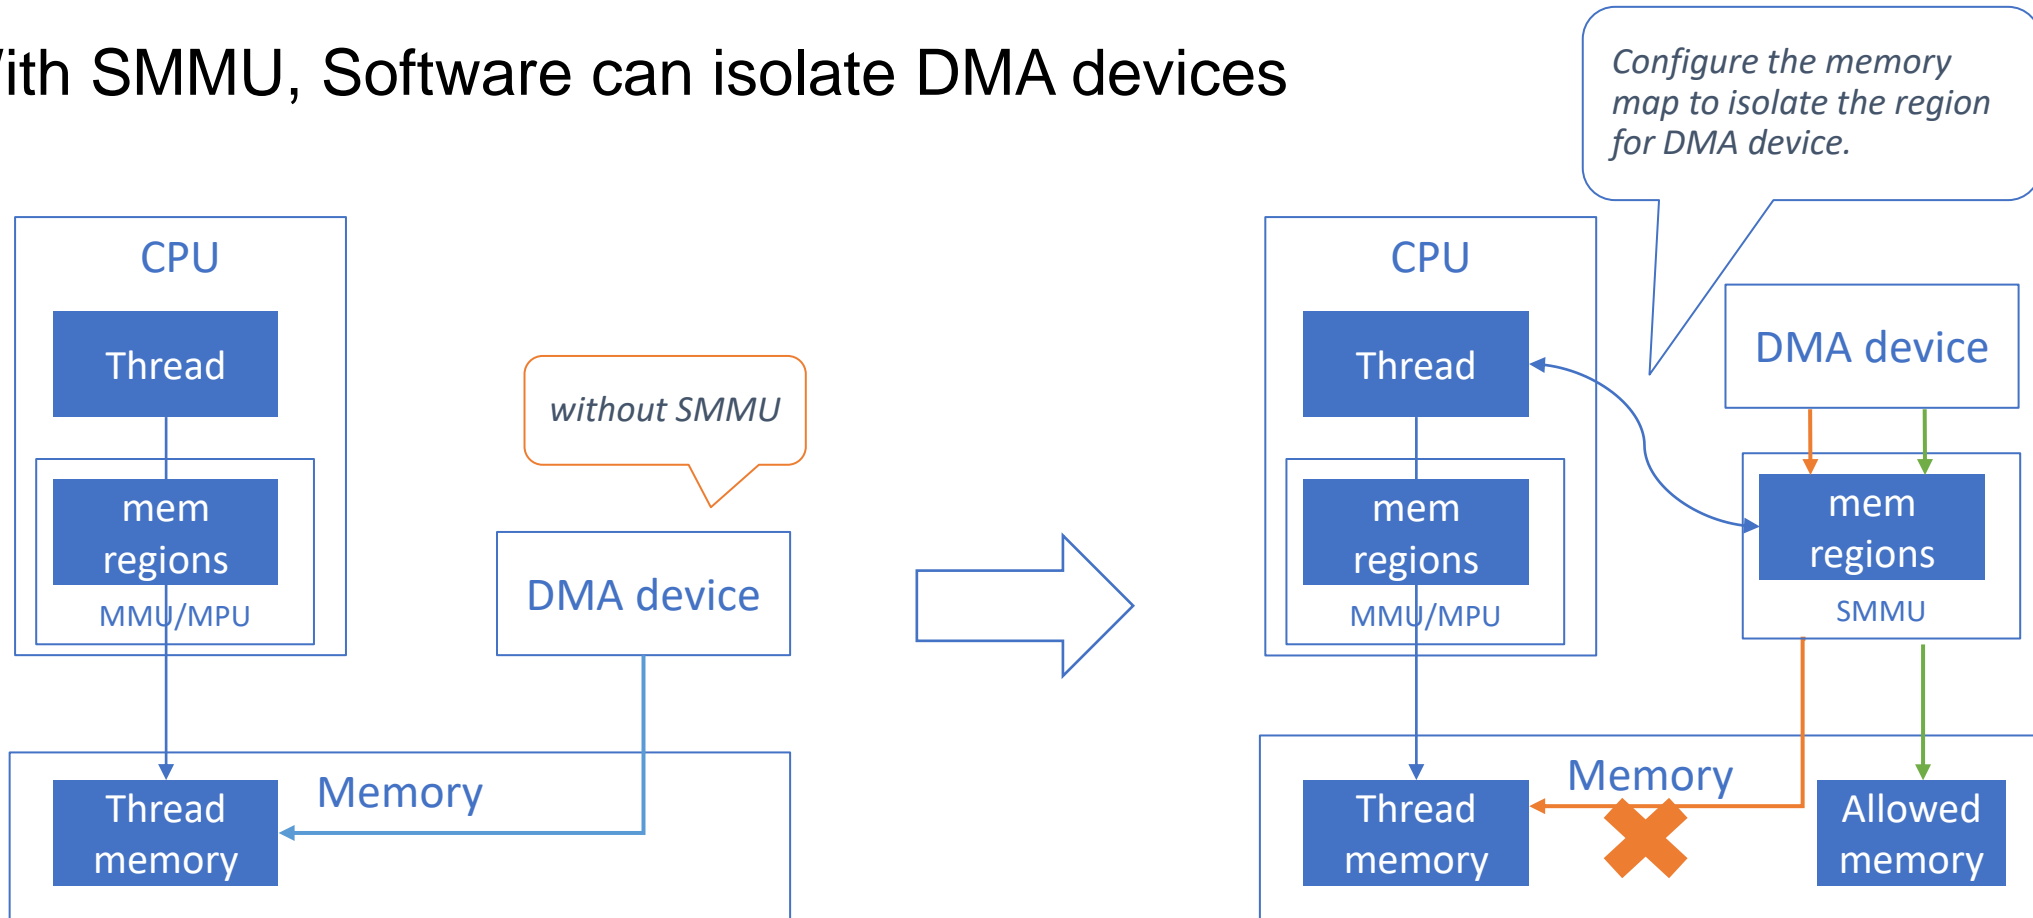
Configurable capabilities

- Bypass the translation (thus, VA = PA)
- Bypass Substream ID
- Reduce translation table level by using block attr(similar with MMU)
- Supports the 2nd stage translation (VA -> IPA -> PA)
 - for hypervisor (virtualization)



Isolate DMA devices

- With SMMU, Software can isolate DMA devices

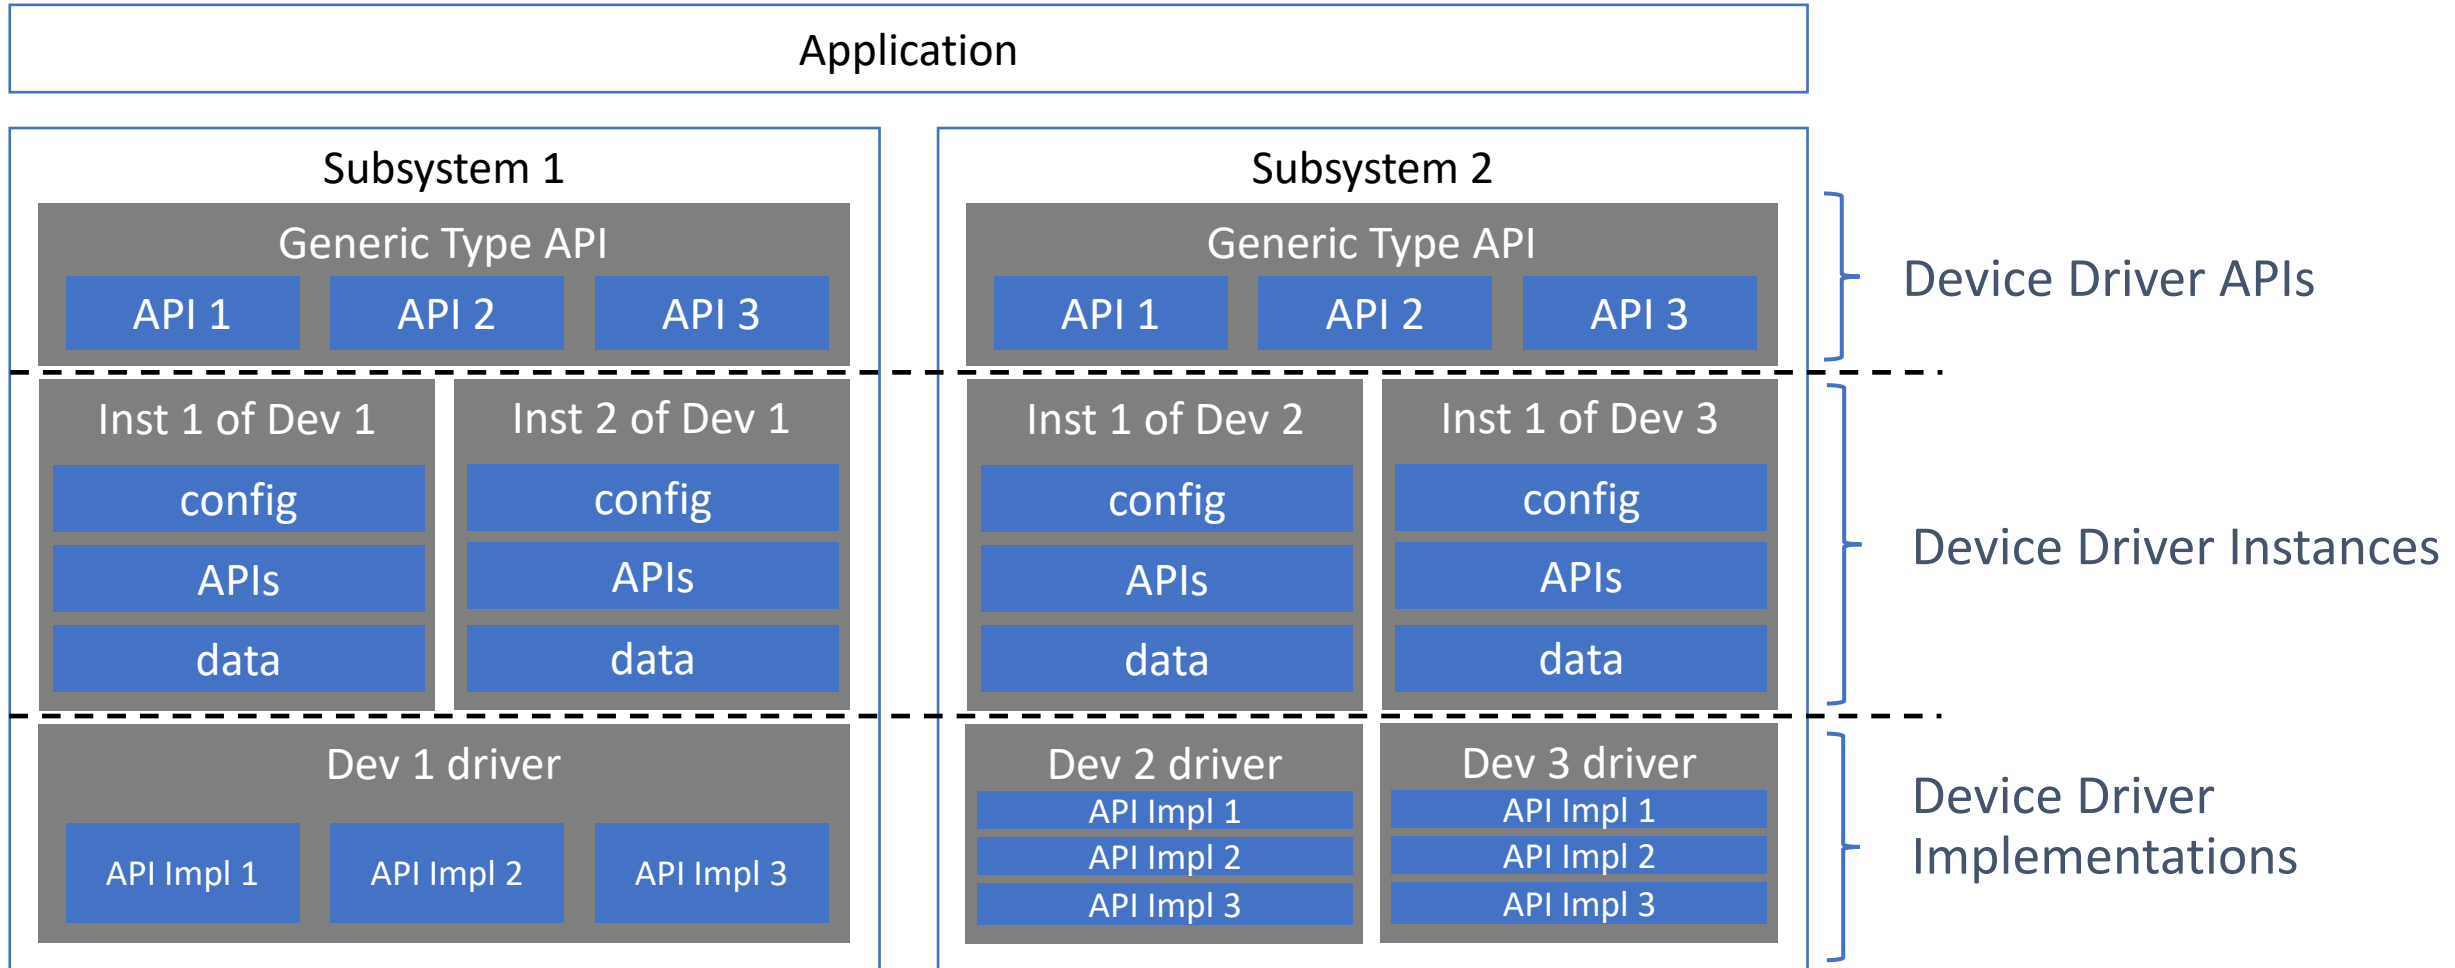


Contents

- Background
- SMMU
- **Zephyr Device Model**
- Zephyr HW-level device isolation
- Conclusion

Zephyr Device Model

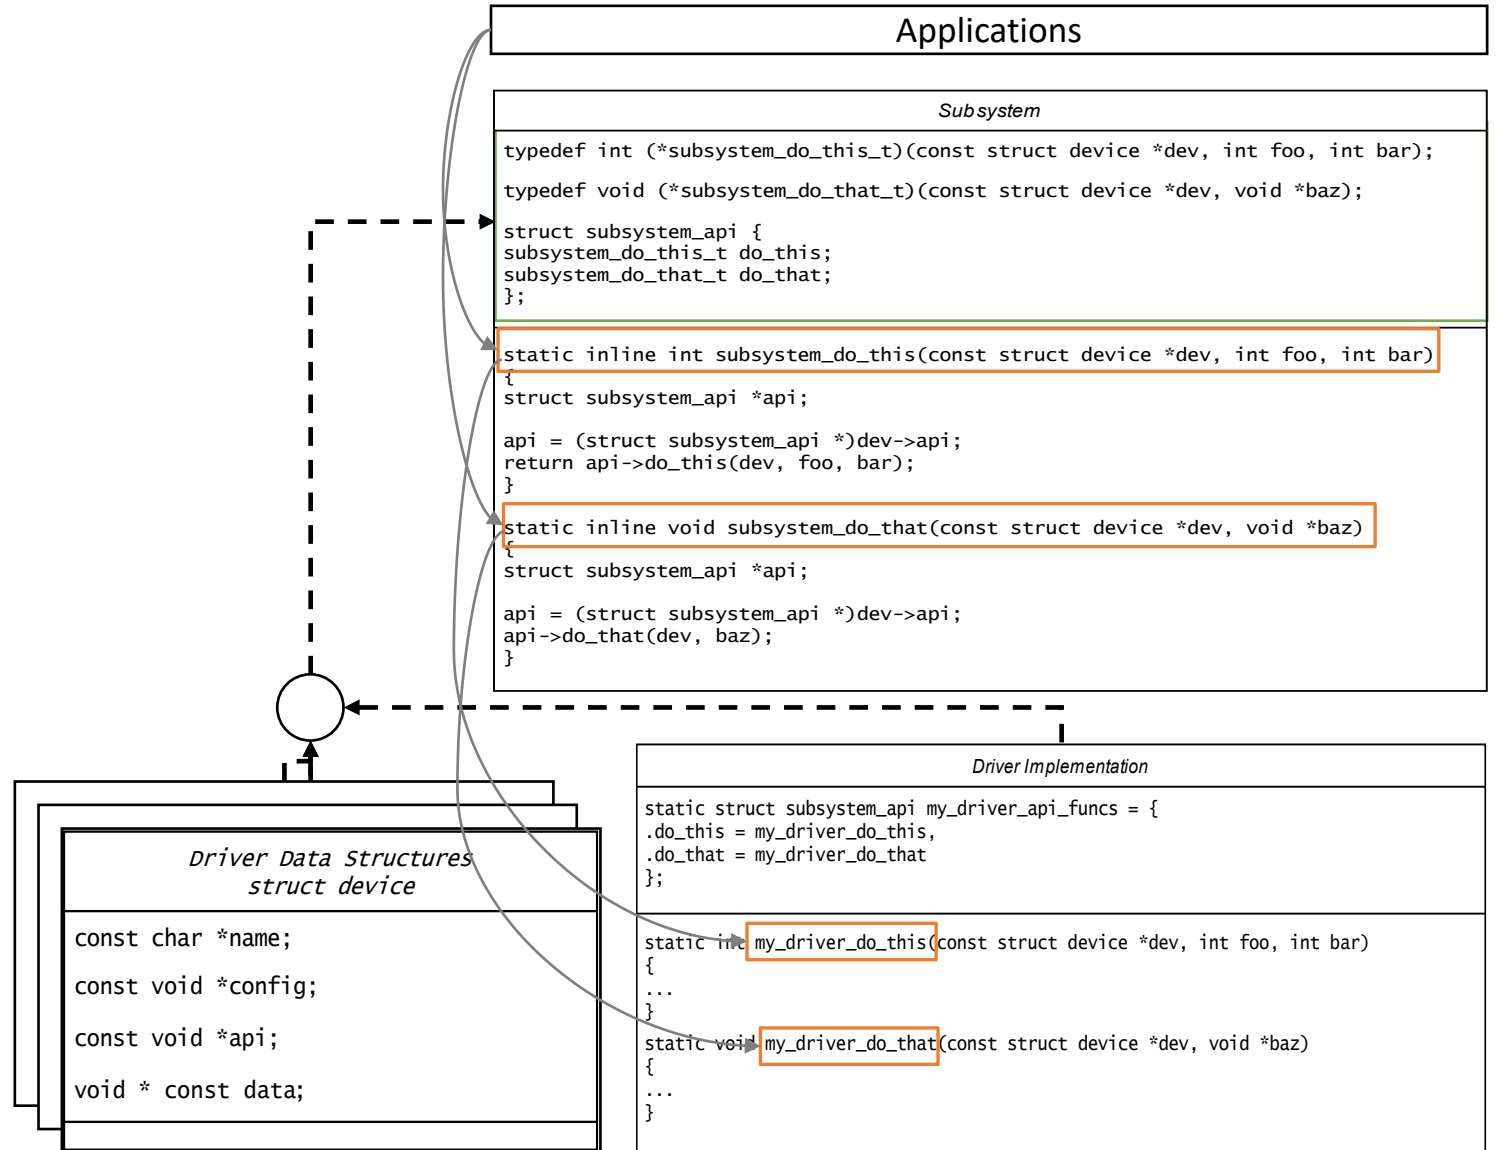
Overview



Zephyr Device Model

Overview

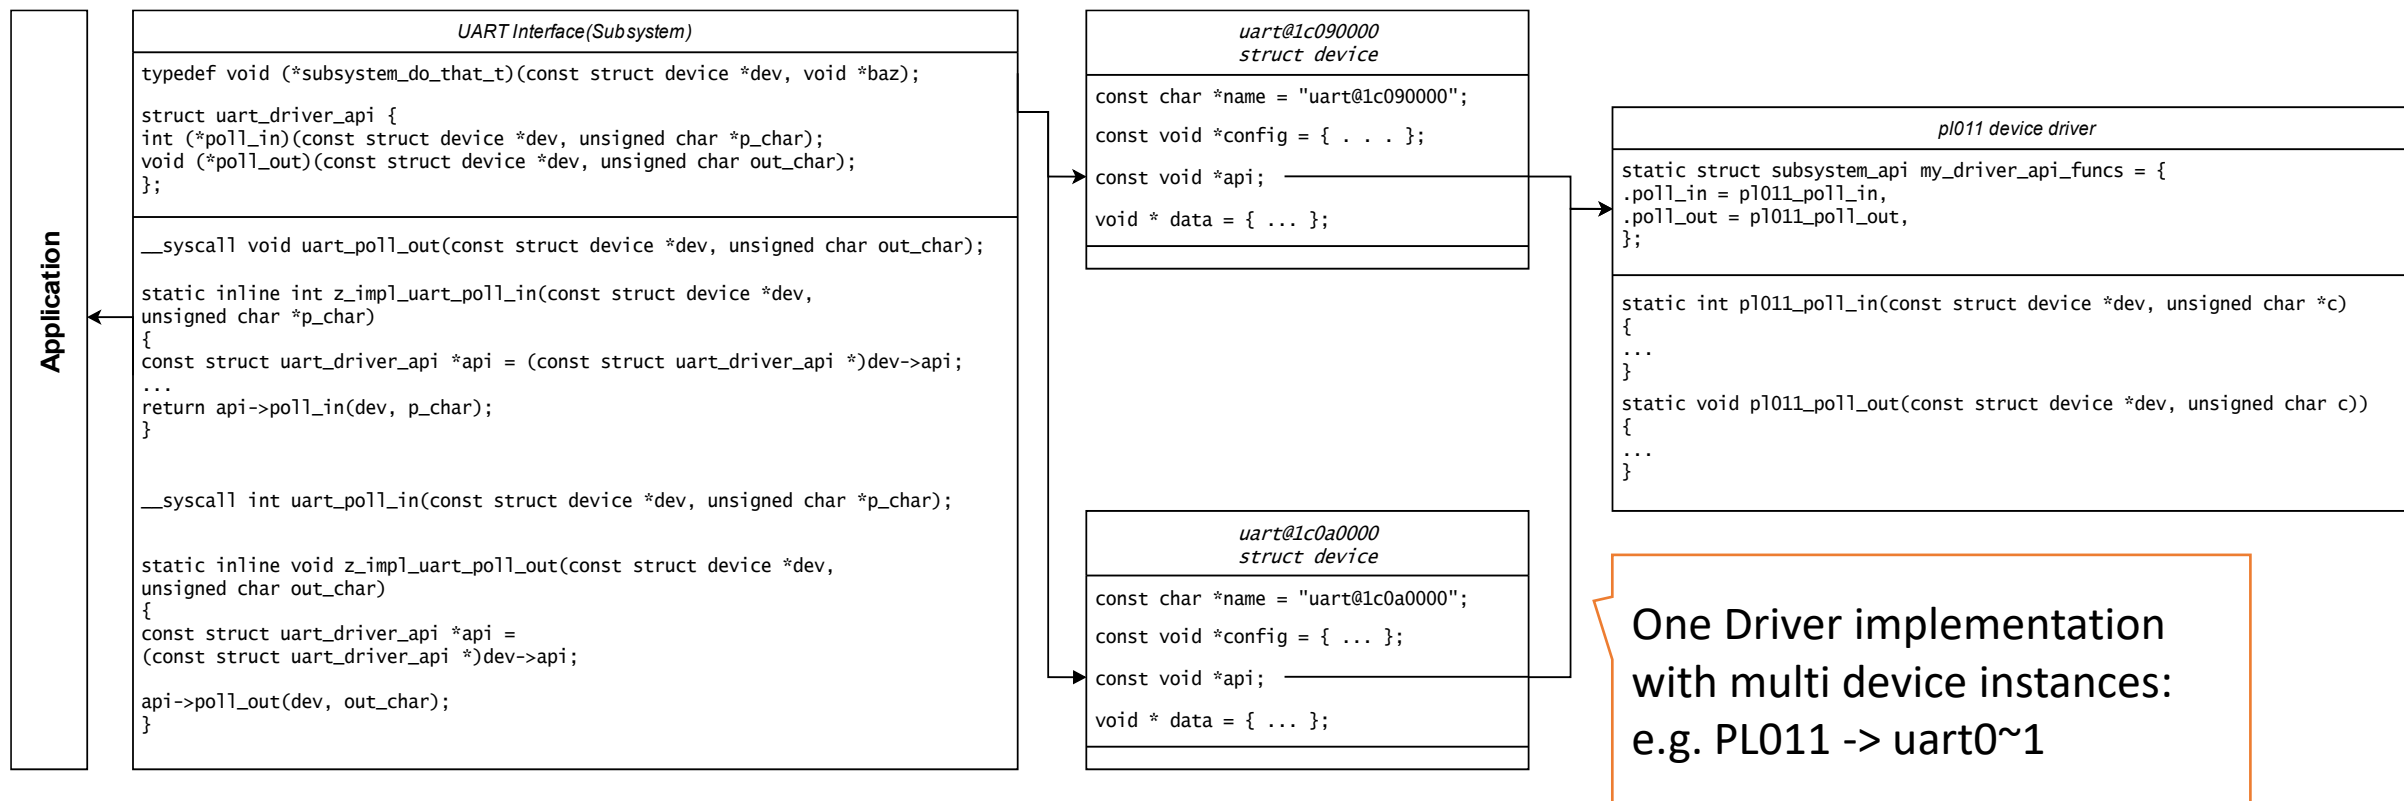
- **Subsystem**
 - Defines generic Device Driver APIs
 - simply program for applications
 - not specific to any particular driver implementation.
- **Driver Data Structures**
 - holds Devices data to support multiple devices driver instances
- **Device Driver Implementations**
 - implements a device-independent subsystem API
 - fills in the pointer to the Driver Data Structures.



Zephyr Device Model

A real example

- Single Driver, Multiple Instances



Contents

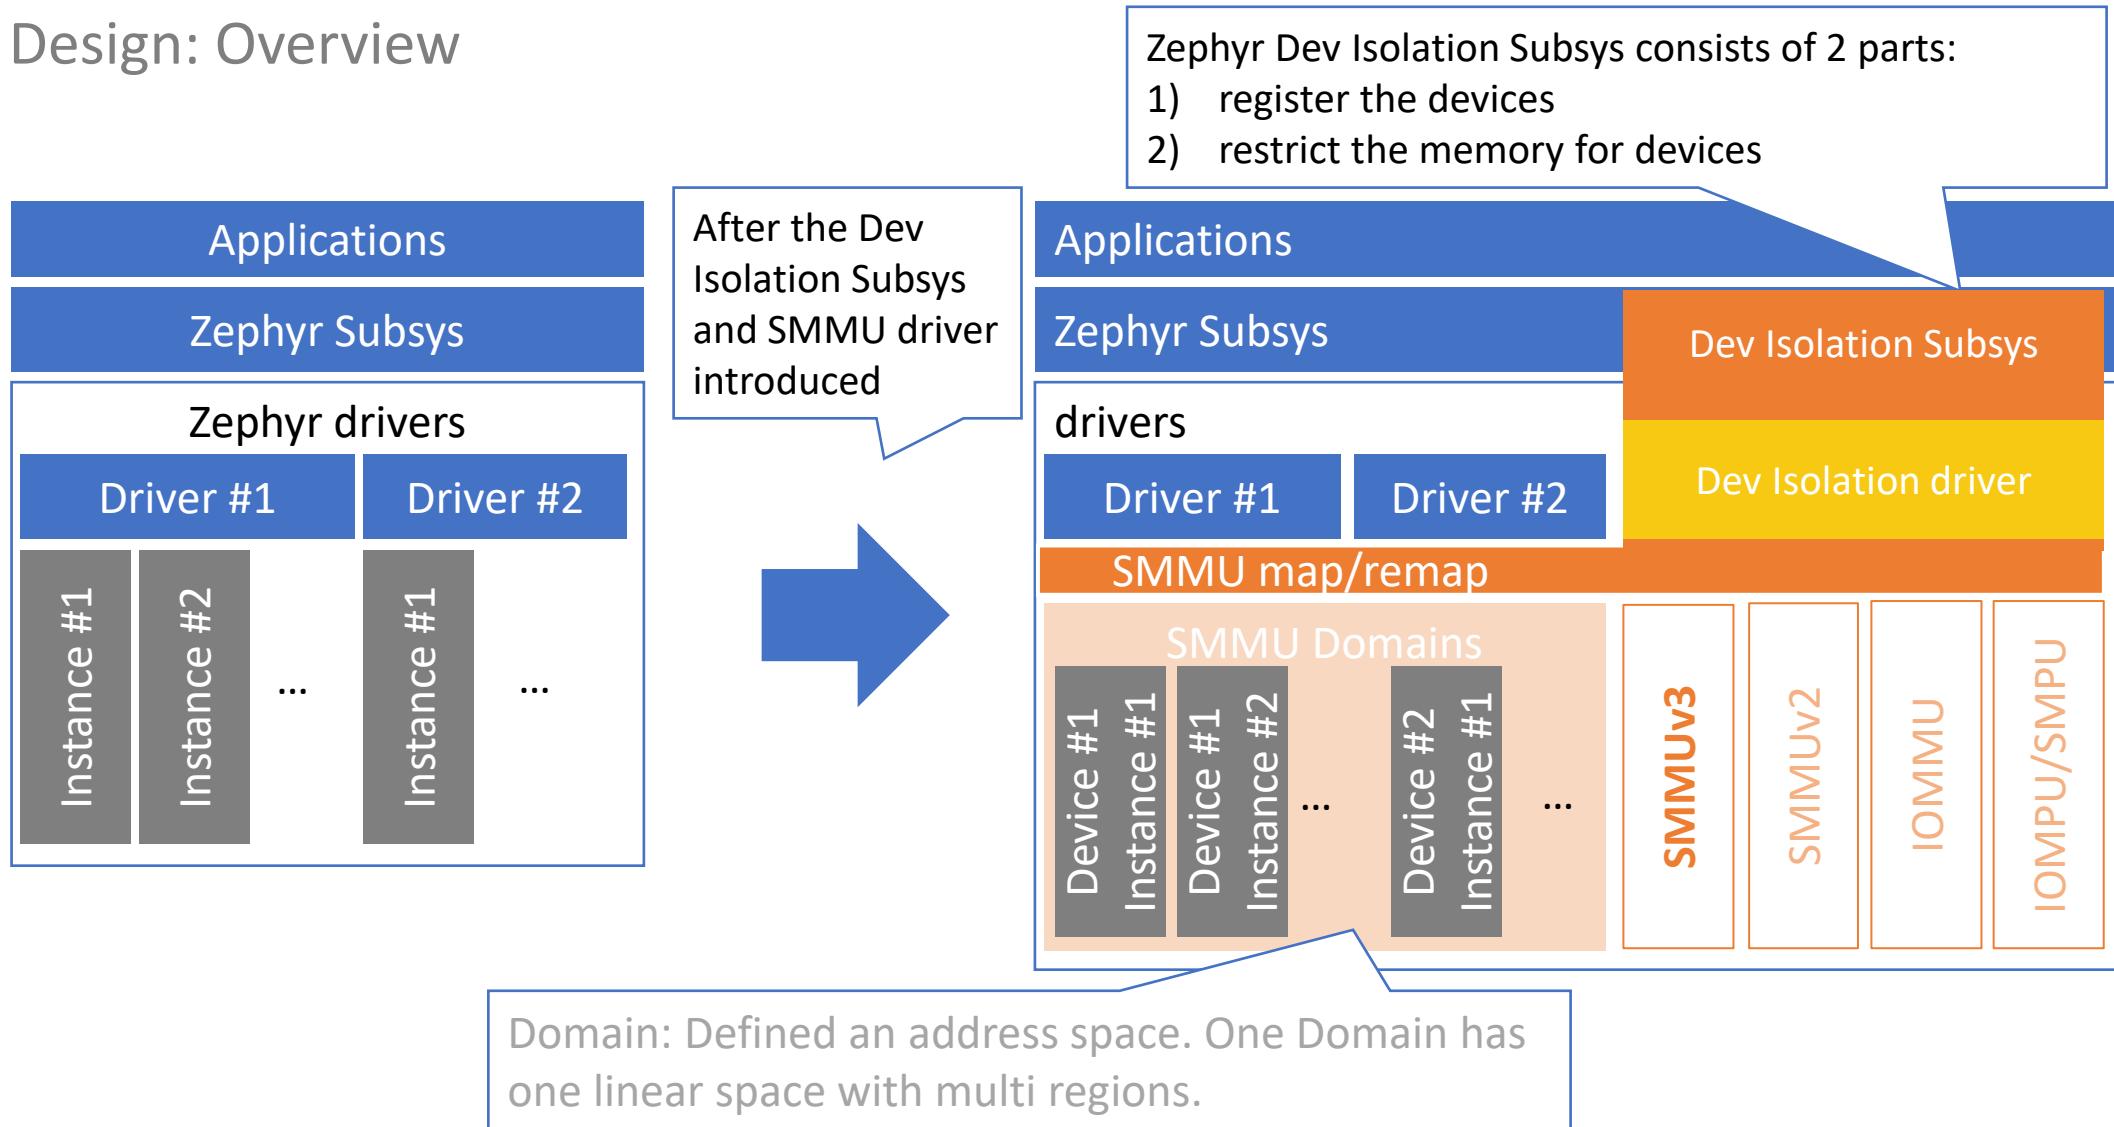
- Background
- SMMU
- Zephyr device model
- **Zephyr HW-level device isolation**
- Conclusion

Sub Contents

- The Design: Overview
- The design: DTS
- The Subsystem interface
- The implementation
- Add latency for DMA operations?

Zephyr HW-level device Isolation

The Design: Overview



Zephyr HW-level device Isolation

The design: DTS: PCI device + SMMUv3

- To make them work, the DTS should provide :

- *SMMU nodes:*

- *compatible*
- *base address*
- *length/size*

- *The DMA devices:*

- *the PCI & dev information*
- *SMMU phandle*
- *RID base address(BDF on PCI)*
- *StreamID*

```
pci: pci@40000000 {  
    compatible = "pci-host-ecam-generic";  
    reg = <0x40000000 0x10000000>;  
    msi-parent = <&its>;  
  
    #address-cells = <2>;  
    #size-cells = <1>;  
    ranges = <0x2000000 0x50000000 0x50000000 0x10000000>;
```

```
    smmu-maps = <&smmu 0 0 0x10000>;
```

```
    ahci: ahci0 {  
        compatible = "ata-ahci";  
  
        vendor-id = <0x0abc>;  
        device-id = <0xaced>;  
  
        status = "okay";  
    };  
};
```

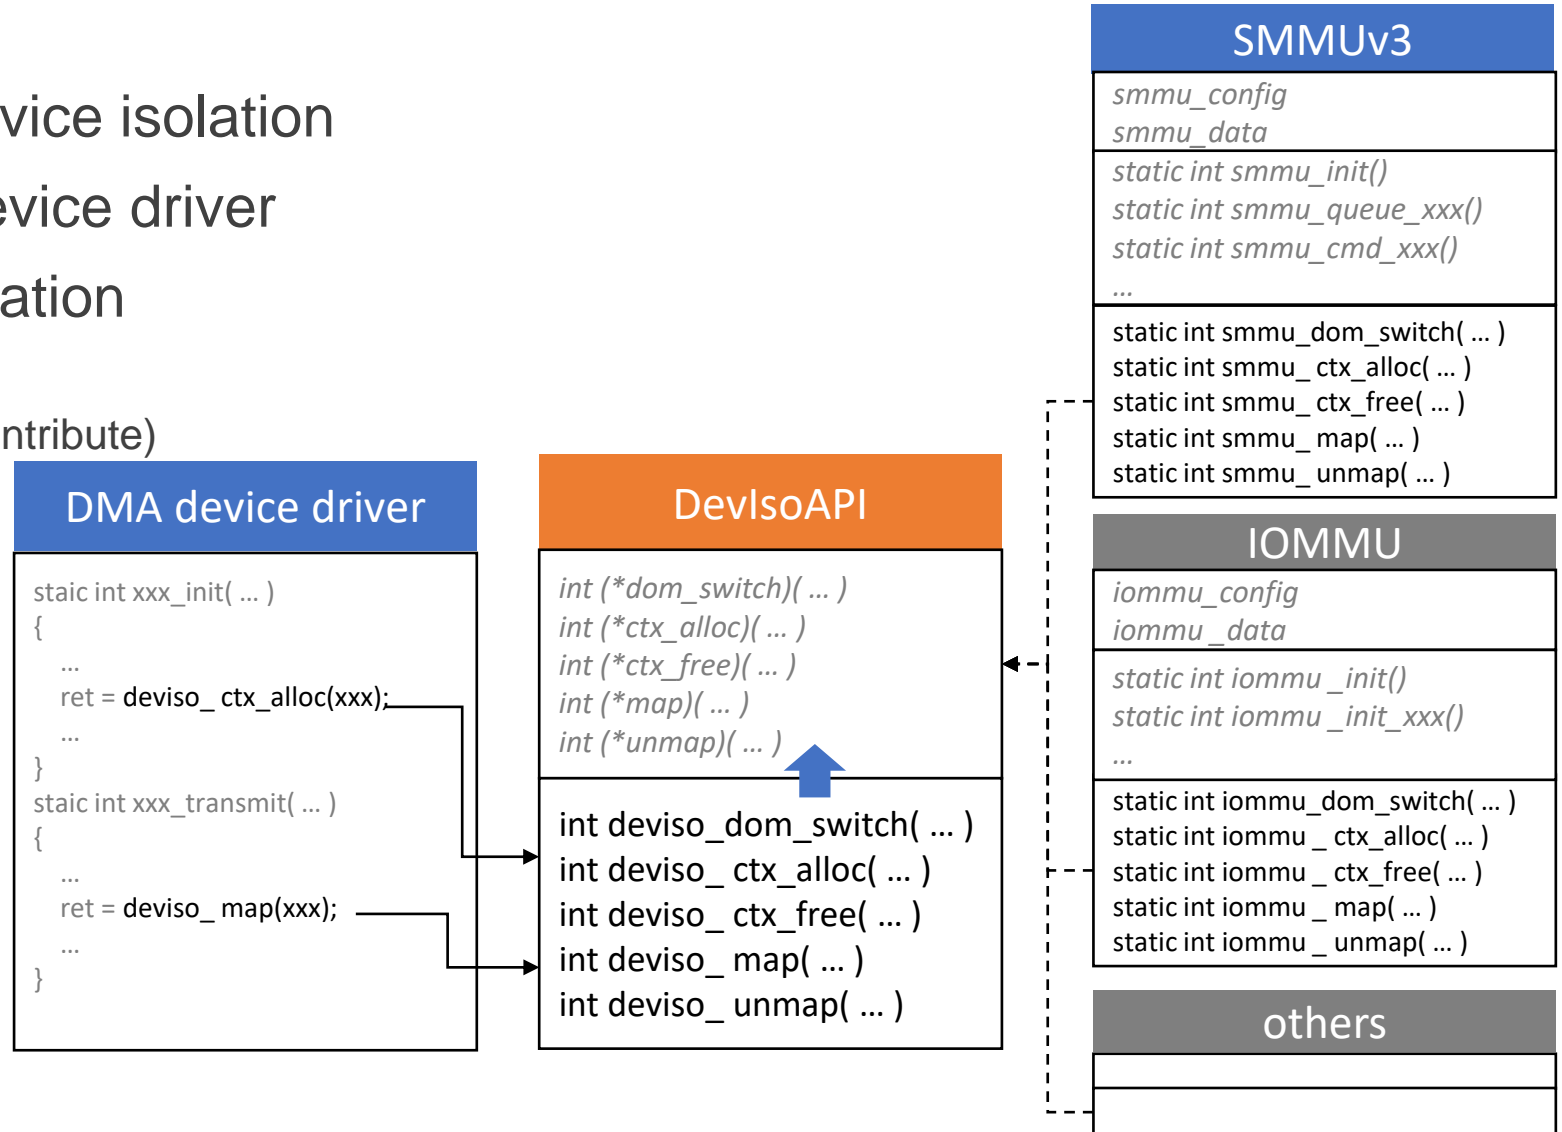
```
smmu: smmu@2b400000 {  
    compatible = "arm,smmu-v3";  
    reg = <0x2b400000 0x100000>;  
    #smmu-map-cells = <3>;  
};
```



Zephyr HW-level device Isolation

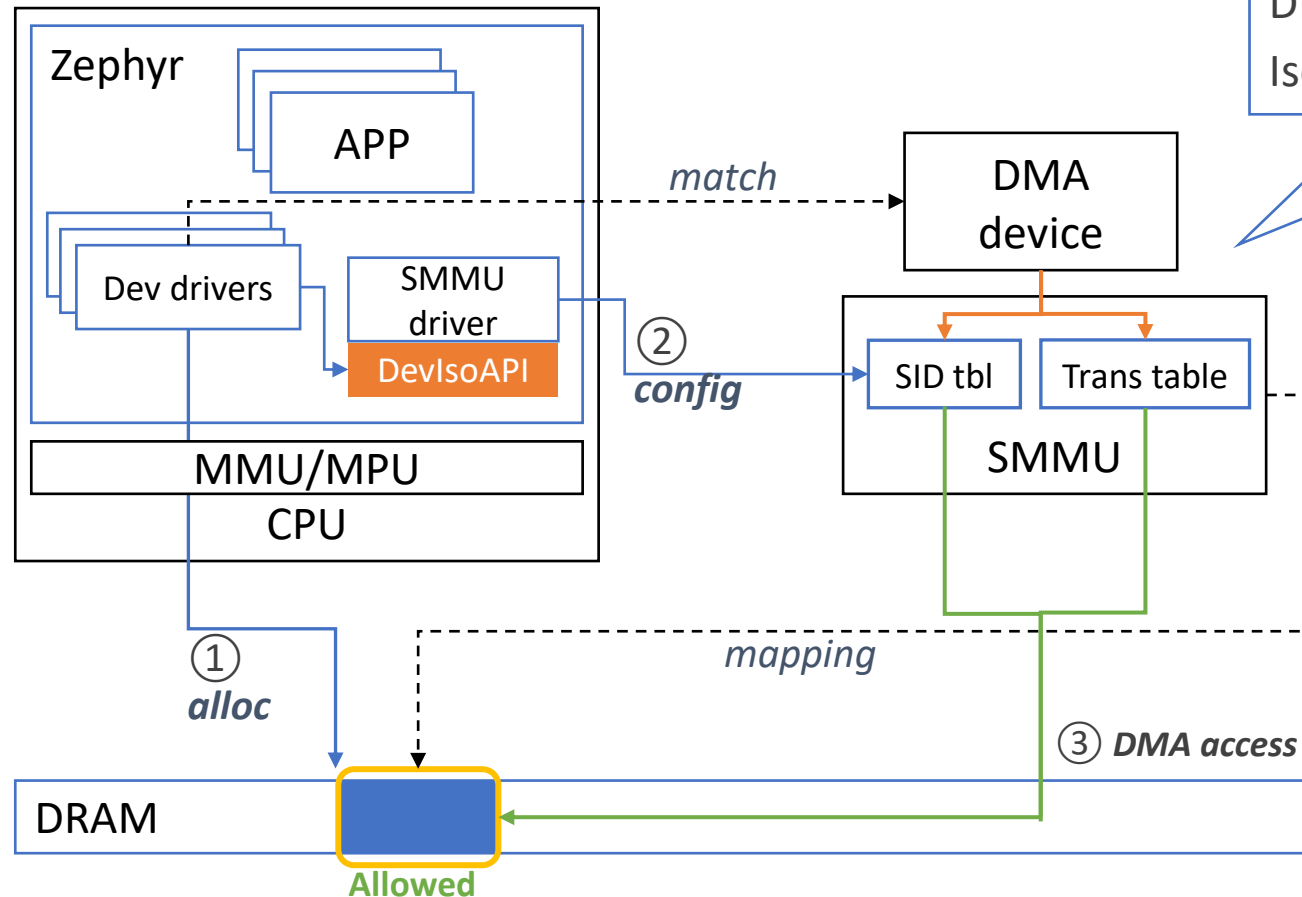
The Subsystem interface

- Defines generic APIs for device isolation
- Simply program for DMA device driver
- Support multiple implementation
 - SMMUv3
 - IOMMU (welcome anyone to contribute)
 - ...
- Easy to add new one



Zephyr HW-level device Isolation

The implementation (PoC)

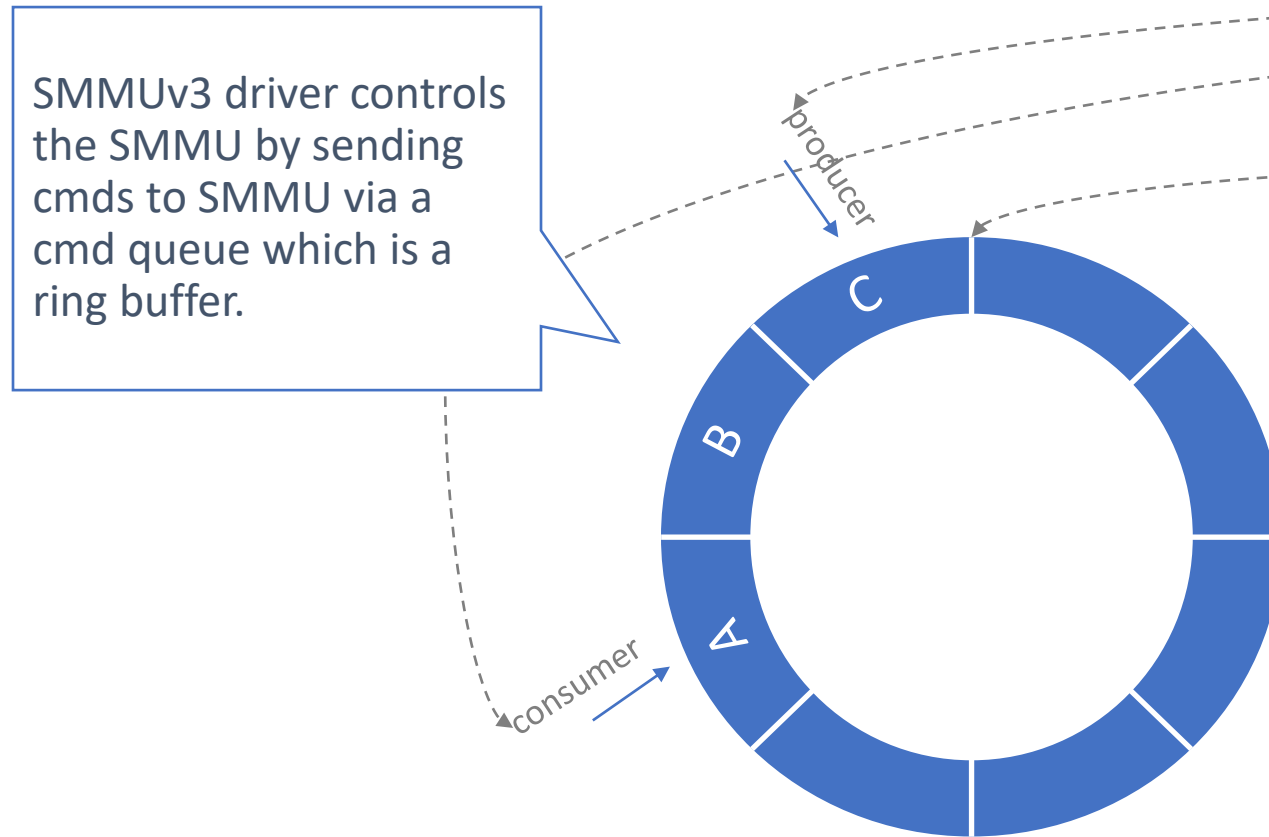


board: fvp_base_revc_2xaemv8a
DMA dev: PCI + AHCI
Isolation driver: SMMUv3

- ① Dev driver allocate a DMA buffer for device.
- ② Dev driver call DevIsoAPI to map the buffer for the device.
- ③ Device copy data to DRAM.

Zephyr HW-level device Isolation

The implementation: SMMUv3 control



```
struct smmu_queue_local_copy
{
    uint32_t prod;
    uint32_t cons;
}
```

```
struct smmu_cmdq_entry
{
    uint8_t opcode;
    struct tlbi;
    struct cfgi;
    struct prefetch;
    struct sync;
}
```

```
struct smmu_queue
{
    struct smmu_queue_local_copy lc;
    void *base;
    mem_addr_t base_dma;
    mm_reg_t prod_reg;
    mm_reg_t cons_reg;
    uint64_t q_base;
    int size_log2;
}
```

```

smmu command queue functions

int smmu_q_has_space(struct smmu_queue *q)
uint32_t smmu_q_inc_prod(struct smmu_queue *q)
void make_cmd(uint64_t *cmd, struct smmu_cmdq_entry *entry)
void smmu_cmdq_enqueue_cmd(struct smmu_device_data *data,
    struct smmu_cmdq_entry *entry)

int smmu_sync(struct smmu_device_data *data)
void smmu_invalidate_sid(struct smmu_device_data *data,
    uint32_t sid)

void smmu_prefetch_sid(struct smmu_device_data *data,
    uint32_t sid)

void smmu_invalidate_all_sid(struct smmu_device_data *data)
void smmu_tlbi_all(struct smmu_device_data *data)

```

Add latency for DMA operations?

- SMMU (lower down the tlb miss):
 - statically linear address mapping
 - use block page
 - skip sub stream id
 - appropriately use ATS (Address translation services, platform specific)
 - try not switch translation tables

Contents

- Background
- SMMU/IOMMU
- Zephyr device model
- Zephyr HW-level device isolation
- **Conclusion**

Conclusion

- DMA can break the system
 - buggy
 - malicious
- DMA device + Zephyr RTOS risk is increasing
 - more DMA on low-power platforms
 - more RTOS on high-performance platforms
- Add HW-level device isolation
 - enable SMMU but lower the overhead
 - add subsys for easy extension
- Future work:
 - send to upstream
 - measure the latency

Thank you!

Q&A