



Zephyr[®] Project

Developer Summit



Zephyr® Project
Developer Summit

Pigweed Tokenizer

AI Semjonovs, *Google*

Email: asemjonovs@google.com

Discord: asemjonovs#8232

Github: asemjonovs



#EMBEDDEDOSSUMMIT

Agenda

- What is Pigweed's tokenizer?
- Zephyr's Dictionary Logging Comparison
- How does tokenizing help us?
- How to integrate Pigweed's tokenization into your project
- Future improvements with tokenization



What is [Pigweed's tokenizer](#)?

- Not related to string parsing
- Replaces whole string literal with a 32-bit hash token.
- Why?
 - Reduce binary size by removing string literals from binaries
 - Reduce I/O traffic, RAM, and flash usage.
 - Reduce CPU usage by replacing sprintf calls with simple tokenization code.





	Zephyr Dictionary Logging	Pigweed Tokenizer								
String mapping	Maps to string address	Maps to 32-bit hash generated by string literal								
Probability of Collisions	Strings have a 1:1 mapping to address. Guaranteed no collisions.	<table> <tr> <th rowspan="2">Token bits</th><th colspan="2">Collision probability by string count</th></tr> <tr> <th>50%</th><th>1%</th></tr> <tr> <td>32</td><td>77000</td><td>9300</td></tr> </table>	Token bits	Collision probability by string count		50%	1%	32	77000	9300
Token bits	Collision probability by string count									
	50%	1%								
32	77000	9300								
Database Format	JSON dictionary	CSV, Binary, Directory based								
Database Portability	Addresses are not guaranteed between builds. Only works with build it was compiled against.	The token hash will be the same between builds and boards using same source. Able to merge tokens from multiple boards and versions to a single database.								


Token Database

- Stores mapping of hash tokens to strings they represent
- Generated from ELF file
- [Database Format Types](#)
 - CSV
 - Binary
 - Directory Based
- Update an existing database
 - Captures removal date of unused tokens
- Integrates with CMake and GN builds



Detokenizing

- Language support
 - Python
 - C/C++
 - Typescript
- [Web Console](#)
- [System Console](#)
 - ```
pigweed$ python3 -m pw_system.console --device /dev/ttyUSB0 --token-databases database.csv
```





# Tokenized Logging Example

Before: Plain text logging

| Location         | Logging Content                                                                | Num bytes |
|------------------|--------------------------------------------------------------------------------|-----------|
| Source contains  | <code>LOG("Battery state: %s; battery voltage: %d mV", state, voltage);</code> |           |
| Binary contains  | <code>"Battery state: %s; battery voltage: %d mV"</code>                       | 41        |
| Device Transmits | <code>"Battery state: CHARGING; battery voltage: 3989 mV"</code>               | 49        |
| When Viewed      | <code>"Battery state: CHARGING; battery voltage: 3989 mV"</code>               |           |





# Tokenized Logging Example

After: Tokenized logging

| Location         | Logging Content                                                                                                                                                         | Tokenized Num bytes | Plain Text Num bytes |
|------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------|----------------------|
| Source contains  | <code>LOG("Battery state: %s; battery voltage: %d mV", state, voltage);</code>                                                                                          |                     |                      |
| Binary contains  | <code>d9 28 47 8e</code> (0x8e4728d9)<br>(log statement is called with "CHARGING" and 3989 as arguments)                                                                | 4                   | 41                   |
| Device Transmits | <code>d9 28 47 8e</code> → Token (4 bytes)<br><code>43 48 41 52 47 49 4e 47 00</code> → "CHARGING" argument (9 bytes)<br><code>0f 95</code> → 3989, as varint (2 bytes) | 15                  | 49                   |
| When Viewed      | "Battery state: CHARGING; battery voltage: 3989 mV"                                                                                                                     |                     |                      |


## How is tokenizing helping us?

- We're applying tokenization to our logging module.
- Within our Embedded Controller, we have seen a reduction in 14KBs of flash memory, a 6% reduction in our image size.
- Debug logging can be more verbose with tokenization.



## Database Management

- How do you want to use your token database?
  - 1:1
  - 1 to many
- Single database supporting multiple EC boards and versions of firmware
  - Chromebooks 8 year support cycle.
- Database size
  - Per EC board: 40 KBs
  - Multiple EC board support: 66 KBs
  - Approximately 1 KB growth per board



## How to integrate Pigweed tokenization into your Zephyr project

- It's simple!
- Include
  - `pigweed/Kconfig.zephyr`
- Enable
  - `CONFIG_STD_CPP20=y`
  - `CONFIG_ZEPHYR_PIGWEED_MODULE=y`
  - `CONFIG_PIGWEED_LOG_TOKENIZE=y`
- Update CMake
  - Add database creation dependency



## Future improvements with tokenization in EC

- Tokenized strings as %s arguments
  - State names
  - Enumerations
- Tokenized RPC logging



# Thanks!

