

Using the Thrift RPC Framework in Zephyr

2023-06-28: Embedded Open Source Summit

Chris Friedt
Embedded SWE, Meta
Thrift Module Maintainer






Word cloud containing various technical terms and project names:

- pthread
- kernel
- libc
- net
- dns_sd
- arm
- i2c
- gpio
- logging
- Devicetree
- boards
- mdns
- clock
- sensor
- thread
- riscv
- atomsics
- cc13XX_cc26XX
- timers
- Kconfig
- sched
- spi
- bluetooth
- greybus
- ring_buffer
- perf
- release
- pinctrl
- emul
- soc
- cc1352r1_launchxl
- cc1352r_sensortag
- boards
- gpio
- logging
- Devicetree
- mdns
- clock
- sensor
- thread
- riscv
- atomsics
- cc13XX_cc26XX
- timers
- Kconfig
- sched

yocto.
PROJECT





<https://bit.ly/442kJrJ>

Agenda

01 What is Thrift?

02 Thrift in Zephyr

03 Status Update & Additional Work



01 What is Thrift?

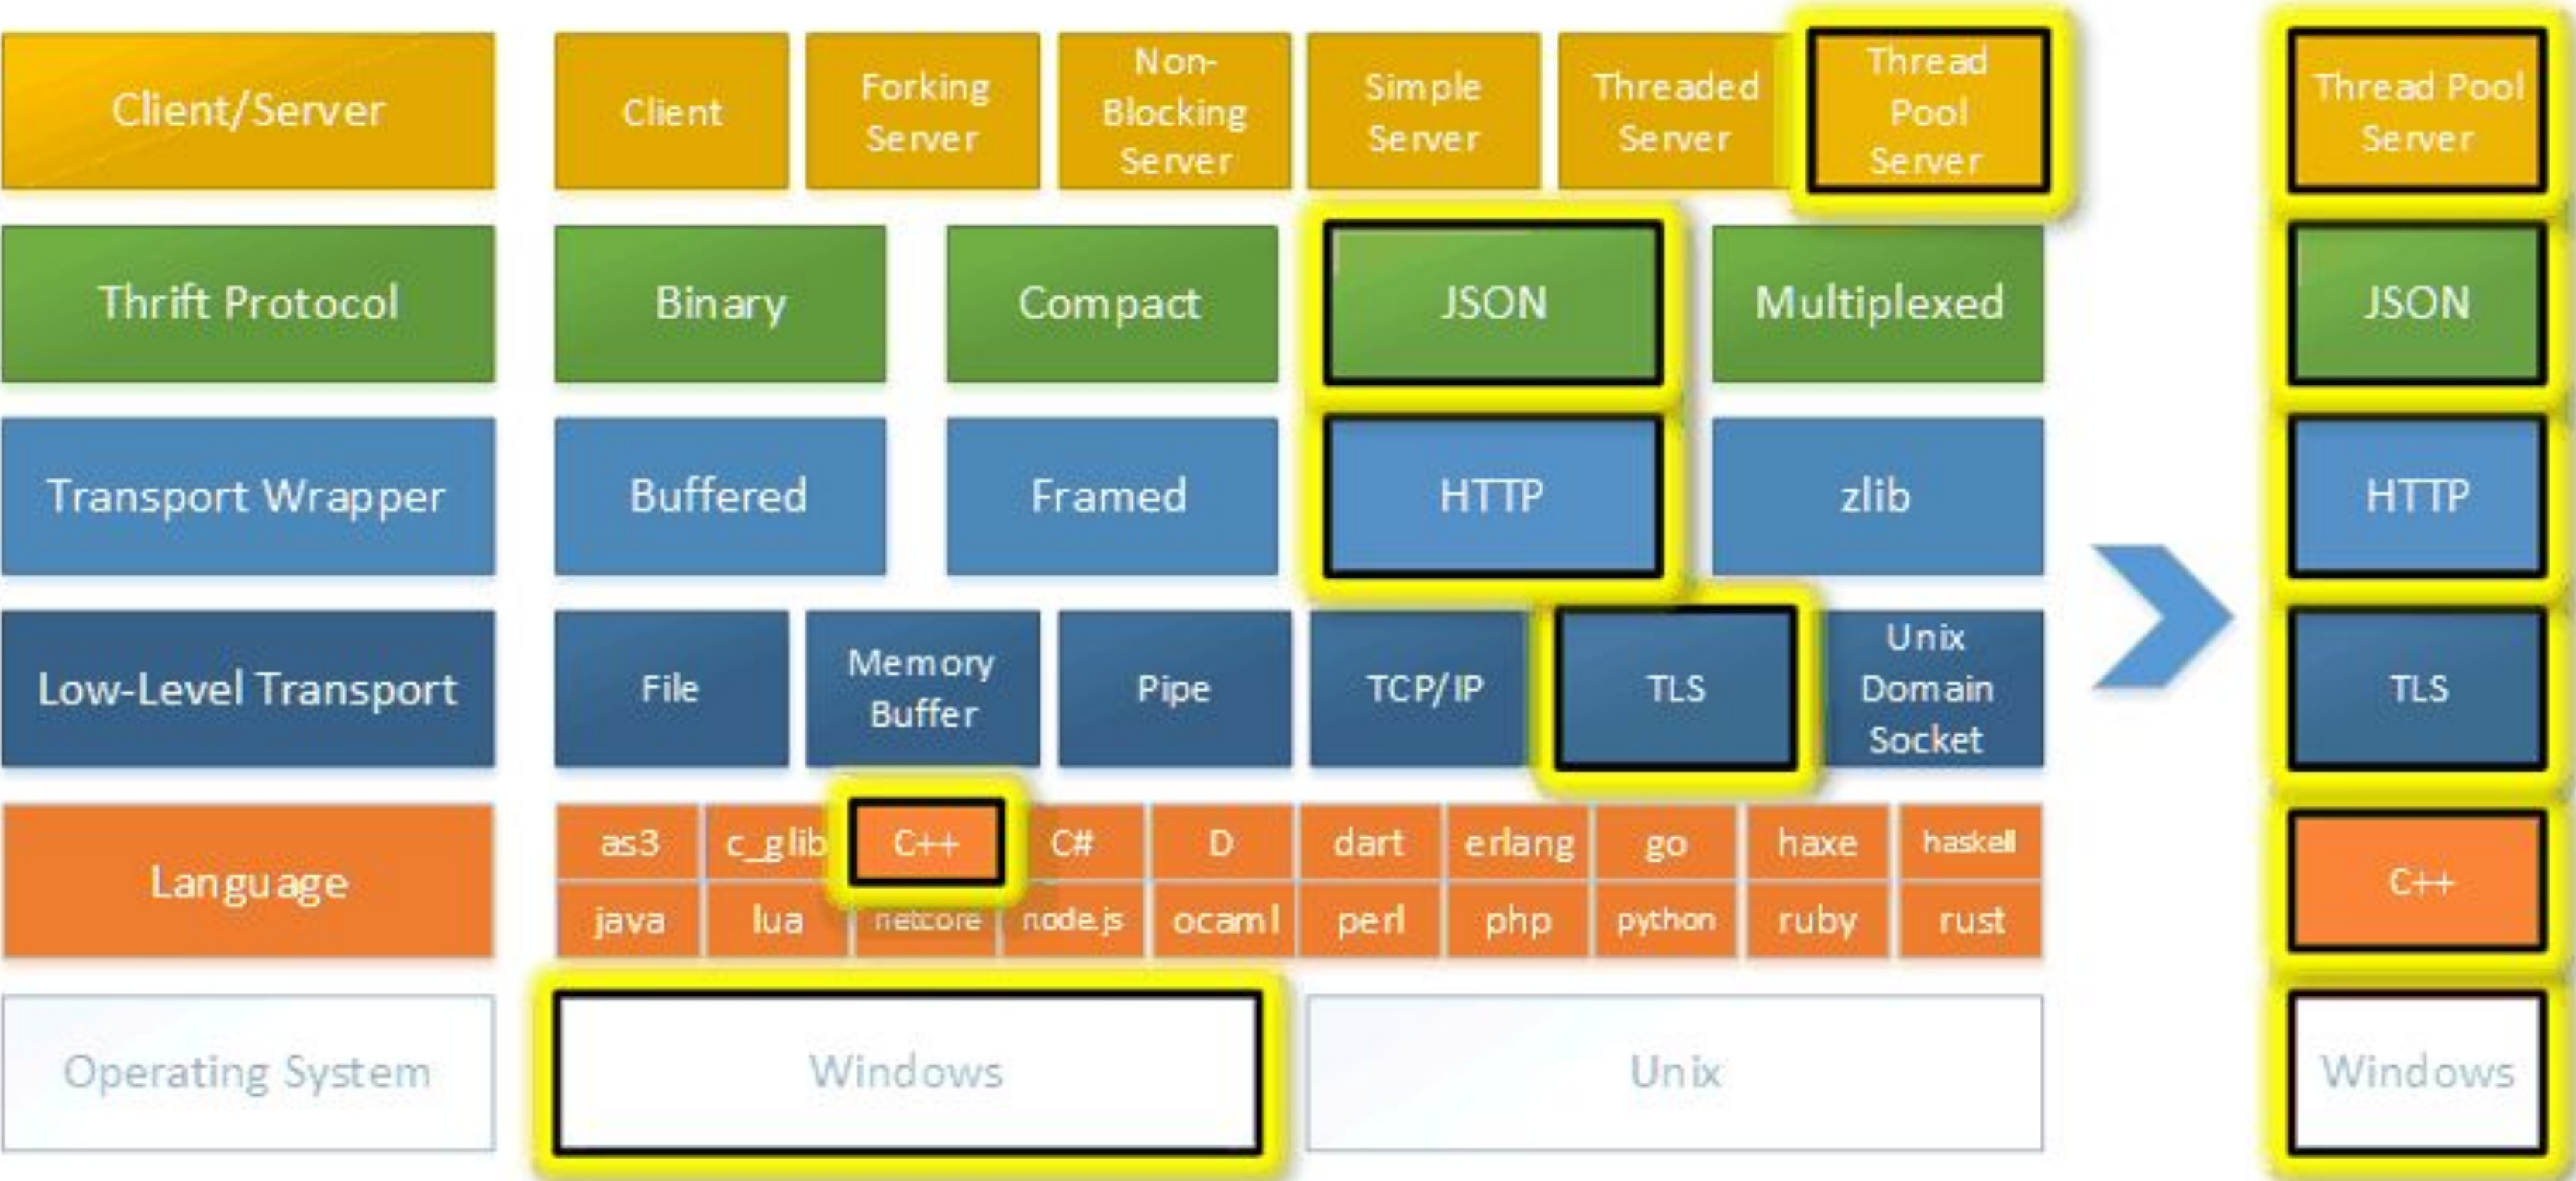


<https://bit.ly/44nXnwq>



<https://bit.ly/433nA2p>

Apache Thrift Layered Architecture



Remote Procedure Call Frameworks

- RPC is nothing new
 - Solution for ad-hoc, hand-rolled, bit-stuffed “protocols”
 - Solution for re-inventing TCP 1000x times
 - Reliable asynchronous communication is hard
 - [ASN.1](#) (1984, ITU-T) now [ISO/IEC 8825-1:2021](#)
 - [Sun RPC](#) (1984-1986, part of [NFS](#))
 - Part of glibc (and others) to this day ([rpcgen](#))
 - TLV (Tag-Length-Value) representation
 - E.g. [Netlink](#), [MISB](#) (oddly I did [MISB in Thrift!](#))
- Some issues with SunRPC
 - every message (procedure) required a version
 - absolute nightmare for inter-version compatibility
- Protobufs, Thrift, and gRPC started around the “[downfall of Sun](#)”
 - Drop message versions (just use new Tag)
 - Do not reuse Tags
 - Guaranteed consistence +/- one rev (i.e. commit)

```
/*
 * The directory program definition
 */
program DIRPROG {
    version DIRVERS {
        readdir_res
        READDIR(nametype) = 1;
    } = 1;
} = 0x20000076;
```

<https://bit.ly/4411BdP>



01 WHAT IS THRIFT?



Thrift History / Internals

- Created at Facebook (now Meta) 2006-2007
- Inspired by [Protocol Buffers](#)
- Used extensively throughout Meta infra
- Released to the [ASF in 2007](#), original [whitepaper](#)
- Adopted by several cloud companies, some OSS projects
- Not only primitive types such as *bool*, *i32*, and *double*, but also rich types such as *enum*, *list*, *map*, and *struct*
- Not only [Binary](#), [Compact Binary](#), but also e.g. [JSON](#)
- Procedures (methods) can *throw* any kind of *struct* and take any number of arguments
- All OS's, 27 Languages,
- Facebook [forked Thrift in 2014](#) 😊 to create [fbthrift](#)

- `BOOL`, encoded as `2`
- `I8`, encoded as `3`
- `DOUBLE`, encoded as `4`
- `I16`, encoded as `6`
- `I32`, encoded as `8`
- `I64`, encoded as `10`
- `BINARY`, used for binary and string fields, encoded as `11`
- `STRUCT`, used for structs and union fields, encoded as `12`
- `MAP`, encoded as `13`
- `SET`, encoded as `14`
- `LIST`, encoded as `15`
- `UUID`, encoded as `16`

List and Set

List and sets are encoded the same: a header indicating the size and the element-t

Binary protocol list (5+ bytes) and elements:

```

+-----+-----+-----+-----+-----+-----+-----+-----+
| t t t t t t | size | elements |
+-----+-----+-----+-----+-----+-----+

```

<https://bit.ly/3CS4PV8>



<https://github.com/facebook/fbthrift>

```

enum PhoneType {
    HOME,
    WORK,
    MOBILE,
    OTHER
}

struct Phone {
    1: i32 id,
    2: string number,
    3: PhoneType type
}

service PhoneService {
    Phone findById(1: i32 id),
    list<Phone> findAll()
}
https://bit.ly/443gpse

```



<https://bit.ly/3JxrVUs>

02 Thrift in Zephyr



<https://zephyrproject.org/store/>



<https://bit.ly/3JvDyeP>




<https://bit.ly/3Xv60D9>

Google Summer of Code, 2022

- Thrift [ported to Zephyr](#) by [Young Mei](#) ([StdElectronics](#))
 - Zephyr [Blog Post](#)
 - C++ Code Generation
- Supported Features
 - [Binary Protocol](#)
 - [Compressed Binary Protocol](#)
 - [Zlib](#) Transport (via [uzlib](#), [muzic](#))
 - TLS sockets via [MbedTLS](#)
- All Zephyr Architectures Supported
- All Thrift features ([ThriftTest.thrift](#), Zephyr [testsuite](#))
- Multi-OS, Multi-language [samples](#)
- Released in [Zephyr v3.3.0](#) (-zlib) [EXPERIMENTAL]



<https://bit.ly/43aC5le>





How Can I use Thrift in Zephyr?

Nice features about Thrift is it is Transport Agnostic, but also ways to stitch things together

- [TFDTransport](#)
 - anything that can be represented as a file descriptor
 - socket (TCP, UDP, socketpair), text file,
 - In Linux, “everything is a file”
 - /dev/mem, /dev/gpu0, /dev/accel0, ...
- [TMemoryBuffer](#)
 - e.g. dual-port SRAM, DMA buffer...
 - In Zephyr, I wonder if it's possible to wrap an arbitrary *const struct device** with an integer file descriptor... 🤔
 - [fdtable](#)
 - What if I wanted to use Thrift over a UART? 🤔



Integration Challenges

- Partial C++ Support in Zephyr
 - See C++ Roadmap ([#45785](#))
 - Support for [std::thread](#) and [std::this_thread](#) ([#25569](#))
 - Get [std::mutex](#), [std::condition_variable](#), ... free
 - No async (due to missing `std::thread` support)
- C++ Requirements
 - [Zephyr SDK](#) needs [gthr-posix.h](#) ([#43729](#))
 - Support [pthread_create\(\)](#) dynamic stacks ([#25973](#))
 - [socketpair\(\)](#) (testing) needs subsys/net ([#51211](#))
- POSIX requirements
 - Well, dynamic thread stacks benefit all languages
 - Fix at kernel level
- Several workarounds (no Mutex synchronization)



Integration Challenges

- We used [eventfd\(\)](#) for notification in [TFDServer](#)
 - [eventfd_read\(\)](#) / [eventfd_write\(\)](#) deadlock ([#58790](#))
 - Fixed (with 10x performance improvement) 👍
- No Compression / Decompression subsys yet 🙄
- Code size is huge - can run / test in Qemu but very difficult to fit onto practically size
- SRAM requirements were quite large (for zlib) but reduced dramatically by Young during GSoC 2022



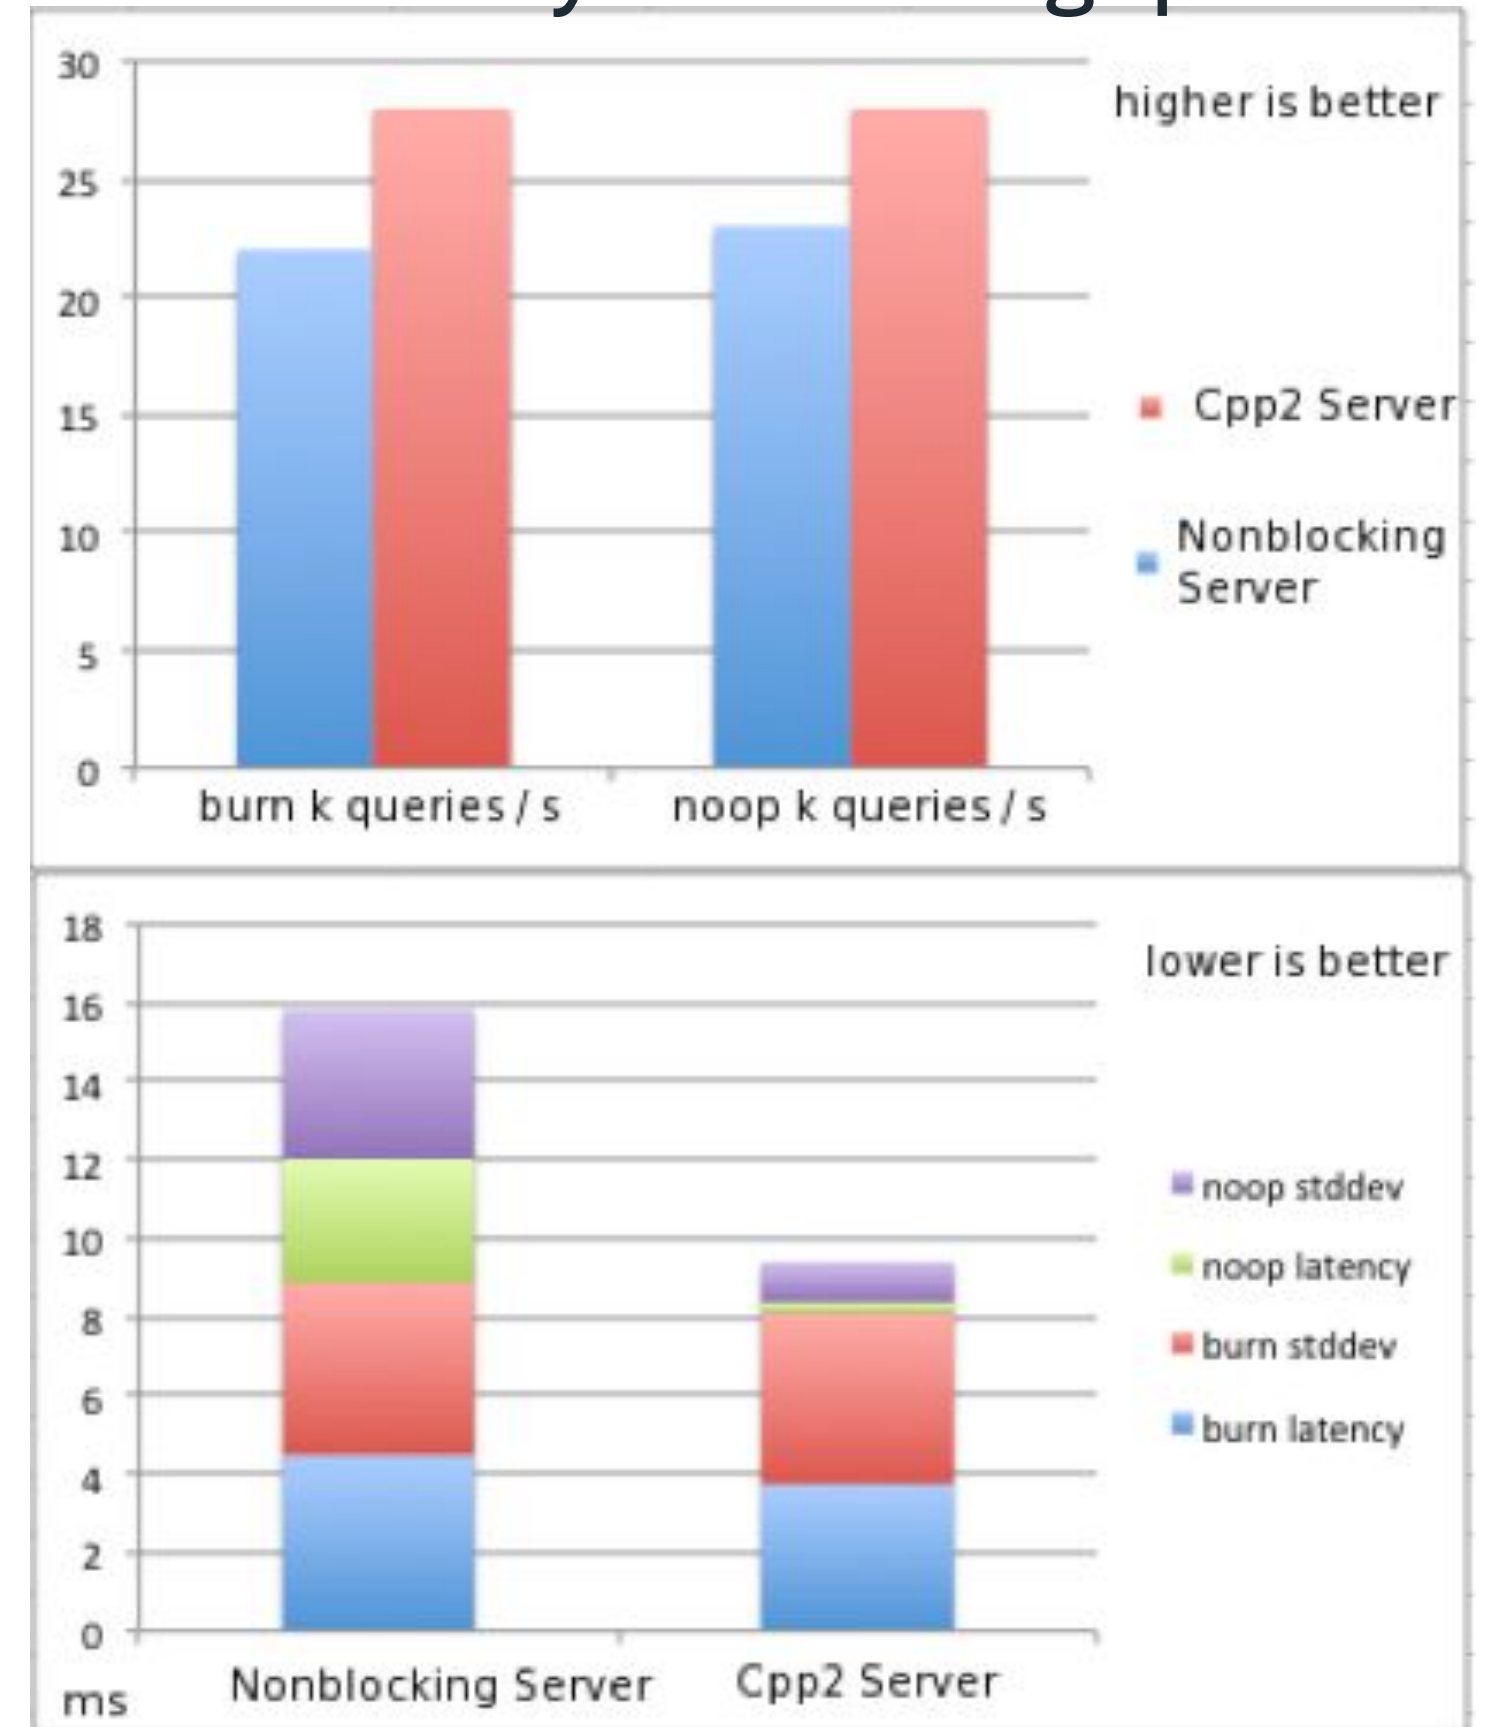
03 Status Update & Additional Work

fbthrift

- <https://github.com/facebook/fbthrift>
- Major C++ enhancements over Apache Thrift
- Fully asynchronous
- Server performance gains of 5-10k queries / s
- Switch to [Google Test Framework](#)
- Now uses [std::coroutine](#) (C++20)
- Dramatically lower latency across the board
- Code size reduction
- Use [std::string_view](#) (C++17) where possible
- Zero-copy buffers,





latency and throughput



<https://bit.ly/3r67vf6>

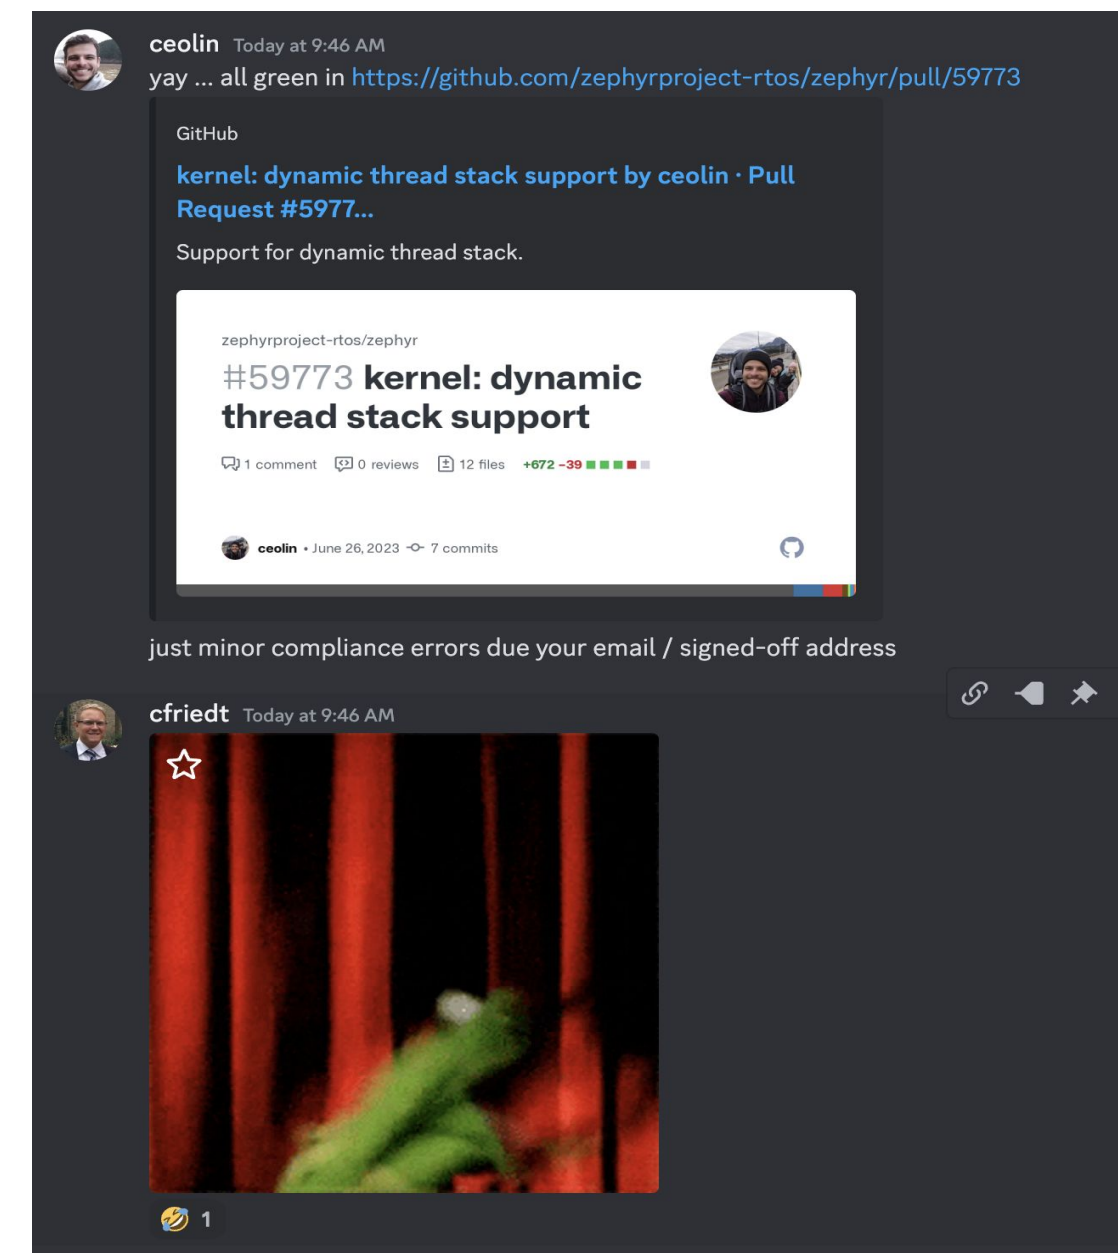
fbthrift

- Even more recent improvements..



Dynamic Kernel Thread Stacks Working!!11!

- After 4 years of “hacking in my spare time”
- Received this from [Flavio](#) literally this morning!!
- With / without MMU, with / without Userspace
- Collaboration from Andy, Flavio, Daniel, Anas and others at Intel, Stephanos Ioannis, Keith Packard



Questions? / Feedback

